

Community Regularization of Visually-Grounded Dialog

Akshat Agarwal, Swaminathan Gurumurthy, Vasu Sharma, Mike Lewis, Katia Sycara*

School of Computer Science
Carnegie Mellon University
Pittsburgh, PA

aa7@cmu.edu, sgurumur@andrew.cmu.edu, vasus@andrew.cmu.edu, ml@sis.pitt.edu, katia@cs.cmu.edu

ABSTRACT

The task of conducting visually grounded dialog involves learning goal-oriented cooperative dialog between autonomous agents who exchange information about a scene through several rounds of questions and answers in natural language. We posit that requiring artificial agents to adhere to the rules of human language, while also requiring them to maximize information exchange through dialog is an ill-posed problem. We observe that humans do not stray from a common language because they are social creatures who live in communities, and have to communicate with many people everyday, so it is far easier to stick to a common language even at the cost of some efficiency loss. Using this as inspiration, we propose and evaluate a multi-agent community-based dialog framework where each agent interacts with, and learns from, multiple agents, and show that this community-enforced regularization results in more relevant and coherent dialog (as judged by human evaluators) without sacrificing task performance (as judged by quantitative metrics).

KEYWORDS

Visual Dialog; Multi Agent Reinforcement Learning; Curriculum Learning; Emergent Communication

ACM Reference Format:

Akshat Agarwal, Swaminathan Gurumurthy, Vasu Sharma, Mike Lewis, Katia Sycara. 2019. Community Regularization of Visually-Grounded Dialog. In *Proc. of the 18th International Conference on Autonomous Agents and Multiagent Systems (AAMAS 2019), Montreal, Canada, May 13-17, 2019*, IFAAMAS, 9 pages.

1 INTRODUCTION

AI is increasingly becoming an important part of our daily lives, be it in the household, the workplace or in public places. In order for humans to interact with and understand the AI system, it needs to learn how to communicate with us about our environment using the languages that we speak. This requires the AI system to visually interpret the world, and communicate descriptions of the physical world. While such a task would have been considered impossible a few years ago, the recent progress in the fields of Computer Vision and Natural Language Processing, which are important building blocks for this task, have reinvigorated interest in the community. Several problems like image captioning ([12], [36], [30], [10], [20], [37]), image classification ([14], [29], [8], [34]), object detection

*AA, SG and VS contributed equally to this paper

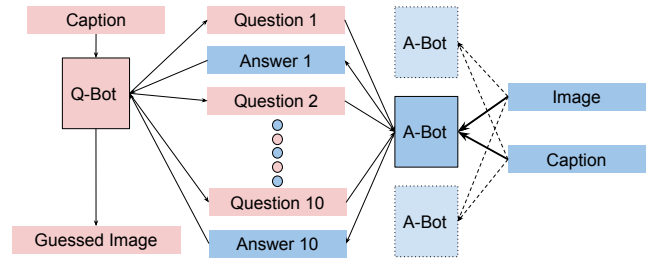


Figure 1: Multi-Agent (with 1 Q-Bot, 3 A-Bots) Dialog Framework. The diagram shows multiple A-bots interacting with a single Q-bot using natural language questions and answers. The Q-bot asks relevant questions to improve its understanding of the image and one of the randomly chosen A-bots answers the question. This ensures that the Q-bot can't overfit to responses provided by any one A-bot.

([17], [23], [24]), image segmentation ([18], [9], [22]), dialog ([28], [31], [5]), question answering ([38], [27], [35]) etc. have received immense amounts of attention from the research community. The paradigm of reinforcement learning has also shown promising results in several problems including learning to play Go [26] and Atari games [21], among others, at superhuman levels.

Capitalizing on the growth in all these different domains, it now seems plausible to build more advanced dialog systems capable of reasoning over multiple modalities while also learning from one another. Such systems will allow humans to have a meaningful dialog with intelligent systems containing visual as well as textual content. Use cases include assistive systems for the visually impaired, smart multimodal dialog agents (unlike current versions of Siri and Alexa which are primarily audio based and cannot make effective use of multimodal data) and even large scale visual retrieval systems. However, as these systems become more advanced, it will become increasingly common to have two agents interact with each other to achieve a particular goal [15]. We want these conversations to be interpretable to humans for the sake of transparency and ease of debugging. This motivates our work on goal-driven agents which interact in coherent language understandable to humans.

This paper presents work on Goal driven Visual Dialog Agents. Most prior work on visual dialog ([2], [4]) has approached the problem using supervised learning where, conditioned on the question-answer pair dialog history, a caption c and the image I , the agent is required to answer a given question q . The model is trained in a supervised learning framework using ground truth supervision from a human-human dialog dataset.

Some recent work [3] has tried to approach the problem using reinforcement learning, with two agents, namely the Question (Q-)

Bot and the Answer (A-) Bot. While the A-Bot still has the image, caption and the dialog history to answer any question, the Question Bot only has access to the caption and the dialog history. The two agents are initially trained with supervision using the VisDial v0.9 dataset [2], which consists of 80k images, each with a caption and 10 human generated question-answer pairs discussing the image. Under supervision, the agents are trained in an isolated manner to maximize the likelihood of generating the ground truth answers. The agents are then made to interact and talk to each other, with a common goal of trying to improve the Q-Bot's understanding of the image. The agents learn from their conversation with each other via reinforcement learning. While the supervised training *in isolation* helps the agents to learn to interpret the images and communicate information, it is the *interactive* training phase which leads to richer dialog with more informative questions and answers as the agents learn to adapt to each others' strengths and weaknesses. However, it is important to note that **the optimization problem in this conversational setting does not make the agents stick to the domain of grammatically correct and coherent natural language.** Indeed, **if the two agents are allowed to communicate and learn from each other for too long, they quickly start generating non-grammatical and semantically meaningless sentences.** While the generated sentences stop making sense to human observers, the two agents are able to understand each other much better, and the Q-Bot's understanding of the image improves. This is similar to how twins often develop a private language [25], an idiosyncratic and exclusive form of communication understandable only to them. **This, however, reduces transparency of the agents' dialog to any observer (human or AI), and is hence undesirable.** Prior work [2, 3] which has focused on improving performance as measured by the Q-Bot's image retrieval rank has suffered from incoherent dialog. We address this problem of improving the agents' performance while increasing dialog quality by taking inspiration from humans. **We observe that humans continue to speak in commonly spoken languages, and hypothesize that this is because they need to communicate with an entire community, and having a private language for each person would be extremely inefficient.** With this idea, **we let our agents learn in a similar setting, by making them talk to (ask questions of, get answers from) multiple agents, one by one.** We claim that if, instead of allowing a single pair of agents to interact, **we were to make the agents more social, and make them interact and learn from multiple other agents, they would be disincentivized to develop a private language, and would have to conform to the common language. We call this Community Regularization.**

In the subsequent sections we describe the Visual Dialog task and the neural network architectures of our Q-Bots and A-Bots in detail. We then describe **the training process of the agents sequentially: (a) in isolation (via supervision), (b) while interacting with one partner agent (via reinforcement), and (c) our proposed multi-agent setup where each agent interacts with multiple other agents (via reinforcement).** **We compare the performance of the agents trained in these different settings, both quantitatively using image retrieval ranks, and qualitatively evaluating the overall coherence, grammar and relevance of the dialog generated, as judged by impartial human evaluators.** We make the following contributions: We propose a multi-agent dialog setup in a natural

language setting and show that it results in community regularization which ensures that the interactions between the agents remain grounded in the rules and grammar of natural language, are coherent and human-interpretable without compromising on task performance. The code can be found in the following repository <https://github.com/agakshat/visualdialog-pytorch>.

2 PROBLEM STATEMENT

We begin by defining the problem of Visually Grounded Dialog for the co-operative image guessing game on the VisDial dataset.

Players and Roles: The game involves two collaborative agents – a question bot (Q-bot) and an answer bot (A-bot). The A-bot has access to an image and caption, while the Q-bot has access to the image's caption, but not the image itself. Both the agents share a common objective, which is for the Q-bot to form a good "mental representation" of the unseen image which can be used to retrieve, rank or generate that image. This is facilitated by the exchange of 10 pairs of questions and answers between the two agents, using a shared vocabulary, where the Q-bot asks the A-bot a question about the image, and the A-bot answers the question, hence improving the Q-Bot's understanding of the image scene.

General Game Objective: At each round, in addition to communicating with the A-bot, the Q-bot also provides the learning algorithm with its best estimate y_t of the unknown image I based only on the dialog history and caption. Both agents receive a common reward from the environment which is inversely proportional to the error in this description under some metric $L(y_t, y_{gt})$. We note that this is a general setting where the 'description' y_t can take on varying levels of specificity – from image feature embeddings extracted by deep neural networks to textual descriptions and pixel-level image re-generations.

Specific Instantiation: In our experiments, we focus on the setting where the Q-bot is tasked with estimating a vector embedding of the image I , which is later used to retrieve a similar image from the dataset. Given a feature extractor (say, a pretrained CNN model like VGG [29]), the target 'description' y_{gt} of the image, can be obtained by simply forward propagating through the VGG model, without the requirement of any human annotation. Reward/error can be measured by the Euclidean distance between the target description y_{gt} and the predicted description y_t , and any image may be used as the visual grounding for a dialog. Thus, an unlimited number of games may be simulated without human supervision, motivating the use of reinforcement learning in this framework.

Our primary focus for this work is to ensure that the agents' dialog remains coherent and understandable while also being informative and improving task performance. For concreteness, consider an example of dialog that is informative yet incoherent: **question:** "do you recognize the guy and age is the adult?", **answered with:** "you couldn't be late teens, his". The example shows that the bots try to extract and convey as much information as possible in a single question/answer (sometimes by incorporating multiple questions or answers into a single statement). But in doing so they lose basic semantic and syntactic structure. We also provide a sample of the dialogs in Figure 2.

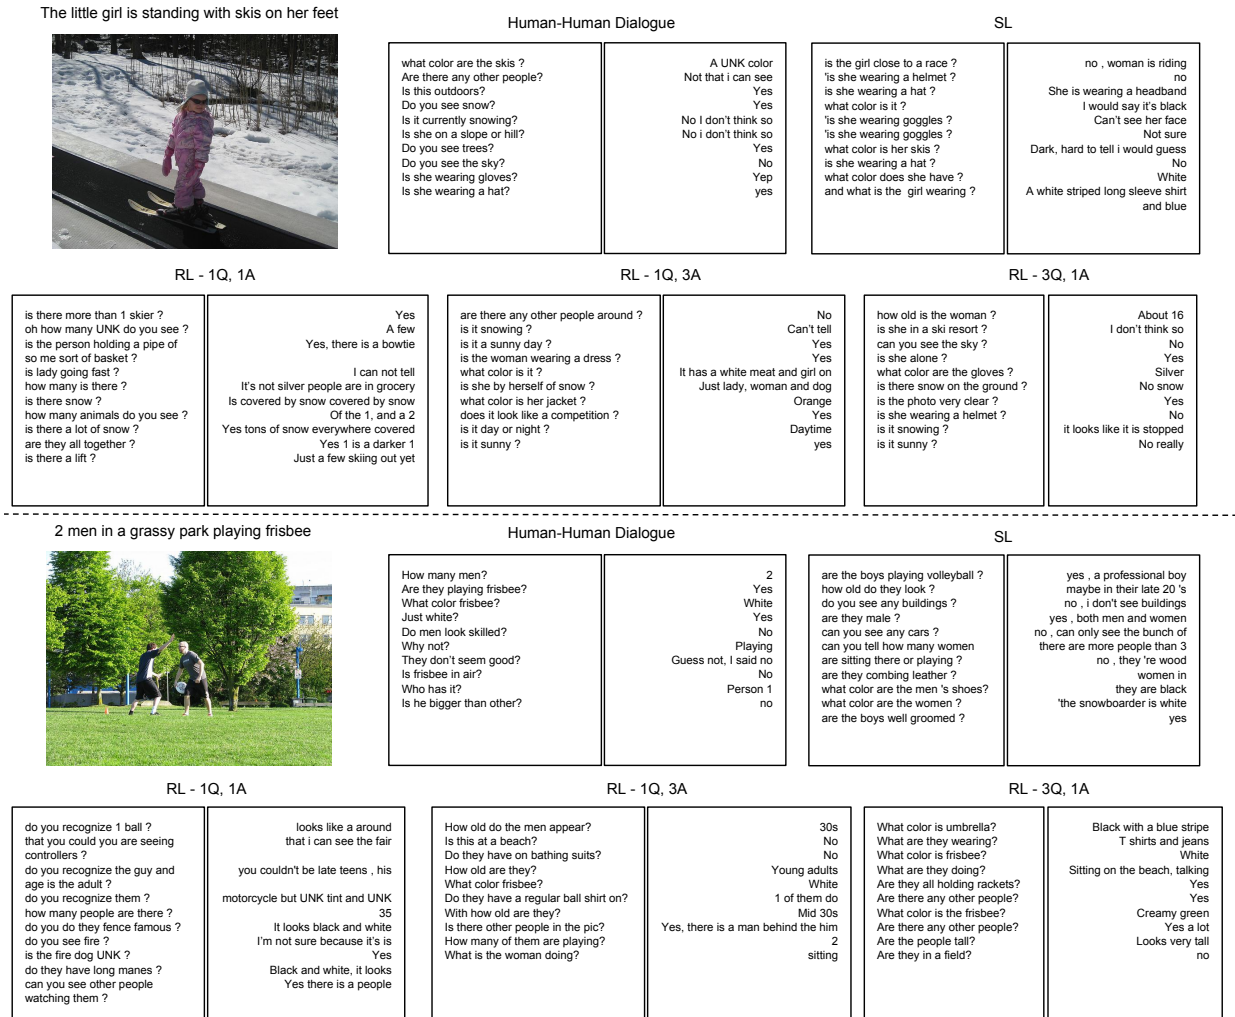


Figure 2: A randomly selected image from the VisDial dataset followed by the ground truth (human) and generated dialog about that image for each of our 4 systems (SL, RL-1Q,1A, RL-1Q,3A, RL-3Q,1A). These examples were 2 of the 102 used in the human evaluation results shown in Table 2

3 RELATED WORK

Most of the major works which combine vision and language have traditionally focused on the problem of image captioning ([12], [36], [30], [10], [20], [37]) and visual question answering ([1], [39], [7]). The problem of visual dialog is relatively new and was first introduced by Das et al. [2] who also created the VisDial dataset to advance the research on visually grounded dialog. The dataset was collected by pairing two annotators on Amazon Mechanical Turk to chat about an image. They formulated the task as a ‘multi-round’ VQA task and evaluated individual responses at each round in an image guessing setup. In subsequent work by Das et al. [3] they proposed a reinforcement learning based setup where they allowed the Question bot and the Answer bot to have a dialog with each other with the goal of correctly predicting the image unseen to the

Question bot. However, in their work they noticed that the reinforcement learning based training quickly led the bots to diverge from natural language. In fact Kottur et al. [13] recently showed that language emerging from two agents interacting with each other might not even be interpretable or compositional. We use community regularization to alleviate this problem. Recent work has also proposed using such goal driven dialog agents for other related tasks including negotiation [16] and collaborative drawing [11]. We believe that our work can easily extend to those settings as well. Lu et al. [19] proposed a generative-discriminative framework for visual dialog where they train only an answer bot to generate informative answers for ground truth questions. These answers were then fed to a discriminator, which was trained to rank the generated answer among a set of candidate answers. This is a major restriction of their model as it can only be trained when this additional information of candidate answers is available, which

restricts it to a supervised setting. Furthermore, since they train only the answer bot and have no question bot, they cannot simulate an entire dialog which also prevents them from learning by self-play via reinforcement. Wu et al. [33] further improved upon this generative-discriminative framework by formulating the discriminator as a more traditional GAN [6], where the adversarial discriminator is tasked to distinguish between human generated and machine generated dialogs.

4 AGENT ARCHITECTURES


We describe all the different components of the agent architectures in this section. Note that the overall architecture is mostly borrowed from Das et al. [3], Lu et al. [19] with slight modifications to individual units and an additional caption encoder. We explain these modifications in detail in this section. We would like to stress that these changes are not the main contribution of our paper. The main contribution is the Multi-agent dialog framework described in section 5.3.

4.1 Question Bot Architecture


The question bot architecture we use is inspired by the answer bot architecture in Das et al. [3], Lu et al. [19] but the individual units have been modified to provide more useful representations. Similar to the original architecture, **our Q-Bot, shown in Fig. 3a, also consists of 4 parts, (a) fact encoder, (b) state-history encoder, (c) question decoder and (d) image regression network.**

- (1) **Fact Encoder:** The fact encoder is a unidirectional LSTM which is given the previous question-answer pair (q_{t-1}, a_{t-1}) as input. The LSTM generates a fact embedding $F_t \in R^{512}$.
- (2) **State/History Encoder:** We modify the state-history encoder to incorporate a two-level hierarchical encoding of the dialog. The encoder first computes the fact embeddings $H_t^Q = (F_1, F_2, F_3 \dots F_{t-1})$, using an LSTM akin to the fact encoder described above. We pass these embeddings and F_t computed by the Fact Encoder through a fully connected layer, generating attention weights which are used to attend over H_t^Q , producing the history embedding $S_t^Q \in R^{512}$. Notice that this results in a two-level hierarchical encoding of the dialog $(q_t, a_t) \rightarrow F_t$ and $(F_1, F_2, F_3, \dots F_t) \rightarrow S_t^Q$.
- (3) **Caption Encoder:** This is a unidirectional LSTM which is given the image caption c as input. The LSTM generates a caption embedding $C^Q \in R^{512}$.
- (4) **Feature Regression Network:** $\{F_t^Q, S_t^Q, C^Q\}$ are concatenated to produce an embedding E_t^Q . This is passed through 2 fully connected layers with dropout to produce \hat{y}_t from the current encoded state $\hat{y}_t = f(S_t^Q)$.
- (5) **Question Decoder:** The hidden state of this LSTM is initialized with the hidden state of the fact encoder. E_t^Q is passed through a fully connected layer to generate e_t^Q , which is used to update the hidden state of the LSTM of the question decoder. The question q_t is then generated by sequentially sampling words (either via teacher forcing during supervised pretraining or via autoregressive generation during RL and evaluation).

Note that we use a dropout of 0.5 in all the LSTMs during training. All LSTM hidden layers sizes are 512, and the image embedding size is 4096. The input word embedding size is 300.



(a) Encoder architecture for Q-Bot



(b) Encoder architecture for A-Bot

Figure 3: Agent Encoder Architectures

4.2 Answer Bot Architecture

The architecture for A-Bot, also inspired from Lu et al. [19], shown in Fig. 3b, is similar to that of the Q-Bot. It has 3 components: (a) question encoder, (b) state-history encoder and (c) answer decoder.

- (1) **Question Encoder:** The question encoder is a unidirectional LSTM which is given the current question q_t generated by the Q-Bot as input. The LSTM generates a question embedding $Q_t^A \in R^{512}$.
- (2) **State/History/Image Encoder:** The encoder first computes the fact embeddings $H_t^A = (F_1, F_2, F_3 \dots F_{t-1})$, using an LSTM akin to the fact encoder described above. By passing these embeddings and the Q_t^A computed by the Question Encoder through a fully connected layer, attention weights are generated which are used to attend over H_t^A , producing the history embedding $S_t^A \in R^{512}$. Notice that this results in a two-level hierarchical encoding of the dialog $(q_t, a_t) \rightarrow F_t$ and $(F_1, F_2, F_3 \dots F_t) \rightarrow S_t^A$. $\{Q_t^A, S_t^A, y_{gt}\}$ are then concatenated to produce an embedding E_t^A .
- (3) **Answer Decoder:** The hidden state of this LSTM is initialized with the hidden state of the question encoder. E_t^A is passed through a fully connected layer to generate e_t^A , which is used to update the hidden state of the LSTM of the answer decoder. The answer a_t is then generated by sequentially sampling words (either via teacher forcing during supervised pretraining or via autoregressive generation during RL and evaluation).

5 TRAINING

We follow the training process proposed in Das et al. [3]. Two agents, a Q-Bot and an A-Bot are first trained in isolation, by supervision from the VisDial dataset. After this supervised pretraining for 15 epochs over the data, we smoothly transition the agents to learn from each other via reinforcement learning. The individual phases of training will be described in more detail below. Note that the key novelty of the work is the multi-agent dialog framework proposed in Section 5.3

5.1 Supervised pre-training

In the supervised part of training, both the Q-Bot and A-Bot are trained separately, using a Maximum Likelihood Estimation (MLE) loss computed using the ground truth questions and answers, respectively, for every round of dialog. The Q-Bot simultaneously optimizes another objective: minimizing the Mean Squared Error (MSE) loss between the true (y_{gt}) and predicted (y_t) image embeddings. The ground truth dialogs and image embeddings are from the VisDial dataset.

Given the true dialog history, image features and a question from the dataset, the A-Bot generates an answer to that question. Given the true dialog history and the previous question-answer pair from the dataset, the Q-Bot is made to generate the next question to ask the A-Bot. Both agents receive only ground truth questions and answers, never what the other agent generated - so the two agents never actually interact during this phase of training. However, there are multiple problems with this training scheme. First, MLE is known to result in models that generate repetitive dialogs and often produce generic responses. Second, since the agents are never allowed to interact during training, they end up encountering out-of-distribution questions and answers when made to interact during evaluation, which reduces the task performance. This can be observed in Figure 4. The performance of the agents trained via supervised learning dips after each successive dialog round.

5.2 Reinforcement Learning Setup

To alleviate the issues pointed out with supervised training, we let the two bots interact with each other via self-play (no ground-truth except images and captions). This interaction is also in the form of questions asked by the Q-Bot, and answered in turn by the A-Bot, using a shared vocabulary. The state space is partially observed and asymmetric, with the Q-Bot observing $\{c, q_1, a_1 \dots q_{10}, a_{10}\}$ and the A-Bot observing the same, plus the image I . Here, c is the caption, and q_i, a_i is the i^{th} dialog pair exchanged where $i = 1 \dots 10$. The action space for both bots consists of all possible output sequences of a specified maximum length (Q-Bot: 16, A-Bot: 9) under a fixed vocabulary (size 8645). Each action involves predicting words sequentially until a stop token is predicted, or the generated statement reaches the maximum length. Additionally, Q-Bot also produces a guess of the visual representation of the input image (VGG fc-7 embedding of size 4096). Both Q-Bot and A-Bot share the same objective and get the same reward to encourage cooperation. Information gain in each round of dialog is incentivized by setting the reward as the change in distance of the predicted image embedding to the ground-truth image representation. This means that a QA pair is of high quality only if it helps the Q-Bot make a better

prediction of the image representation. Both policies are modeled by neural networks, as discussed in Section 4.

A dialog round at time t consists of the following steps: 1) the Q-Bot, conditioned on the state encoding, generates a question q_t , 2) A-Bot updates its state encoding with q_t and then generates an answer a_t , 3) Both Q-Bot and A-Bot encode the completed exchange as a fact embedding, 4) Q-Bot updates its state encoding to incorporate this fact and finally 5) Q-Bot predicts the image representation for the unseen image conditioned on its updated state encoding.

Similar to Das et al. [2], we use the REINFORCE [32] algorithm that updates policy parameters in response to experienced rewards. The per-round rewards that are used to calculate the discounted returns follow:

$$r_t(s_t^Q, (q_t, a_t, y_t)) = l(y_{t-1}, y^{gt}) - l(y_t, y^{gt}) \quad (1)$$

This reward is positive if the distance between image representation generated at time t is closer to the ground truth than the representation at time $t - 1$, hence incentivizing information gain at each round of dialog. The REINFORCE update rule ensures that a (q_t, a_t) exchange that is informative has its probabilities pushed up. Do note that the image feature regression network f is trained directly via supervised gradient updates on the L-2 loss.

However, as noted above, this RL optimization problem is ill-posed, since rewarding the agents for information exchange does not motivate them to stick to the rules and conventions of the English language. Thus, we follow an elaborate curriculum scheme described in [2]. Specifically, for the first K rounds of dialog for each image, the agents are trained using supervision from the VisDial dataset. For the remaining $10-K$ rounds, however, they are trained via reinforcement learning. K starts at 9 and is linearly annealed to 0 over 10 epochs. Despite these modifications the bots are still observed to diverge from natural language and produce non-grammatical and incoherent dialog. Thus, we propose a multi bot architecture to help the agents interact in diverse and coherent, yet informative, dialog.

5.3 Multi-Agent Dialog Framework (MADF)

In this section we describe our proposed Multi-Agent Dialog architecture in detail. We claim that if, instead of allowing a single pair of agents to interact, we were to make the agents more social, and make them interact and learn from multiple other agents, they would be disincentivized to develop a private language, and would have to conform to the common language. We call this Community Regularization.

In particular, we create either multiple Q-bots to interact with a single A-bot, or multiple A-bots to interact with a single Q-bot. All these agents are initialized with the learned parameters from the supervised pretraining as described in Section 5.1. Then, for each batch of images from the VisDial dataset, we randomly choose a Q-bot to interact with the A-bot, or randomly choose an A-bot to interact with the Q-bot, as the case may be. The two chosen agents then have a complete dialog consisting of 10 question-answer pairs about each of those images, and update their respective weights based on the rewards received (as per Equation 1) during the conversation, using the REINFORCE algorithm. This process is repeated for each batch of images, over the entire VisDial dataset. It is important to note that histories are *not shared* across batches. MADF

Table 1: Comparison of answer retrieval metrics with previously published work. SL has the best scores. The scores drop drastically in RL-1Q,1A, but MADF agents (RL-3Q,1A and RL-1Q,3A) are able to retain the same language quality as the SL agent.

Model	MRR	Mean Rank	R@10
Answer Prior [2]	0.3735	26.50	53.23
MN-QIH-G [2]	0.5259	17.06	68.88
HCIAE-G-DIS [19]	0.547	14.23	71.55
Frozen-Q-Multi [3]	0.437	21.13	60.48
CoAtt-GAN [33]	0.5578	14.4	71.74
SL(Ours)	0.610	5.323	72.68
RL - 1Q,1A(Ours)	0.459	7.097	72.34
RL - 1Q,3A(Ours)	0.601	5.495	72.48
RL - 3Q,1A(Ours)	0.590	5.56	72.61

can be understood in detail using the pseudocode in Algorithm 1.

Connection to Regularization: It is interesting to note that the MADF setting can actually be seen as a regularizer for the model. To establish this more formally, we look at the total loss minimized by each agent. The Total loss (TL) being minimized during the RL phase = $A1_t + A2_t$ at time t , where $A1_t$ is negative of the RL reward as described in section 5.2, and $A2_t$ is the L-2 loss between predicted and true image embeddings. Consider a setting with N Abots and 1 QBot. The $A2_t$ Loss can be written as:

$$A2_t = \sum_{i=1}^N (y_t^{(i)} - y^{gt})^2 = (y_t^{(1)} - y^{gt})^2 + \sum_{i=2}^N (y_t^{(i)} - y^{gt})^2$$

From the equation, we observe that $A2_t$ is a sum of 2 terms, where the first term is the standard regression loss which would apply for the 1Q,1A case. The second term can be viewed as a regularization imposed by pairing the other A-bots with the Q-bot, hence we can rewrite $A2$ as:

$$A2_t = (y_t^{(1)} - y^{gt})^2 + R_{A2_t}(\theta) \quad (2)$$

where $R_{A2_t}(\theta)$ represents regularization imposed by the other agents. Similarly, $A1_t$ can also be broken down into a likelihood and a regularization term as follows:

$$A1_t = -G_t^{(1)} \log \pi(q_t^{(1)}, a_t^{(1)}) + R_{A1_t}(\theta) \quad (3)$$

where $G_t^{(1)}$ is the monte-carlo return calculated using the first pair of agents at time t . Thus, both the terms in the total loss can be broken down into a loss term akin to the 1Q, 1A case and a regularization term. This regularization term comes from the regularization imposed by pairing each Q-Bot with multiple A-bots or vice versa. This clearly shows that the multi-bot framework can be seen as a form of regularization. In the experiments we show that the regularization helps with the language quality by ensuring that the bots don't deviate much from natural language.

6 EXPERIMENTS AND RESULTS

6.1 Dataset description

We use the VisDial 0.9 dataset for our task introduced by Das et al. [2]. The data is collected using Amazon Mechanical Turk by pairing 2 annotators and asking them to chat about the image as a multi round VQA setup. One of the annotators acts as the questioner

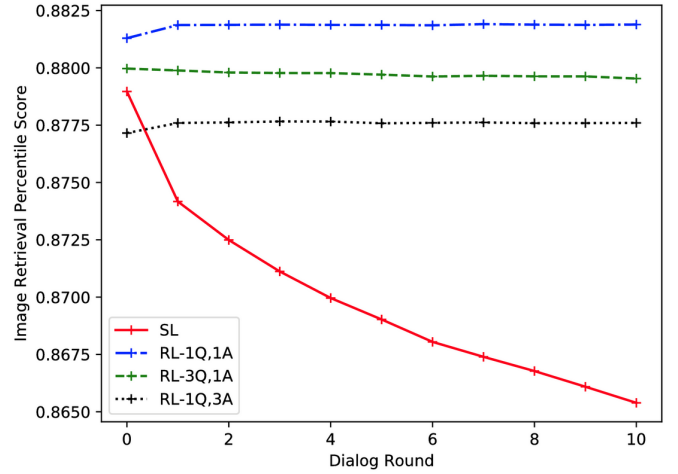


Figure 4: Comparison of Task Performance: Image Retrieval Percentile scores. This refers to the percentile scores of the ground truth image compared to the entire test set of 40k images, as ranked by distance from the Q-Bot's estimate of the image. The X-axis denotes the dialog round number (from 1 to 10), while the Y-axis denotes the image retrieval percentile score. The percentile score decreases monotonically for SL, but is stable for all versions using RL. This shows that the MADF agents are able to capitalize on the benefits of interactive learning.

and has access to only the caption of the image and has to ask questions from the other annotator who acts as the 'answerer' and must answer the questions based on the visual information from the actual image. This dialog repeats for 10 rounds at the end of which the questioner has to guess what the image was. We perform our experiments on VisDial v0.9 (the latest available release) containing 83k dialogs on COCO-train and 40k on COCO-val images, for a total of 1.2M dialog question-answer pairs. We split the 83k into 82k for train, 1k for validation, and use the 40k as test, in a manner consistent with [2]. The caption is considered to be the first round in the dialog history.

6.2 Evaluation Metrics

We evaluate the performance of our model's individual responses by using 4 metrics, proposed by [3]: 1) **Mean Reciprocal Rank (MRR)**, 2) **Mean Rank**, 3) **Recall@10** and 4) **Image Retrieval Percentile**. Mean Rank and MRR compute the average rank (and its reciprocal, respectively) assigned to the ground truth answer, over a set of 100 candidate answers for each question (also averaged over all the 10 rounds). Recall@10 computes the percentage of answers with rank less (better) than 10. Intuitively all these language metrics are trying to measure similar things, i.e, how highly does the model rank the ground truth answer over a set of 100 candidate responses. Thus, if the model gives a lower rank to the ground truth answer, then we can say that the model is highly likely to produce the ground truth response to the question. But at the same time they do have some qualitative differences. For example, MRR and Rank@10 are more robust to outliers but Mean Rank is not (but it

Algorithm 1 Multi-Agent Dialog Framework (MADF)

```

1: procedure MULTIBOTTRAIN
2:   while train_iter < max_train_iter do                                     ▶ Main Training loop over batches
3:      $Q_{bot} \leftarrow \text{random\_select}(Q_1, Q_2, Q_3 \dots Q_q)$ 
4:      $A_{bot} \leftarrow \text{random\_select}(A_1, A_2, A_3 \dots A_a)$ 
5:      $history \leftarrow (0, 0, \dots 0)$ 
6:      $fact \leftarrow (0, 0, \dots 0)$ 
7:      $\Delta image\_pred \leftarrow 0$ 
8:      $Qz_1 \leftarrow \text{Ques\_enc}(Q_{bot}, fact, history, caption)$ 
9:     for t in 1:10 do                                                         ▶ Have 10 rounds of dialog
10:       $ques_t \leftarrow \text{Ques\_gen}(Q_{bot}, Qz_t)$ 
11:       $Az_t \leftarrow \text{Ans\_enc}(A_{bot}, fact, history, image, ques_t, caption)$ 
12:       $ans_t \leftarrow \text{Ans\_gen}(A_{bot}, Az_t)$ 
13:       $fact \leftarrow [ques_t, ans_t]$ 
14:       $history \leftarrow \text{concat}(history, fact)$ 
15:       $Qz_t \leftarrow \text{Ques\_enc}(Q_{bot}, fact, history, caption)$ 
16:       $image\_pred \leftarrow \text{image\_regress}(Q_{bot}, fact, history, caption)$ 
17:       $R_t \leftarrow (target\_image - image\_pred)^2 - \Delta image\_pred$ 
18:       $\Delta image\_pred \leftarrow (target\_image - image\_pred)^2$ 
19:    end for
20:     $\Delta(w_{Qbot}) \leftarrow \frac{1}{10} \sum_{t=1}^{10} \nabla_{\theta_{Qbot}} [G_t \log p(ques_t, \theta_{Qbot}) - \Delta image\_pred]$ 
21:     $\Delta(w_{Abot}) \leftarrow \frac{1}{10} \sum_{t=1}^{10} G_t \nabla_{\theta_{Abot}} \log p(ans_t, \theta_{Abot})$ 
22:     $w_{Qbot} \leftarrow w_{Qbot} + \Delta(w_{Qbot})$ 
23:     $w_{Abot} \leftarrow w_{Abot} + \Delta(w_{Abot})$ 
24:  end while
25: end procedure

```

▶ Either q or a is equal to 1
 ▶ History initialized with zeros
 ▶ Fact encoding initialized with zeros
 ▶ Tracks change in Image Embedding
 ▶ Fact encoder stores the last dialog pair
 ▶ History stores all previous dialog pairs
 ▶ REINFORCE and Image Loss update for Qbot
 ▶ REINFORCE update for Abot

Table 2: Human Evaluation Results - Mean Rank (Lower is better) : Results show that RL-3Q,1A outperforms RL-1Q,3A for A-relevance and overall coherence but otherwise SL (Supervised Learning), RL-1Q,3A, and RL-3Q,1A showed equivalent performance indicating that community regularization can effectively eliminate any losses to human intelligibility introduced through RL.

	Metric	N	SK	RL 1Q,1A	RL 1Q,3A	RL 3Q,1A
1	Question Relevance	49	1.97	3.57	2.20	2.24
2	Question Grammar	49	2.16	3.67	2.24	1.91
3	Overall Dialog Coherence: Q	49	2.08	3.73	2.34	1.83
4	Answer Relevance	53	2.09	3.77	2.28	1.84
5	Answer Grammar	53	2.20	3.75	2.05	1.98
6	Overall Dialog Coherence: A	53	2.09	3.64	2.35	1.90

is more interpretable). The fourth metric, i.e., the Image Retrieval percentile, is different from the first 3 metrics. It is a measure of how close the image prediction generated by the Q-bot is to the ground truth. All the images in the test set are ranked according to their distance from the predicted image embedding, and the rank of the ground truth embedding is used to calculate the image retrieval percentile. This gives a measure of the quality of the information exchange. All results are reported after 15 epochs of supervised learning and 10 epochs of curriculum learning as described in Section 5. Consequently, the training time of all 3 systems are equal.

Table 1 compares the Mean Rank, MRR, and Recall@10 of our agent architecture and dialog framework (below the horizontal line) with previously proposed architectures (above the line). SL refers to the agents after only the isolated, supervised training of Section 5.1. RL-1Q,1A refers to a single, idiosyncratic pair of agents trained via reinforcement as in Section 5.2. RL-1Q,3A and RL-3Q,1A refer to social agents trained via our Multi-Agent Dialog framework

in Section 5.3, with 1Q,3A referring to 1 Q-Bot and 3 A-Bots, and 3Q,1A referring to 3 Q-Bots and 1 A-Bot.

It can be seen that our agent architectures clearly outperform all previously published results using generative architectures in MRR, Mean Rank and R@10. This indicates that our approach produces consistently good answers (as measured by MRR, Mean Rank and R@10). It is important to note that the point here is not to demonstrate the superiority of our architecture compared to other architectures. The point here is instead to show that the MADF framework (RL-3Q,1A and RL-1Q,3A) is able to maintain the same language quality as the supervised agent while improving the image retrieval scores. In fact, community regularization (in the form of the proposed MADF setup) can be integrated with any of the visual dialog algorithms in Table 1. Notice that SL has the best scores. The scores drop drastically in RL-1Q,1A, but RL-3Q,1A and RL-1Q,3A obtain scores comparable to SL. This shows that the agents trained by MADF are able to maintain the language quality of SL

agents without sacrificing much on the task performance (image retrieval percentile). Fig. 4 shows the change in image retrieval percentile scores over dialog rounds. The percentile score decreases monotonically for SL, however it is stable for all versions using RL. The decrease in image retrieval score over dialog rounds for SL is because the test set questions and answers are not perfectly in-distribution (compared to the training set), and the SL system can't adapt to these samples as well as the systems trained with RL. As the dialog rounds increase, the out-of-distribution nature of dialog exchange increases, hence leading to a decrease in SL scores. Interestingly, despite having strictly more information in later rounds, the scores of RL agents do not increase - which we think is because of the nature of recurrent networks to forget.

The results in Fig. 4 and Table 1 show that the MADF setup allows the agents to achieve consistent task performance without sacrificing on language quality. We further support this claim in the next section where we show that human evaluators rank the language quality of MADF agents to be much better than the agents trained via reinforcement without community regularization.

6.3 Human Evaluation

There are no quantitative metrics to comprehensively evaluate dialog quality, hence we do a human evaluation of the generated dialog. There are 6 metrics we evaluate on: 1) Q-Bot Relevance, 2) Q-Bot Grammar, 3) A-Bot Relevance, 4) A-Bot Grammar, 5) Q-Bot Overall Dialog Coherence and 6) A-Bot Overall Dialog Coherence. We evaluate 4 Visual Dialog systems, trained via: 1) **Supervised Learning (SL)**, 2) **Reinforce for 1 Q-Bot, 1 A-Bot (RL-1Q,1A)**, 3) **Reinforce for 1 Q-Bot, 3 A-Bots (RL-1Q,3A)** and 4) **Reinforce for 3 Q-Bots, 1 A-Bot (RL-3Q,1A)**. We asked a total of 61 people to evaluate the 10 QA-pairs generated by each system for a total of 102 randomly chosen images, requiring them to give an ordinal ranking (from 1 to 4) for each metric. All the evaluators were provided with the caption from the dataset. Evaluators taking the perspective of the A-Bot were provided with the image and asked to evaluate answer relevance and grammar, while those taking the perspective of the Q-Bot evaluated question relevance and grammar. Both groups rated dialogs for overall coherence. Table 2 contains the average ranks obtained on each metric (lower is better).

The results convincingly validate our hypothesis that having multiple A-Bots/Q-Bots improves the language quality as compared with single Q-Bot and A-Bot. Kruskal-Wallis tests found strong differences among rankings ($p < .0001$) across all measures. Pairwise comparisons using the Mann-Whitney U test found a consistent pattern in which RL 1Q,1A performed substantially worse than other methods across all measures: for

- (1) **Q-relevance**: SL: $U=348$, $p < .0001$; RL-1Q3A: $U=2235$, $p < .0001$; RL-3Q1A $U=2209$, $p < .0001$;
- (2) **Q-grammar**: SL: $U=319$, $p < .0001$; RL-1Q3A $U=2280$, $p < .0001$; RL-3Q1A $U=2221$, $p < .0001$;
- (3) **A-relevance**: SL $U=256$, $p < .0001$; RL-1Q3A $U=2741$, $p < .0001$; RL-3Q1A $U=2909$, $p < .0001$;
- (4) **A-grammar**: SL $U=305$, $p < .0001$; RL-1Q3A $U=2857$, $p < .0001$; RL-3Q1A $U=2673$, $p < .0001$;
- (5) **Overall (both groups)**: SL $U=1206$, $p < .0001$; RL-1Q3A $U=9458$, $p < .0001$; RL-3Q1A $U=10052$, $P < .0001$.

Results showed that RL 3Q,1A outperformed RL 1Q,3A for A-relevance $U=1889$, $p < .02$ and overall coherence $U=6543$, $p < .006$ but otherwise SL, RL-1Q,3A, and RL-3Q,1A showed equivalent performance indicating that community regularization can effectively eliminate any losses to human intelligibility introduced through reinforcement learning. These results further support the claims made in the previous section that the MADF setup allows the agents to show consistent task performance (image retrieval percentile) while maintaining the language quality of the supervised agents.

We show a couple of randomly chosen examples from the set shown to the human evaluators in Fig. 2. The trends observed in the scores given by human evaluators are also clearly visible in this example. MADF agents are able to model the human responses much better than RL 1Q,1A and are about as well as (if not better) than SL trained agents. It can also be seen that although the RL-1Q,1A system has greater diversity in its responses, the quality of those responses is greatly degraded, with the A-Bot's answers especially being both non-grammatical and irrelevant.

7 DISCUSSION AND CONCLUSION

In this paper we propose a novel community regularization technique, the Multi-Agent Dialog Framework (MADF), to improve the dialog quality of artificial agents. We show that training 2 agents with supervised learning does not ensure good task performance (measured by the image retrieval percentile scores) at test time, and it only deteriorates as the agents exchange more information about the image. We hypothesize that this is because the agents were trained in isolation and never allowed to interact during supervised learning, which leads to failure during testing when they encounter out of distribution samples (generated by the other agent, instead of ground truth) for the first time. We show how allowing a single pair of agents to interact and learn from each other via reinforcement learning dramatically improves their percentile scores, which additionally does not deteriorate over multiple rounds of dialog, since the agents have interacted with one another and been exposed to the other's generated questions or answers. However, in an attempt to improve task performance, the agents end up developing their own private language which does not adhere to the rules and conventions of human languages. As a result, the dialog system loses interpretability and sociability. To alleviate this issue, we propose a multi-agent dialog framework to provide regularization. In this framework, a single A-Bot is allowed to interact with multiple Q-Bots and vice versa. We first show mathematically that this has direct connections to regularization. We then back it up with multiple empirical experiments including a human evaluation study, and show that MADF leads to significant improvements in dialog quality measured by relevance, grammar and overall coherence, without compromising the task performance.

ACKNOWLEDGMENTS

This work has been funded by AFOSR award FA9550-15-1-0442 and AFOSR/AFRL award FA9550-18-1-0251. We would like to thank William Guss, Abhishek Das, Satwik Kottur and Fei Fang for their insightful and fruitful discussions and feedback.

REFERENCES

- [1] Aishwarya Agrawal, Jiasen Lu, Stanislaw Antol, Margaret Mitchell, C. Lawrence Zitnick, Devi Parikh, and Dhruv Batra. 2017. VQA: Visual Question Answering. *Int. J. Comput. Vision* 123, 1 (May 2017), 4–31. <https://doi.org/10.1007/s11263-016-0966-6>
- [2] Abhishek Das, Satwik Kottur, Khushi Gupta, Avi Singh, Deshraj Yadav, José M. F. Moura, Devi Parikh, and Dhruv Batra. 2016. Visual Dialog. *CoRR* abs/1611.08669 (2016). [arXiv:1611.08669](http://arxiv.org/abs/1611.08669) <http://arxiv.org/abs/1611.08669>
- [3] Abhishek Das, Satwik Kottur, José M. F. Moura, Stefan Lee, and Dhruv Batra. 2017. Learning Cooperative Visual Dialog Agents with Deep Reinforcement Learning. *CoRR* abs/1703.06585 (2017). [arXiv:1703.06585](http://arxiv.org/abs/1703.06585) <http://arxiv.org/abs/1703.06585>
- [4] Harm de Vries, Florian Strub, Sarath Chandar, Olivier Pietquin, Hugo Larochelle, and Aaron C. Courville. 2016. GuessWhat?! Visual object discovery through multi-modal dialogue. *CoRR* abs/1611.08481 (2016). [arXiv:1611.08481](http://arxiv.org/abs/1611.08481) <http://arxiv.org/abs/1611.08481>
- [5] Mihail Eric and Christopher D. Manning. 2017. A Copy-Augmented Sequence-to-Sequence Architecture Gives Good Performance on Task-Oriented Dialogue. *CoRR* abs/1701.04024 (2017). <http://arxiv.org/abs/1701.04024>
- [6] Ian Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, and Yoshua Bengio. 2014. Generative Adversarial Nets. In *Advances in Neural Information Processing Systems 27*, Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger (Eds.). Curran Associates, Inc., 2672–2680. <http://papers.nips.cc/paper/5423-generative-adversarial-nets.pdf>
- [7] Yash Goyal, Tejas Khot, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. 2017. Making the V in VQA Matter: Elevating the Role of Image Understanding in Visual Question Answering. In *Conference on Computer Vision and Pattern Recognition (CVPR)*.
- [8] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2015. Deep Residual Learning for Image Recognition. *CoRR* abs/1512.03385 (2015). <http://arxiv.org/abs/1512.03385>
- [9] Wei-Chih Hung, Yi-Hsuan Tsai, Yan-Ting Liou, Yen-Yu Lin, and Ming-Hsuan Yang. 2018. Adversarial Learning for Semi-Supervised Semantic Segmentation. *CoRR* abs/1802.07934 (2018).
- [10] Justin Johnson, Andrej Karpathy, and Fei-Fei Li. 2015. DenseCap: Fully Convolutional Localization Networks for Dense Captioning. *CoRR* abs/1511.07571 (2015). [arXiv:1511.07571](http://arxiv.org/abs/1511.07571) <http://arxiv.org/abs/1511.07571>
- [11] Jin-Hwa Kim, Devi Parikh, Dhruv Batra, Byoung-Tak Zhang, and Yuandong Tian. 2017. CoDraw: Visual Dialog for Collaborative Drawing. *arXiv preprint arXiv:1712.05558* (2017).
- [12] Ryan Kiros, Ruslan Salakhutdinov, and Richard S. Zemel. 2014. Unifying Visual-Semantic Embeddings with Multimodal Neural Language Models. *CoRR* abs/1411.2539 (2014). [arXiv:1411.2539](http://arxiv.org/abs/1411.2539) <http://arxiv.org/abs/1411.2539>
- [13] Satwik Kottur, José M. F. Moura, Stefan Lee, and Dhruv Batra. 2017. Natural Language Does Not Emerge ‘Naturally’ in Multi-Agent Dialog. *CoRR* abs/1706.08502 (2017). [arXiv:1706.08502](http://arxiv.org/abs/1706.08502) <http://arxiv.org/abs/1706.08502>
- [14] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. 2012. ImageNet Classification with Deep Convolutional Neural Networks. In *Proceedings of the 25th International Conference on Neural Information Processing Systems - Volume 1 (NIPS’12)*, Curran Associates Inc., USA, 1097–1105. <http://dl.acm.org/citation.cfm?id=2999134.2999257>
- [15] Yaniv Leviathan. 2018. Google Duplex: An AI System for Accomplishing Real-World Tasks Over the Phone. (May 2018). <https://ai.googleblog.com/2018/05/duplex-ai-system-for-natural-conversation.html>
- [16] Mike Lewis, Denis Yarats, Yann N Dauphin, Devi Parikh, and Dhruv Batra. 2017. Deal or no deal? end-to-end learning for negotiation dialogues. *arXiv preprint arXiv:1706.05125* (2017).
- [17] Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian Szegedy, Scott E. Reed, Cheng-Yang Fu, and Alexander C. Berg. 2015. SSD: Single Shot MultiBox Detector. *CoRR* abs/1512.02325 (2015). <http://arxiv.org/abs/1512.02325>
- [18] Jonathan Long, Evan Shelhamer, and Trevor Darrell. 2015. Fully Convolutional Networks for Semantic Segmentation. *CVPR* (2015).
- [19] Jiasen Lu, Anitha Kannan, Jianwei Yang, Devi Parikh, and Dhruv Batra. 2017. Best of Both Worlds: Transferring Knowledge from Discriminative Learning to a Generative Visual Dialog Model. *CoRR* abs/1706.01554 (2017). [arXiv:1706.01554](http://arxiv.org/abs/1706.01554) <http://arxiv.org/abs/1706.01554>
- [20] Jiasen Lu, Caiming Xiong, Devi Parikh, and Richard Socher. 2016. Knowing When to Look: Adaptive Attention via A Visual Sentinel for Image Captioning. *CoRR* abs/1612.01887 (2016). [arXiv:1612.01887](http://arxiv.org/abs/1612.01887) <http://arxiv.org/abs/1612.01887>
- [21] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex Graves, Ioannis Antonoglou, Daan Wierstra, and Martin A. Riedmiller. 2013. Playing Atari with Deep Reinforcement Learning. *CoRR* abs/1312.5602 (2013). <http://arxiv.org/abs/1312.5602>
- [22] Geraldin Nanfack, Azeddine Elhassouny, and Rachid Oulad Haj Thami. [n. d.]. Squeeze-SegNet: A new fast Deep Convolutional Neural Network for Semantic Segmentation. *CoRR* abs/1711.05491 ([n. d.]).
- [23] Joseph Redmon, Santosh Kumar Divvala, Ross B. Girshick, and Ali Farhadi. 2015. You Only Look Once: Unified, Real-Time Object Detection. *CoRR* abs/1506.02640 (2015). <http://arxiv.org/abs/1506.02640>
- [24] Shaoqing Ren, Kaiming He, Ross Girshick, and Jian Sun. 2015. Faster R-CNN: Towards Real-time Object Detection with Region Proposal Networks. In *Proceedings of the 28th International Conference on Neural Information Processing Systems - Volume 1 (NIPS’15)*. MIT Press, Cambridge, MA, USA, 91–99. <http://dl.acm.org/citation.cfm?id=2969239.2969250>
- [25] Michael Rutter, Karen Thorpe, Rosemary Greenwood, Kate Northstone, and Jean Golding. 2003. Twins as a natural experiment to study the causes of mild language delay: I: Design; twin-singleton differences in language, and obstetric risks. *Journal of Child Psychology and Psychiatry* 44, 3 (2003), 326–341.
- [26] Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. 2017. Mastering the game of Go without human knowledge. *Nature* 550 (2017), 354–359. <https://doi.org/10.1038/nature24270>
- [27] Min Joon Seo, Aniruddha Kembhavi, Ali Farhadi, and Hannaneh Hajishirzi. 2016. Bidirectional Attention Flow for Machine Comprehension. *CoRR* abs/1611.01603 (2016). <http://arxiv.org/abs/1611.01603>
- [28] Iulian Vlad Serban, Chinnadhurai Sankar, Mathieu Germain, Saizheng Zhang, Zhouhan Lin, Sandeep Subramanian, Taesup Kim, Michael Pieper, Sarath Chandar, Nan Rosemary Ke, Sai Mudumba, Alexandre de Brébisson, Jose Sotelo, Dendi Suhubdy, Vincent Michalski, Alexandre Nguyen, Joelle Pineau, and Yoshua Bengio. 2017. A Deep Reinforcement Learning Chatbot. *CoRR* abs/1709.02349 (2017). <http://arxiv.org/abs/1709.02349>
- [29] Karen Simonyan and Andrew Zisserman. 2014. Very deep convolutional networks for large-scale image recognition. *arXiv preprint arXiv:1409.1556* (2014).
- [30] Oriol Vinyals, Alexander Toshev, Samy Bengio, and Dumitru Erhan. 2014. Show and Tell: A Neural Image Caption Generator. *CoRR* abs/1411.4555 (2014). [arXiv:1411.4555](http://arxiv.org/abs/1411.4555) <http://arxiv.org/abs/1411.4555>
- [31] Gellért Weisz, Paweł Budzianowski, Pei-Hao Su, and Milica Gasic. 2018. Sample Efficient Deep Reinforcement Learning for Dialogue Systems with Large Action Spaces. *CoRR* abs/1802.03753 (2018). <http://arxiv.org/abs/1802.03753>
- [32] Ronald J Williams. 1992. Simple statistical gradient-following algorithms for connectionist reinforcement learning. In *Reinforcement Learning*. Springer, 5–32.
- [33] Qi Wu, Peng Wang, Chunhua Shen, Ian D. Reid, and Anton van den Hengel. 2017. Are You Talking to Me? Reasoned Visual Dialog Generation through Adversarial Learning. *CoRR* abs/1711.07613 (2017). [arXiv:1711.07613](http://arxiv.org/abs/1711.07613) <http://arxiv.org/abs/1711.07613>
- [34] Saining Xie, Ross B. Girshick, Piotr Dollár, Zhuowen Tu, and Kaiming He. 2016. Aggregated Residual Transformations for Deep Neural Networks. *CoRR* abs/1611.05431 (2016). <http://arxiv.org/abs/1611.05431>
- [35] Caiming Xiong, Victor Zhong, and Richard Socher. 2016. Dynamic Coattention Networks For Question Answering. *CoRR* abs/1611.01604 (2016). <http://arxiv.org/abs/1611.01604>
- [36] Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaron C. Courville, Ruslan Salakhutdinov, Richard S. Zemel, and Yoshua Bengio. 2015. Show, Attend and Tell: Neural Image Caption Generation with Visual Attention. *CoRR* abs/1502.03044 (2015). [arXiv:1502.03044](http://arxiv.org/abs/1502.03044) <http://arxiv.org/abs/1502.03044>
- [37] Ting Yao, Yingwei Pan, Yehao Li, Zhaofan Qiu, and Tao Mei. 2016. Boosting Image Captioning with Attributes. *CoRR* abs/1611.01646 (2016). [arXiv:1611.01646](http://arxiv.org/abs/1611.01646) <http://arxiv.org/abs/1611.01646>
- [38] Adams Wei Yu, David Dohan, Minh-Thang Luong, Rui Zhao, Kai Chen, Moham-mad Norouzi, and Quoc V. Le. 2018. QANet: Combining Local Convolution with Global Self-Attention for Reading Comprehension. *International Conference on Learning Representations, ICLR-2018* (2018).
- [39] Peng Zhang, Yash Goyal, Douglas Summers-Stay, Dhruv Batra, and Devi Parikh. 2016. Yin and Yang: Balancing and Answering Binary Visual Questions. In *Conference on Computer Vision and Pattern Recognition (CVPR)*.