

AN INTRODUCTION TO COLLECTIVE INTELLIGENCE

David H. Wolpert
 NASA Ames Research Center
 Moffett Field, CA 94035
 dhw@ptolemy.arc.nasa.gov

Kagan Tumer
 NASA Ames Research Center
 Caelum Research
 Moffett Field, CA 94035
 kagan@ptolemy.arc.nasa.gov

http://ic.arc.nasa.gov/ic/people/kagan/coin_pubs.html

Tech Report: NASA-ARC-IC-99-63

February 1, 2008

Abstract

This paper surveys the emerging science of how to design a “Collective Intelligence” (COIN). A COIN is a large multi-agent system where:

- i) There is little to no centralized communication or control.
- ii) There is a provided world utility function that rates the possible histories of the full system.

In particular, we are interested in COINs in which each agent runs a reinforcement learning (RL) algorithm. The conventional approach to designing large distributed systems to optimize a world utility does not use agents running RL algorithms. Rather, that approach begins with explicit modeling of the dynamics of the overall system, followed by detailed hand-tuning of the interactions between the components to ensure that they “cooperate” as far as the world utility is concerned. This approach is labor-intensive, often results in highly nonrobust systems, and usually results in design techniques that have limited applicability.

In contrast, we wish to solve the COIN design problem implicitly, via the “adaptive” character of the RL algorithms of each of the agents. This approach introduces an entirely new, profound design problem: Assuming the RL algorithms are able to achieve high rewards, what reward functions for the individual agents will, when pursued by those agents, result in high world utility? In other words, what reward functions will best ensure that we do not have phenomena like the tragedy of the commons, Braess’s paradox, or the liquidity trap?

Although still very young, research specifically concentrating on the COIN design problem has already resulted in successes in artificial domains, in particular in packet-routing, the leader-follower problem, and in variants of Arthur’s El Farol bar problem. It is expected that as it matures and draws upon other disciplines related to COINs, this research will greatly expand the range of tasks addressable by human engineers. Moreover, in addition to drawing on them, such a fully developed science of COIN design may provide much insight into other already established scientific fields, such as economics, game theory, and population biology.

1 INTRODUCTION

Over the past decade or so two separate developments have occurred in computer science whose intersection promises to open a vast new area of research, an area extending far beyond the current boundaries of computer science. The first of these developments is the growing realization of how useful it would be to be able to control distributed systems that have little (if any) centralized communication, and to do so “adaptively”, with minimal reliance on detailed knowledge of the system’s small-scale dynamical behavior. The second development is the maturing of the discipline of reinforcement learning (RL). This is the branch of machine learning that is concerned with an agent who periodically receives “reward” signals from the environment that partially reflect the value of that agent’s private utility function. The goal of an RL algorithm is to determine how, using those reward signals, the agent should update its action policy to maximize its utility [142, 258, 272]. (Until our detailed discussions below, we will use the term “reinforcement learning” broadly, to include any algorithm of this sort, including ones that rely on detailed Bayesian modeling of underlying Markov processes [225, 48, 98].

Intuitively, one might hope that RL would help us solve the distributed control problem, since RL is adaptive, and, in particular, since it is not restricted to domains having sufficient breadths of communication. However, by itself, conventional single-agent RL does not provide a means for controlling large, distributed systems. This is true even if the system *does* have centralized communication. The problem is that the space of possible action policies for such systems is too big to be searched. We might imagine as a variant using a large set of agents, each controlling only part of the system. Since the individual action spaces of such agents would be relatively small, we could realistically deploy conventional RL on each one. However, now we face the central question of how to map the world utility function concerning the overall system into private utility functions for each of the agents. In particular, how should we design those private utility functions so that each agent can realistically hope to optimize its function, and at the same time the collective behavior of the agents will optimize the world utility?

We use the term “COllective INtelligence” (COIN) to refer to any pair of a large, distributed collection of interacting computational processes among which there is little to no centralized communication or control, together with a ‘world utility’ function that rates the possible dynamic histories of the collection. The central COIN design problem we consider arises when the computational processes run RL algorithms: How, without any detailed modeling of the overall system, can one set the utility functions for the RL algorithms in a COIN to have the overall dynamics reliably and robustly achieve large values of the provided world utility? The benefits of an answer to this question would extend beyond the many branches of computer science, having major ramifications for many other sciences as well. Section 2 discusses some of those benefits. Section 3 reviews previous work that has bearing on the COIN design problem. Section 4 section constitutes the core of this chapter. It presents a quick outline of a promising

mathematical framework for addressing this problem in its most general form, and then experimental illustrations of the prescriptions of that framework. Throughout, we will use italics for emphasis, single quotes for informally defined terms, and double quotes to delineate colloquial terminology.

2 Background

There are many design problems that involve distributed computational systems where there are strong restrictions on centralized communication (“we can’t all talk”); or there is communication with a central processor, but that processor is not sufficiently powerful to determine how to control the entire system (“we aren’t smart enough”); or the processor is powerful enough in principle, but it is not clear what algorithm it could run by itself that would effectively control the entire system (“we don’t know what to think”). Just a few of the potential examples include:

- i) Designing a control system for constellations of communication satellites or for constellations of planetary exploration vehicles (world utility in the latter case being some measure of quality of scientific data collected);
- ii) Designing a control system for routing over a communication network (world utility being some aggregate quality of service measure)
- iii) Construction of parallel algorithms for solving numerical optimization problems (the optimization problem itself constituting the world utility);
- iv) Vehicular traffic control, *e.g.*, air traffic control, or high-occupancy toll-lanes for automobiles. (In these problems the individual agents are humans and the associated utility functions must be of a constrained form, reflecting the relatively inflexible kinds of preferences humans possess.);
- v) Routing over a power grid;
- vi) Control of a large, distributed chemical plant;
- vii) Control of the elements of an amorphous computer;
- viii) Control of the elements of a ‘noisy’ phased array radar;
- ix) Compute-serving over an information grid.

Such systems may be best controlled with an artificial COIN. However, the potential usefulness of deeper understanding of how to tackle the COIN design problem extends far beyond such engineering concerns. That’s because the COIN design problem is an inverse problem, whereas essentially all of the scientific fields that are concerned with naturally-occurring distributed systems analyze them purely as a “forward problem.” That is, those fields analyze what global behavior would arise from provided local dynamical laws, rather than grapple with the inverse problem of how to configure those laws to induce desired global behavior. (Indeed, the COIN design problem could almost be defined as decentralized adaptive control theory for massively distributed stochastic

environments.) It seems plausible that the insights garnered from understanding the inverse problem would provide a trenchant novel perspective on those fields. Just as tackling the inverse problem in the design of steam engines led to the first true understanding of the macroscopic properties of physical bodies (aka thermodynamics), so may the cracking of the COIN design problem may improve our understanding of many naturally-occurring COINs. In addition, although the focuses of those other fields are not on the COIN design problem, in that they are related to the COIN design problem, that problem may be able to serve as a “touchstone” for all those fields. This may then reveal novel connections between the fields.

As an example of how understanding the COIN design problem may provide a novel perspective on other fields, consider countries with capitalist human economies. Although there is no intrinsic world utility in such systems, they can still be viewed from the perspective of COINs, as naturally occurring COINs. For example, one can declare world utility to be a time average of the Gross Domestic Product (GDP) of the country in question. (World utility per se is not a construction internal to a human economy, but rather something defined from the outside.) The reward functions for the human agents in this example could then be the achievements of their personal goals (usually involving personal wealth to some degree).

Now in general, to achieve high world utility in a COIN it is necessary to avoid having the agents work at cross-purposes. Otherwise the system is vulnerable to economic phenomena like the Tragedy of the Commons (TOC), in which individual avarice works to lower world utility [115], or the liquidity trap, where behavior that helps the entire system when employed by some agents results in poor global behavior when employed by all agents [158]. One way to avoid such phenomena is by modifying the agents’ utility functions. In the context of capitalist economies, this kind of effect can be achieved via punitive legislation that modifies the rewards the agents receive for engaging in certain kinds of activity. A real world example of an attempt to make just such a modification was the creation of anti-trust regulations designed to prevent monopolistic practices.

In designing a COIN we usually have more freedom than anti-trust regulators though, in that there is no base-line “organic” private utility function over which we must superimpose legislation-like incentives. Rather, the entire “psychology” of the individual agents is at our disposal when designing a COIN. This obviates the need for honesty-elicitation (‘incentive compatible’) mechanisms, like auctions, which form a central component of conventional economics. Accordingly, COINs can differ in certain crucial respects from human economies. The precise differences — the subject of current research — seem likely to present many insights into the functioning of economic structures like anti-trust regulators.

To continue with this example, consider the usefulness, as far as the world utility is concerned, of having (commodity, or especially fiat) money in the COIN. Formally, from a COIN perspective, the use of ‘money’ for trading between agents constitutes

a particular class of couplings between the states and utility functions of the various agents. For example, if one agent’s ‘bank account’ variable goes up in a ‘trade’ with another agent, then a corresponding ‘bank account’ variable in that other agent must decrease to compensate. In addition to this coupling between the agents’ states, there is also a coupling between their utilities, if one assume that both agents will prefer to have more money rather than less, everything else being equal. However one might formally define such a ‘money’ structure, we can consider what happens if it does (or does not) obtain for an arbitrary dynamical system, in the context of an arbitrary world utility. For some such dynamical systems and world utilities, a money structure will improve the value of that world utility. But for the same dynamics, the use of a money structure will simultaneously induce *low levels* of other world utilities (a trivial example being a world utility that equals the negative of the first one). This raises a host of questions, like how to formally specify the most general set of world utilities that benefits significantly from using money-based private utility functions. If one is provided a world utility that is not a member of that set, then an “economics-like” configuration of the system is likely to result in poor performance. Such a characterization of how and when money helps improve world utilities of various sorts might have important implications for conventional human economics, especially when one chooses world utility to be one of the more popular choices for social welfare function. (See [251, 74] and references therein for some of the standard economics work that is most relevant to this issue.)

There are many other scientific fields that are currently under investigation from a COIN-design perspective. Some of them are, like economics, part of (or at least closely related to) the social sciences. These fields typically involve RL algorithms under the guise of human agents. An example of such a field is game theory, especially game theory of bounded rational players. As illustrated in our money example, viewing such systems from the perspective of a non-endogenous world utility, *i.e.*, from a COIN-design perspective, holds the potential for providing novel insight into them. (In the case of game theory, it holds the potential for leading to deeper understanding of many-player inverse stochastic game theory.)

However there are other scientific fields that might benefit from a COIN-design perspective even though they study systems that don’t even involve RL algorithms. The idea here is that if we viewed such systems from an “artificial” teleological perspective, both in concentrating on a non-endogenous world utility and in casting the nodal elements of the system as RL algorithms, we could learn a lot about the form of the ‘design space’ in which such systems live. (Just as in economics, where the individual nodal elements *are* RL algorithms, investigating the system using an externally imposed world utility might lead to insight.) Examples here are ecosystems (individual genes, individuals, or species being the nodal elements) and cells (individual organelles in Eukaryotes being the nodal elements). In both cases, the world utility could involve robustness of the desired equilibrium against external perturbation, efficient exploitation of free energy in the environment, etc.

3 Review of Literature Related to COINs

The following list elaborates what we mean by a COIN:

1) There are many processors running concurrently, performing actions that affect one another’s behavior.

2) There is little to no centralized personalized communication, *i.e.*, little to no behavior in which a small subset of the processors not only communicates with all the other processors, but communicates differently with each one of those other processors. Any single processor’s “broadcasting” the same information to all other processors is not precluded.

3) There is little to no centralized personalized control, *i.e.*, little to no behavior in which a small subset of the processors not only controls all the other processors, but controls each one of those other processors differently. “Broadcasting” the same control signal to all other processors is not precluded.

4) There is a well-specified task, typically in the form of extremizing a utility function, that concerns the behavior of the entire distributed system. So we are confronted with the inverse problem of how to configure the system to achieve the task.

The following elements characterize the sorts of approaches to COIN design we are concerned with here:

5) The approach for tackling (4) is scalable to very large numbers of processors.

6) The approach for tackling (4) is very broadly applicable. In particular, it can work when little (if any) “broadcasting” as in (2) and (3) is possible.

7) The approach for tackling (4) involves little to no hand-tailoring.

8) The approach for tackling (4) is robust and adaptive, with minimal need to “get the details exactly right or else,” as far as the stochastic dynamics of the system is concerned.

9) The individual processors are running RL algorithms. Unlike the other elements of this list, this one is not an *a priori* engineering necessity. Rather, it is a reflection of the fact that RL algorithms are currently the best-understood and most mature technology for addressing the points (8) and (9).

There are many approaches to COIN design that do not have every one of those features. These approaches constitute part of the overall field of COIN design. As discussed below though, not having every feature in our list, no single one of those approaches can be extended to cover the entire breadth of the field of COIN design. (This is not too surprising, since those approaches are parts of fields whose focus is not the COIN design problem per se.)

The rest of this section consists of brief presentations of some of these approaches, and in particular characterizes them in terms of our list of nine characteristics of COINs and of our desired data for their design. Of the approaches we discuss, at present it is probably

the ones in Artificial Intelligence and Machine Learning that are most directly applicable to COIN design. However it is fairly clear how to exploit those approaches for COIN design, and in that sense relatively little needs to be said about them. In contrast, as currently employed, the toolsets in the social sciences are not as immediately applicable to COIN design. However, it seems likely that there is more yet to be discovered about how to exploit them for COIN design. Accordingly, we devote more space to those social science-based approaches here.

We present an approach that holds promise for covering all nine of our desired features in Section 4.

3.1 AI and Machine Learning

There is an extensive body of work in AI and machine learning that is related to COIN design. Indeed, one of the most famous speculative works in the field can be viewed as an argument that AI should be approached as a COIN design problem [188]. Much work of a more concrete nature is also closely related to the problem of COIN design.

3.1.1 Reinforcement Learning

As discussed in the introduction, the maturing field of reinforcement learning provides a much needed tool for the types of problems addressed by COINs. Because RL generally provides model-free¹ and “online” learning features, it is ideally suited for the distributed environment where a “teacher” is not available and the agents need to learn successful strategies based on “rewards” and “penalties” they receive from the overall system at various intervals. It is even possible for the learners to use those rewards to modify *how* they learn [234, 235].

Although work on RL dates back to Samuel’s checker player [226], relatively recent theoretical [272] and empirical results [60, 263] have made RL one of the most active areas in machine learning. Many problems ranging from controlling a robot’s gait to controlling a chemical plant to allocating constrained resource have been addressed with considerable success using RL [113, 133, 196, 215, 294]. In particular, the RL algorithms $TD(\lambda)$ (which rates potential states based on a *value function*) [258] and Q -learning (which rates action-state pairs) [272] have been investigated extensively. A detailed investigation of RL is available in [142, 259, 272].

Although powerful and widely applicable, solitary RL algorithms will not perform well on large distributed heterogeneous problems in general. This is due to the very big size of the action-policy space for such problems. In addition, without centralized communication and control, how a solitary RL algorithm could run the full system at all, poorly or well, becomes a major concern.² For these reasons, it is natural to con-

¹There exist some model-based variants of traditional RL. See for example [7].

²One possible solution would be to run the RL off-line on a simulation of the full system and then

sider deploying many RL algorithms rather than a single one for these large distributed problems. We will discuss the coordination issues such an approach raises in conjunction with multi-agent systems in Section 3.1.3 and with learnability in COINs in Section 4.

3.1.2 Distributed Artificial Intelligence

The field of Distributed Artificial Intelligence (DAI) has arisen as more and more traditional Artificial Intelligence (AI) tasks have migrated toward parallel implementation. The most direct approach to such implementations is to directly parallelize AI production systems or the underlying programming languages [93, 221]. An alternative and more challenging approach is to use distributed computing, where not only are the individual reasoning, planning and scheduling AI tasks parallelized, but there are *different modules* with different such tasks, concurrently working toward a common goal [137, 138, 166].

In a DAI, one needs to ensure that the task has been modularized in a way that improves efficiency. Unfortunately, this usually requires a central controller whose purpose is to allocate tasks and process the associated results. Moreover, designing that controller in a traditional AI fashion often results in brittle solutions. Accordingly, recently there has been a move toward both more autonomous modules and fewer restrictions on the interactions among the modules [229].

Despite this evolution, DAI maintains the traditional AI concern with a pre-fixed set of *particular* aspects of intelligent behavior (*e.g.* reasoning, understanding, learning etc.) rather than on their *cumulative* character. As the idea that intelligence may have more to do with the interaction among components started to take shape [44, 45], focus shifted to concepts (*e.g.*, multi-agent systems) that better incorporated that idea [140].

3.1.3 Multi-Agent Systems

The field of Multi-Agent Systems (MAS) is concerned with the interactions among the members of such a set of agents [42, 108, 140, 239, 260], as well as the inner workings of each agent in such a set (*e.g.*, their learning algorithms) [38, 39, 40]. As in computational ecologies and computational markets (see below), a well-designed MAS is one that achieves a global task through the actions of its components. The associated design steps involve [140]:

1. Decomposing a global task into distributable subcomponents, yielding tractable tasks for each agent;
2. Establishing communication channels that provide sufficient information to each of the agents for it to achieve its task, but are not too unwieldy for the overall

convey the results to the components of the system at the price of a single centralized initialization (*e.g.*, [195]). In general though, this approach will suffer from being extremely dependent on “getting the details right” in the simulation.

system to sustain; and

3. Coordinating the agents in a way that ensures that they cooperate on the global task, or at the very least does not allow them to pursue conflicting strategies in trying to achieve their tasks.

Step (3) is rarely trivial; one of the main difficulties encountered in MAS design is that agents act selfishly and artificial cooperation structures have to be imposed on their behavior to enforce cooperation [12]. An active area of research, which holds promise for addressing parts the COIN design problem, is to determine how selfish agents’ “incentives” have to be engineered in order to avoid the tragedy of the commons (TOC) [244]. (This work draws on the economics literature, which we review separately below.) When simply providing the right incentives is not sufficient, one can resort to strategies that actively induce agents to cooperate rather than act selfishly. In such cases coordination [240], negotiations [156], coalition formation [228, 230, 296] or contracting [2] among agents may be needed to ensure that they do not work at cross purposes.

Unfortunately, all of these approaches share with DAI and its offshoots the problem of relying excessively on hand-tailoring, and therefore being difficult to scale and often nonrobust. In addition, except as noted in the next subsection, they involve no RL, and therefore the constituent computational elements are usually not as adaptive and robust as we would like.

3.1.4 Reinforcement Learning-Based Multi-Agent Systems

Because it neither requires explicit modeling of the environment nor having a “teacher” that provides the “correct” actions, the approach of having the individual agents in a MAS use RL is well-suited for MAS’s deployed in domains where one has little knowledge about the environment and/or other agents. There are two main approaches to designing such MAS’s:

- (i) One has ‘solipsistic agents’ that don’t know about each other and whose RL rewards are given by the performance of the entire system (so the joint actions of all other agents form an “inanimate background” contributing to the reward signal each agent receives);
- (ii) One has ‘social agents’ that explicitly model each other and take each others’ actions into account.

Both (i) and (ii) can be viewed as ways to (try to) coordinate the agents in a MAS in a robust fashion.

Solipsistic Agents: MAS’s with solipsistic agents have been successfully applied to a multitude of problems [60, 112, 122, 227, 233]. Generally, these schemes use RL algorithms similar to those discussed in Section 3.1.1. However much of this work lacks a well-defined global task or broad applicability (*e.g.*, [227]). More generally, none of the work with solipsistic agents scales well. (As illustrated in our experiments on the “bar problem”, recounted below.) The problem is that each agent must be able to discern

the effect of its actions on the overall performance of the system, since that performance constitutes its reward signal. As the number of agents increases though, the effects of any one agent’s actions (signal) will be swamped by the effects of other agents (noise), making the agent unable to learn well, if at all. (See the discussion below on learnability.) In addition, of course, solipsistic agents cannot be used in situations lacking centralized calculation and broadcast of the single global reward signal.

Social agents: MAS’s whose agents take the actions of other agents into account synthesize RL with game theoretic concepts (*e.g.*, Nash equilibrium). They do this to try to ensure that the overall system both moves toward achieving the overall global goal and avoids often deleterious oscillatory behavior [57, 99, 130, 131, 132]. To that end, the agents incorporate internal mechanisms that actively model the behavior of other agents. In Section 3.3.1, we discuss a situation where such modeling is necessarily self-defeating. More generally, this approach usually involves extensive hand-tailoring for the problem at hand.

3.2 Social Science–Inspired Systems

Some human economies provides examples of naturally occurring systems that can be viewed as a (more or less) well-performing COIN. The field of economics provides much more though. Both empirical economics (*e.g.*, economic history, experimental economics) and theoretical economics (*e.g.*, general equilibrium theory [3], theory of optimal taxation [189]) provide a rich literature on strategic situations where many parties interact. In fact, much of the entire field of economics can be viewed as concerning how to maximize certain constrained kinds of world utilities, when there are certain (very strong) restrictions on the individual agents and their interactions, and in particular when we have limited freedom in setting either the utility functions of those agents or modifying their RL algorithms in any other way.

In this section we summarize just two economic concepts, both of which are very closely related to COINs, in that they deal with how a large number of interacting agents can function in a stable and efficient manner: general equilibrium theory and mechanism design. We then discuss general attempts to apply those concepts to distributed computational problems. We follow this with a discussion of game theory, and then present a particular celebrated toy-world problem that involves many of these issues.

3.2.1 General Equilibrium Theory

Often the first version of “equilibrium” that one encounters in economics is that of supply and demand in single markets: the price of the market’s good is determined by where the supply and demand curves for that good intersect. In cases where there is interaction among multiple markets however, even when there is no production but only trading, one cannot simply determine the price of each market’s good individually, as

both the supply and demand for each good depends on the supply/demand of other goods. Considering the price fluctuations across markets leads to the concept of ‘general equilibrium’, where prices for each good are determined in such a way to ensure that all markets ‘clear’ [3, 250]. Intuitively, this means that prices are set so the total supply of each good is equal to the demand for that good.³ The existence of such an equilibrium, proven in [3], was first postulated by Leon Walras [271]. A mechanism that calculates the equilibrium (*i.e.*, ‘market-clearing’) prices now bears his name: the Walrasian auctioneer.

In general, for an arbitrary goal for the overall system, there is no reason to believe that having markets clear achieves that goal. In other words, there is no *a priori* reason why the general equilibrium point should maximize one’s provided world utility function. However, consider the case where one’s goal for the overall system is in fact that the markets clear. In such a context, examine the case where the interactions of real-world agents will induce the overall system to adopt the general equilibrium point, so long as certain broad conditions hold. Then if we can impose those conditions, we can cause the overall system to behave in the manner we wish. However general equilibrium theory is not sufficient to establish those “broad conditions”, since it says little about real-world agents. In particular, general equilibrium theory suffers from having no temporal aspect (*i.e.*, no dynamics) and from assuming that all the agents are perfectly rational.

Another shortcoming of general equilibrium theory as a model of real-world systems is that despite its concerning prices, it does not readily accommodate the full concept of money [96]. Of the three main roles money plays in an economy (medium of exchange in trades, store of value for future trades, and unit of account) none are essential in a general equilibrium setting. The unit of account aspect is not needed as the bookkeeping is performed by the Walrasian auctioneer. Since the supplies and demands are matched directly there is no need to facilitate trades, and thus no role for money as a medium of exchange. And finally, as the system reaches an equilibrium in one step, through the auctioneer, there is no need to store value for future trading rounds [209].

The reason that money is not needed can be traced to the fact that there is an “overseer” with global information who guides the system. If we remove the centralized communication and control exerted by this overseer, then (as in a real economy) agents

³More formally, each agent’s utility is a function of that agent’s allotment of all the possible goods. In addition, every good has a price. (Utility functions are independent of money.) Therefore, for any set of prices for the goods, every agent has a ‘budget’, given by their initial allotment of goods. We pool all the agents’ goods together. In the ‘tatonnement’ (single step) version of market clearing, we next allocate the goods back among the agents in such a way that each agent is given a total value of goods (as determined by the prices) equal to that agent’s budget (as determined by the prices and by that agent’s initial allotment). As a (formally identical) alternative, we can have a two-step process, in which first each agent is given funds equal to its budget, and then each agent decides how to use those funds to purchase goods from the central pool. In either case, the ‘market clearing’ prices are those prices for which exactly all of the goods in the pool are reallocated back among the agents (no “excess supply”), and for which each agent views its allocation of goods as optimizing its utility, subject to its budget and to those prices for the goods (no “excess demand”). Similar definitions hold for a ‘production’ rather than ‘endowment’ economy.

will no longer know the exact details of the overall economy. They will be forced to make guesses as in any learning system, and the differences in those guesses will lead to differences in their actions [160, 161].

Such a decentralized learning-based system more closely resembles a COIN than does a conventional general equilibrium system. In contrast to general equilibrium systems, the three main roles money plays in a human economy are crucial to the dynamics of such a decentralized system [13]. This comports with the important effects in COINs of having the agents' utility functions involve money (see Background section above).

3.2.2 Mechanism Design

Even if there exists centralized communication so that we aren't considering a full-blown COIN, if there is no centralized Walras-like control, it is usually highly non-trivial to induce the overall system to adopt the General Equilibrium point. One way to try to do so is via an auction. (This is the approach usually employed in computational markets — see below.) Along with optimal taxation and public good theory [157], the design of auctions is the subject of the field of mechanism design. More generally, mechanism design is concerned with the incentives that must be applied to any set of agents that interact and exchange goods [189, 267] in order to get those agents to exhibit desired behavior. Usually that desired behavior concerns pre-specified utility functions of some sort for each of the individual agents. In particular, mechanism design is usually concerned with incentive schemes which induce '(Pareto) efficient' (or 'Pareto optimal') allocations in which no agent can be made better off without hurting another agent [100, 101].

One particularly important type of such an incentive scheme is an auction. When many agents interact in a common environment often there needs to be a structure that supports the exchange of goods or information among those agents. Auctions provide one such (centralized) structure for managing exchanges of goods. For example, in the English auction all the agents come together and 'bid' for a good, and the price of the good is increased until only one bidder remains, who gets the good in exchange for the resource bid. As another example, in the Dutch auction the price of a good is decreased until one buyer is willing to pay the current price.

All auctions perform the same task: match supply and demand. As such, auctions are one of the ways in which price equilibration among a set of interacting agents (perhaps an equilibration approximating general equilibrium, perhaps not) can be achieved. However, an auction mechanism that induces Pareto efficiency does not necessarily maximize some other world utility. For example, in a transaction in an English auction both the seller and the buyer benefit. They may even have arrived at an allocation which is efficient. However, in that the winner may well have been willing to pay more for the good, such an outcome may confound the goal of the market designer, if that designer's goal is to maximize revenue. This point is returned to below, in the context of computational

economics.

3.2.3 Computational Economics

‘Computational economies’ are schemes inspired by economics, and more specifically by general equilibrium theory and mechanism design theory, for managing the components of a distributed computational system. They work by having a ‘computational market’, akin to an auction, guide the interactions among those components. Such a market is defined as any structure that allows the components of the system to exchange information on relative valuation of resources (as in an auction), establish equilibrium states (*e.g.*, determine market clearing prices) and exchange resources (*i.e.*, engage in trades).

Such computational economies can be used to investigate real economies and biological systems [32, 36, 37, 149]. They can also be used to design distributed computational systems. For example, such computational economies are well-suited to some distributed resource allocation problems, where each component of the system can either directly produce the “goods” it needs or acquire them through trades with other components. Computational markets often allow for far more heterogeneity in the components than do conventional resource allocation schemes. Furthermore, there is both theoretical and empirical evidence suggesting that such markets are often able to settle to equilibrium states. For example, auctions find prices that satisfy both the seller and the buyer which results in an increase in the utility of both (else one or the other would not have agreed to the sale). Assuming that all parties are free to pursue trading opportunities, such mechanisms move the system to a point where all possible bilateral trades that could improve the utility of both parties are exhausted.

Now restrict attention to the case, implicit in much of computational market work, with the following characteristics: First, world utility can be expressed as a monotonically increasing function F where each argument i of F can in turn be interpreted as the value of a pre-specified utility function f_i for agent i . Second, each of those f_i is a function of an i -indexed ‘goods vector’ x_i of the non-perishable goods “owned” by agent i . The components of that vector are $x_{i,j}$, and the overall system dynamics is restricted to conserve the vector $\sum_i x_{i,j}$. (There are also some other, more technical conditions.) As an example, the resource allocation problem can be viewed as concerning such vectors of “owned” goods.

Due to the second of our two conditions, one can integrate a market-clearing mechanism into any system of this sort. Due to the first condition, since in a market equilibrium with non-perishable goods no (rational) agent ends up with a value of its utility function lower than the one it started with, the value of the world utility function must be higher at equilibrium than it was initially. In fact, so long as the individual agents are smart enough to avoid all trades in which they do not benefit, any computational market can only improve this kind of world utility, even if it does not achieve the market equilibrium. (See the discussion of “weak triviality” below.)

This line of reasoning provides one of the main reasons to use computational markets when they can be applied. Conversely, it underscores one of the major limitations of such markets: Starting with an arbitrary world utility function with arbitrary dynamical restrictions, it may be quite difficult to cast that function as a monotonically increasing F taking as arguments a set of agents' goods-vector-based utilities f_i , if we require that those f_i be well-enough behaved that we can reasonably expect the agents to optimize them in a market setting.

One example of a computational economy being used for resource allocation is Huberman and Clearwater's use of a double-blind auction to solve the complex task of controlling the temperature of a building. In this case, each agent (individual temperature controller) bids to buy or sell cool or warm air. This market mechanism leads to an equitable temperature distribution in the system [135]. Other domains where market mechanisms were successfully applied include purchasing memory in an operating systems [53], allocating virtual circuits [89], "stealing" unused CPU cycles in a network of computers [79, 269], predicting option futures in financial markets [214], and numerous scheduling and distributed resource allocation problems [159, 165, 245, 255, 275, 276].

Computational economics can also be used for tasks not tightly coupled to resource allocation. For example, following the work of Maes [174] and Ferber [88], Baum shows how by using computational markets a large number of agents can interact and cooperate to solve a variant of the blocks world problem [23, 24].

Viewed as candidate COINs, all market-based computational economics fall short in relying on both centralized communication and centralized control to some degree. Often that reliance is extreme. For example, the systems investigated by Baum not only have the centralized control of a market, but in addition have centralized control of all other non-market aspects of the system. (Indeed, the market is secondary, in that it is only used to decide which single expert among a set of candidate experts gets to exert that centralized control at any given moment). There has also been doubt cast on how well computational economies perform in practice [264], and they also often require extensive hand-tailoring in practice.

Finally, return to consideration of a world utility function that is a monotonically increasing function f whose arguments are the utilities of the agents. In general, the maximum of such a world utility function will be a Pareto optimal point. So given the utility functions of the agents, by considering all such f we map out an infinite set S of Pareto optimal points that maximize *some* such world utility function. (S is usually infinite even if we only consider maximizing those world utilities subject to an overall conservation of goods constraint.) Now the market equilibrium is a Pareto optimal point, and therefore lies in S . But it is only one element of S . Moreover, it is usually set in full by the utilities of the agents, in concert with the agents' initial endowments. In particular, it is independent of the world utility. In general then, given the utilities of the agents and a world utility f , there is no *a priori* reason to believe that the particular

element in S picked out by the auction is the point that maximizes that particular world utility. This subtlety is rarely addressed in the work on using computational markets to achieve a global goal. It need not be uncircumventable however. For example, one obvious idea would be to try to distort the agents' *perceptions* of their utility functions and/or initial endowments so that the resultant market equilibrium has a higher value of the world utility at hand.⁴

3.2.4 Perfect Rationality Noncooperative Game Theory

Game theory is the branch of mathematics concerned with formalized versions of “games”, in the sense of chess, poker, nuclear arms races, and the like [31, 101, 87, 242, 75, 170, 20, 10]. It is perhaps easiest to describe it by loosely defining some of its terminology, which we do here and in the next subsection.

The simplest form of a game is that of ‘non-cooperative single-stage extensive-form’ game, which involves the following situation: There are two or more agents (called ‘players’ in the literature), each of which has a pre-specified set of possible actions that it can follow. (A ‘finite’ game has finite sets of possible actions for all the players.) In addition, each agent i has a utility function (also called a ‘payoff matrix’ for finite games). This maps any ‘profile’ of the action choices of all agents to an associated utility value for agent i . (In a ‘zero-sum’ game, for every profile, the sum of the payoffs to all the agents is zero.)

The agents choose their actions in a sequence, one after the other. The structure determining what each agent knows concerning the action choices of the preceding agents is known as the ‘information set’.⁵ Games in which each agent knows exactly what the preceding (‘leader’) agent did are known as ‘Stackelberg games’. (A variant of such a game is considered in our experiments below. See also [153].)

In a ‘multi-stage’ game, after all the agents choose their first action, each agent is provided some information concerning what the other agents did. The agent uses this

⁴In fact, the second theorem of welfare economics [250] states that given any world utility such that: i) its global maximum can be written as a Pareto optimal point for agents’ utilities all of whose level sets are convex; ii) no other maximum of that world utility is Pareto optimal for those agents’ utilities; then one can always set initial endowments of the agents so that that Pareto optimal point corresponds to the price-clearing point for those endowments. Note though that that setting of endowments requires a centralized process. Moreover, even if we are allowed such a process to set the endowments, we still may not be able to successfully exploit this theorem to arrive at the world utility maximum, if we use markets involving iterative trading with dynamic associated prices, like those in the real world. This is because an intermediate trade with “incorrect” prices may have resulted in some particular agent’s having its utility rise beyond the level it has at the point maximizing world utility, and since that agent will never afterward engage in a trade that diminishes its utility, the system will never arrive at the world utility maximum.

⁵While stochastic choices of actions is central to game theory, most of the work in the field assumes the information in information sets is in the form of definite facts, rather than a probability distribution. Accordingly, there has been relatively little work incorporating Shannon information theory into the analysis of information sets.

information to choose its next action. In the usual formulation, each agent gets its payoff at the end of all of the game’s stages.

An agent’s ‘strategy’ is the rule it elects to follow mapping the information it has at each stage of a game to its associated action. It is a ‘pure strategy’ if it is a deterministic rule. If instead the agent’s action is chosen by randomly sampling from a distribution, that distribution is known a ‘mixed strategy’. Note that an agent’s strategy concerns *all* possible sequences of provided information, even any that cannot arise due to the strategies of the other agents.

Any multi-stage extensive-form game can be converted into a ‘normal form’ game, which is a single-stage game in which each agent is ignorant of the actions of the other agents, so that all agents choose their actions “simultaneously”. This conversion is achieved by having the “actions” of each agent in the normal form game correspond to an entire strategy in the associated multi-stage extensive-form game. The payoffs to all the agents in the normal form game for a particular strategy profile is then given by the associated payoff matrices of the multi-stage extensive form-game.

A ‘solution’ to a game, or an ‘equilibrium’, is a profile in which every agent behaves “rationally”. This means that every agent’s choice of strategy optimizes its utility subject to a pre-specified set of conditions. In conventional game theory those conditions involve, at a minimum, perfect knowledge of the payoff matrices of all other players, and often also involve specification of what strategies the other agents adopted and the like. In particular, a ‘Nash equilibrium’ is a a profile where each agent has chosen the best strategy it can, *given the choices of the other agents*. A game may have no Nash equilibria, one equilibrium, or many equilibria in the space of pure strategies. A beautiful and seminal theorem due to Nash proves that every game has at least one Nash equilibrium in the space of mixed strategies [199].

There are several different reasons one might expect a game to result in a Nash equilibrium. One is that it is the point that perfectly rational Bayesian agents would adopt, assuming the probability distributions they used to calculate expected payoffs were consistent with one another [9, 143]. A related reason, arising even in a non-Bayesian setting, is that a Nash equilibrium provides “consistent” predictions, in that if all parties predict that the game will converge to a Nash equilibrium, no one will benefit by changing strategies. Having a consistent prediction does not ensure that all agents’ payoffs are maximized though. The study of small perturbations around Nash equilibria from a stochastic dynamics perspective is just one example of a ‘refinement’ of Nash equilibrium, that is a criterion for selecting a single equilibrium state when more than one is present [176].

In cooperative game theory the agents are able to enter binding contracts with one another, and thereby coordinate their strategies. This allows the agents to avoid being “stuck” in Nash equilibria that are Pareto inefficient, that is being stuck at equilibrium profiles in which all agents would benefit if only they could agree to all adopt different

strategies, with no possibility of betrayal. The *characteristic function* of a game involves subsets (‘coalitions’) of agents playing the game. For each such subset, it gives the sum of the payoffs of the agents in that subset that those agents can guarantee if they coordinate their strategies. An *imputation* is a division of such a guaranteed sum among the members of the coalition. It is often the case that for a subset of the agents in a coalition one imputation *dominates* another, meaning that under threat of leaving the coalition that subset of agents can demand the first imputation rather than the second. So the problem each agent i is confronted with in a cooperative game is which set of other agents to form a coalition with, given the characteristic function of the game and the associated imputations i can demand of its partners. There are several different kinds of solution for cooperative games that have received detailed study, varying in how the agents address this problem of who to form a coalition with. Some of the more popular are the ‘core’, the ‘Shapley value’, the ‘stable set solution’, and the ‘nucleolus’.

In the real world, the actual underlying game the agents are playing does not only involve the actions considered in cooperative game theory’s analysis of coalitions and imputations. The strategies of that underlying game also involve bargaining behavior, considerations of trying to cheat on a given contract, bluffing and threats, and the like. In many respects, by concentrating on solutions for coalition formation and their relation with the characteristic function, cooperative game theory abstracts away these details of the true underlying game. Conversely though, progress has recently been made in understanding how cooperative games can arise from non-cooperative games, as they must in the real world [10].

3.3 Evolution and Learning in Games

Not surprisingly, game theory has come to play a large role in the field of multi-agent systems. In addition, due to Darwinian natural selection, one might expect game theory to be quite important in population biology, in which the “utility functions” of the individual agents can be taken to be their reproductive fitness. As it turns out, there is an entire subfield of game theory concerned with this connection with population biology, called ‘evolutionary game theory’ [178, 181].

To introduce evolutionary game theory, consider a game in which all players share the same space of possible strategies, and there is an additional space of possible ‘attribute vectors’ that characterize an agent, along with a probability distribution g across that new space. (Examples of attributes in the physical world could be things like size, speed, etc.) We select a set of agents to play a game by randomly sampling g . Those agents’ attribute vectors jointly determine the payoff matrices of each of the individual agents. (Intuitively, what benefit accrues to an agent for taking a particular action depends on its attributes and those of the other agents.) However each agent i has limited information concerning both its attribute vector and that of the other players in the game, information encapsulated in an ‘information structure’. The information

structure specifies how much each agent knows concerning the game it is playing.

In this context, we enlarge the meaning of the term “strategy” to not just be a mapping from information sets and the like to actions, but from entire information structures to actions. In addition to the distribution g over attribute vectors, we also have a distribution over strategies, h . A strategy s is a ‘population strategy’ if h is a delta function about s . Intuitively, we have a population strategy when each animal in a population “follows the same behavioral rules”, rules that take as input what the animal is able to discern about its strengths and weakness relative to those other members of the population, and produce as output how the animal will act in the presence of such animals.

Given g , a population strategy centered about s , and its own attribute vector, any player i in the support of g has an expected payoff for any strategy it might adopt. When i ’s payoff could not improve if it were to adopt any strategy other than s , we say that s is ‘evolutionary stable’. Intuitively, an evolutionary stable strategy is one that is stable with respect to the introduction of mutants into the population.

Now consider a sequence of such evolutionary games. Interpret the payoff that any agent receives after being involved in such a game as the ‘reproductive fitness’ of that agent, in the biological sense. So the higher the payoff the agent receives, in comparison to the fitnesses of the other agents, the more “offspring” it has that get propagated to the next game. In the continuum-time limit, where games are indexed by the real number t , this can be formalized by a differential equation. This equation specifies the derivative of g_t evaluated for each agent i ’s attribute vector, as a monotonically increasing function of the relative difference between the payoff of i and the average payoff of all the agents. (We also have such an equation for h .) The resulting dynamics is known as ‘replicator dynamics’, with an evolutionary stable population strategy, if it exists, being one particular fixed point of the dynamics.

Now consider removing the reproductive aspect of evolutionary game theory, and instead have each agent propagate to the next game, with “memory” of the events of the preceding game. Furthermore, allow each agent to modify its strategy from one game to the next by “learning” from its memory of past games, in a bounded rational manner. The field of learning in games is concerned with exactly such situations [100, 11, 17, 27, 80, 146, 205, 202]. Most of the formal work in this field involves simple models for the learning process of the agents. For example, in ‘fictitious play’ [100], in each successive game, each agent i adopts what would be its best strategy if its opponents chose their strategies according to the empirical frequency distribution of such strategies that i has encountered in the past. More sophisticated versions of this work employ simple Bayesian learning algorithms, or re-inventions of some of the techniques of the RL community [222]. Typically in learning in games one defines a payoff to the agent for a sequence of games, for example as a discounted sum of the payoffs in each of the constituent games. Within this framework one can study the long term effects of

strategies such as cooperation and see if they arise naturally and if so, under what circumstances.

Many aspects of real world games that do not occur very naturally otherwise arise spontaneously in these kinds of games. For example, when the number of games to be played is not pre-fixed, it may behoove a particular agent i to treat its opponent better than it would otherwise, since i *may* have to rely on that other agent’s treating it well in the future, if they end up playing each other again. This framework also allows us to investigate the dependence of evolving strategies on the amount of information available to the agents [184]; the effect of communication on the evolution of cooperation [185, 187]; and the parallels between auctions and economic theory [123, 186].

In many respects, learning in games is even more relevant to the study of COINs than is traditional game theory. However it suffers from the same major shortcoming; it is almost exclusively focused on the forward problem rather than the inverse problem. In essence, COIN design is the problem of *inverse* game theory.

3.3.1 El Farol Bar Problem

The “El Farol” bar problem and its variants provide a clean and simple testbed for investigating certain kinds of interactions among agents [4, 50, 241]. In the original version of the problem, which arose in economics, at each time step (each “night”), each agent needs to decide whether to attend a particular bar. The goal of the agent in making this decision depends on the total attendance at the bar on that night. If the total attendance is below a preset capacity then the agent should have attended. Conversely, if the bar is overcrowded on the given night, then the agent should not attend. (Because of this structure, the bar problem with capacity set to 50% of the total number of agents is also known as the ‘minority game’; each agent selects one of two groups at each time step, and those that are in the minority have made the right choice). The agents make their choices by predicting ahead of time whether the attendance on the current night will exceed the capacity and then taking the appropriate course of action.

What makes this problem particularly interesting is that it is impossible for each agent to be perfectly “rational”, in the sense of correctly predicting the attendance on any given night. This is because if most agents predict that the attendance will be low (and therefore decide to attend), the attendance will actually high, while if they predict the attendance will be high (and therefore decide not to attend) the attendance will be low. (In the language of game theory, this essentially amounts to the property that there are no pure strategy Nash equilibria [52, 292].) Alternatively, viewing the overall system as a COIN, it has a Prisoner’s Dilemma-like nature, in that “rational” behavior by all the individual agents thwarts the global goal of maximizing total enjoyment (defined as the sum of all agents’ enjoyment and maximized when the bar is exactly at capacity).

This frustration effect is similar to what occurs in spin glasses in physics, and makes the bar problem closely related to the physics of emergent behavior in distributed sys-

tems [49, 50, 51, 295]. Researchers have also studied the dynamics of the bar problem to investigate economic properties like competition, cooperation and collective behavior and especially their relationship to market efficiency [64, 141, 232].

3.4 Biologically Inspired Systems

Properly speaking, biological systems do not involve utility functions and searches across them with RL algorithms. However it has long been appreciated that there are many ways in which viewing biological systems as involving searches over such functions can lead to deeper understanding of them [238, 290]. Conversely, some have argued that the mechanism underlying biological systems can be used to help design search algorithms [124].⁶

These kinds of reasoning which relate utility functions and biological systems have traditionally focussed on the case of a single biological system operating in some external environment. If we extend this kind of reasoning, to a set of biological systems that are co-evolving with one another, then we have essentially arrived at biologically-based COINs. This section discusses some of how previous work in the literature bears on this relationship between COINs and biology.

3.4.1 Population Biology and Ecological Modeling

The fields of population biology and ecological modeling are concerned with the large-scale “emergent” processes that govern the systems that consist of many (relatively) simple entities interacting with one another [25, 116]. As usually cast, the “simple entities” are members of one or more species, and the interactions are some mathematical abstraction of the process of natural selection as it occurs in biological systems (involving processes like genetic reproduction of various sorts, genotype-phenotype mappings, inter and intra-species competitions for resources, etc.). Population Biology and ecological modeling in this context addresses questions concerning the dynamics of the resultant ecosystem, and in particular how its long-term behavior depends on the details of the interactions between the constituent entities. Broadly construed, the paradigm of ecological modeling can even be broadened to study how natural selection and self-regulating feedback creates a stable planet-wide ecological environment—Gaia [167].

The underlying mathematical models of other fields can often be usefully modified to apply to the kinds of systems population biology is interested in [14]. (See also the discussion in the game theory subsection above.) Conversely, the underlying mathematical models of population biology and ecological modeling can be applied to other non-biological systems. In particular, those models shed light on social issues such as the emergence of language or culture, warfare, and economic competition [81, 82, 102].

⁶See [173, 285] though for some counter-arguments to the particular claims most commonly made in this regard.

They also can be used to investigate more abstract issues concerning the behavior of large complex systems with many interacting components [103, 114, 179, 203, 213].

Going a bit further afield, an approach that is related in spirit to ecological modeling is ‘computational ecologies’. These are large distributed systems where each component of the system’s acting (seemingly) independently results in complex global behavior. Those components are viewed as constituting an “ecology” in an abstract sense (although much of the mathematics is not derived from the traditional field of ecological modeling). In particular, one can investigate how the dynamics of the ecology is influenced by the information available to each component and how cooperation and communication among the components affects that dynamics [134, 136].

Although in some ways the most closely related to COINs of the current ecology-inspired research, the field of computational ecologies has some significant shortcomings if one tries to view it as a full science of COINs. In particular, it suffers from not being designed to solve the inverse problem of how to configure the system so as to arrive at a particular desired dynamics. This is a difficulty endemic to the general program of equating ecological modeling and population biology with the science of COINs. These fields are primarily concerned with the “forward problem” of determining the dynamics that arises from certain choices of the underlying system. Unless one’s desired dynamics is sufficiently close to some dynamics that was previously catalogued (during one’s investigation of the forward problem), one has very little information on how to set up the components and their interactions to achieve that desired dynamics. In addition, most of the work in these fields does not involve RL algorithms, and viewed as a context in which to design COINs suffers from a need for hand-tailoring, and potentially lack of robustness and scalability.

3.4.2 Swarm Intelligence

The field of ‘swarm intelligence’ is concerned with systems that are modeled after social insect colonies, so that the different components of the system are queen, worker, soldier, etc. It can be viewed as ecological modeling in which the individual entities have extremely limited computing capacity and/or action sets, and in which there are very few types of entities. The premise of the field is that the rich behavior of social insect colonies arises not from the sophistication of any individual entity in the colony, but from the interaction among those entities. The objective of current research is to uncover kinds of interactions among the entity types that lead to pre-specified behavior of some sort.

More speculatively, the study of social insect colonies may also provide insight into how to achieve learning in large distributed systems. This is because at the level of the individual insect in a colony, very little (or no) learning takes place. However across evolutionary time-scales the social insect species as a whole functions as if the various individual types in a colony had “learned” their specific functions. The “learning” is the

direct result of natural selection. (See the discussion on this topic in the subsection on ecological modeling.)

Swarm intelligences have been used to adaptively allocate tasks in a mail company [34], solve the traveling salesman problem [69, 70] and route data efficiently in dynamic networks [33, 236, 257] among others. Despite this, such intelligences do not really constitute a general approach to designing COINs. There is no general framework for adapting swarm intelligences to maximize particular world utility functions. Accordingly, such intelligences generally need to be hand-tailored for each application. And after such tailoring, it is often quite a stretch to view the system as “biological” in any sense, rather than just a simple and *a priori* reasonable modification of some previously deployed system.

3.4.3 Artificial Life

The two main objectives of Artificial Life, closely related to one another, are understanding the abstract functioning and especially the origin of terrestrial life, and creating organisms that can meaningfully be called “alive” [163].

The first objective involves formalizing and abstracting the mechanical processes underpinning terrestrial life. In particular, much of this work involves various degrees of abstraction of the process of self-replication [43, 247, 268]. Some of the more real-world-oriented work on this topic involves investigating how lipids assemble into more complex structures such as vesicles and membranes, which is one of the fundamental questions concerning the origin of life [66, 77, 208, 212, 201]. Many computer models have been proposed to simulate this process, though most suffer from overly simplifying the molecular morphology.

More generally, work concerned with the origin of life can constitute an investigation of the functional self-organization that gives rise to life [180]. In this regard, an important early work on functional self-organization is the *lambda calculus*, which provides an elegant framework (recursively defined functions, lack of distinction between object and function, lack of architectural restrictions) for studying computational systems [56]. This framework can be used to develop an artificial chemistry “function gas” that displays complex cooperative properties [92].

The second objective of the field of Artificial Life is less concerned with understanding the details of terrestrial life per se than of using terrestrial life as inspiration for how to design living systems. For example, motivated by the existence (and persistence) of computer viruses, several workers have tried to design an immune system for computers that will develop “antibodies” and handle viruses both more rapidly and more efficiently than other algorithms [94, 148, 248]. More generally, because we only have one sampling point (life on Earth), it is very difficult to precisely formulate the process by which life emerged. By creating an artificial world inside a computer however, it is possible to study far more general forms of life [217, 218, 219]. See also [286] where the

argument is presented that the richest way of approaching the issue of defining “life” is phenomenologically, in terms of self-*dissimilar* scaling properties of the system.

3.4.4 Training cellular automata with genetic algorithms

Cellular automata can be viewed as digital abstractions of physical gases [35, 85, 277, 278]. Formally, they are discrete-time recurrent neural nets where the neurons live on a grid, each neuron has a finite number of potential states, and inter-neuron connections are (usually) purely local. (See below for a discussion of recurrent neural nets.) So the state update rule of each neuron is fixed and local, the next state of a neuron being a function of the current states of it and of its neighboring elements.

The state update rule of (all the neurons making up) any particular cellular automaton specifies the mapping taking the initial configuration of the states of all of its neurons to the final, equilibrium (perhaps strange) attractor configuration of all those neurons. So consider the situation where we have a desired such mapping, and want to know an update rule that induces that mapping. This is a search problem, and can be viewed as similar to the inverse problem of how to design a COIN to achieve a pre-specified global goal, albeit a “COIN” whose nodal elements do not use RL algorithms.

Genetic algorithms are a special kind of search algorithm, based on analogy with the biological process of natural selection via recombination and mutation of a genome [190]. Although genetic algorithms (and ‘evolutionary computation’ in general) have been studied quite extensively, there is no formal theory justifying genetic algorithms as search algorithms [172, 285] and few empirical comparisons with other search techniques. One example of a well-studied application of genetic algorithms is to (try to) solve the inverse problem of finding update rules for a cellular automaton that induce a pre-specified mapping from its initial configuration to its attractor configuration. To date, they have used this way only for extremely simple configuration mappings, mappings which can be trivially learned by other kinds of systems. Despite the simplicity of these mappings, the use of genetic algorithms to try to train cellular automata to exhibit them has achieved little success [61, 62, 191, 192].

3.5 Physics-Based Systems

3.5.1 Statistical Physics

Equilibrium statistical physics is concerned with the stable state character of large numbers of very simple physical objects, interacting according to well-specified local deterministic laws, with probabilistic noise processes superimposed [5, 220]. Typically there is no sense in which such systems can be said to have centralized control, since all particles contribute comparably to the overall dynamics.

Aside from mesoscopic statistical physics, the numbers of particles considered are usu-

ally huge (*e.g.*, 10^{23}), and the particles themselves are extraordinarily simple, typically having only a few degrees of freedom. Moreover, the noise processes usually considered are highly restricted, being those that are formed by “baths”, of heat, particles, and the like. Similarly, almost all of the field restricts itself to deterministic laws that are readily encapsulated in Hamilton’s equations (Schrodinger’s equation and its field-theoretic variants for quantum statistical physics). In fact, much of equilibrium statistical physics isn’t even concerned with the dynamic laws by themselves (as for example is stochastic Markov processes). Rather it is concerned with invariants of those laws (*e.g.*, energy), invariants that relate the states of all of the particles. Trivially then, deterministic laws without such readily-discoverable invariants are outside of the purview of much of statistical physics.

One potential use of statistical physics for COINs involves taking the systems that statistical physics analyzes, especially those analyzed in its condensed matter variant (*e.g.*, spin glasses [252, 253]), as simplified models of a class of COINs. This approach is used in some of the analysis of the Bar problem (see above). It is used more overtly in (for example) the work of Galam [104], in which the equilibrium coalitions of a set of “countries” are modeled in terms of spin glasses. This approach cannot provide a general COIN framework though. In addition to the restrictions listed above on the kinds of systems it considers, this is due to its not providing a general solution to arbitrary COIN inversion problems, and to its not employing RL algorithms.⁷

Another contribution that statistical physics can make is with the mathematical techniques it has developed for its own purposes, like mean field theory, self-averaging approximations, phase transitions, Monte Carlo techniques, the replica trick, and tools to analyze the thermodynamic limit in which the number of particles goes to infinity. Although such techniques have not yet been applied to COINs, they have been successfully applied to related fields. This is exemplified by the use of the replica trick to analyze two-player zero-sum games with random payoff matrices in the thermodynamic limit of the number of strategies in [28]. Other examples are the numeric investigation of iterated prisoner’s dilemma played on a lattice [261], the analysis of stochastic games by expressing of deviation from rationality in the form of a “heat bath” [176], and the use of topological entropy to quantify the complexity of a voting system studied in [182].

Other quite recent work in the statistical physics literature is formally identical to that in other fields, but presents it from a novel perspective. A good example of this is [246], which is concerned with the problem of controlling a spatially extended system with a single controller, by using an algorithm that is identical to a simple-minded proportional

⁷In regard to the latter point however, it’s interesting to speculate about recasting statistical physics as a COIN, by viewing each of the particles in the physical system as running an “RL algorithm” that perfectly optimizes the “utility function” of its Lagrangian, given the “actions” of the other particles. In this perspective, many-particle physical systems are multi-stage games that are at Nash equilibrium in each stage. So for example, a frustrated spin glass is such a system at a Nash equilibrium that is not Pareto optimal.

RL algorithm (in essence, a rediscovery of RL).

3.5.2 Action Extremization

Much of the theory of physics can be cast as solving for the extremization of an actional, which is a functional of the worldline of an entire (potentially many-component) system across all time. The solution to that extremization problem constitutes the actual worldline followed by the system. In this way the calculus of variations can be used to solve for the worldline of a dynamic system. As an example, simple Newtonian dynamics can be cast as solving for the worldline of the system that extremizes a quantity called the ‘Lagrangian’, which is a function of that worldline and of certain parameters (*e.g.*, the ‘potential energy’) governing the system at hand. In this instance, the calculus of variations simply results in Newton’s laws.

If we take the dynamic system to be a COIN, we are assured that its worldline automatically optimizes a “global goal” consisting of the value of the associated actional. If we change physical aspects of the system that determine the functional form of the actional (*e.g.*, change the system’s potential energy function), then we change the global goal, and we are assured that our COIN optimizes that new global goal. Counter-intuitive physical systems, like those that exhibit Braess’ paradox [21], are simply systems for which the “world utility” implicit in our human intuition is extremized at a point different from the one that extremizes the system’s actional.

The challenge in exploiting this to solve the COIN design problem is in translating an arbitrary provided global goal for the COIN into a parameterized actional. Note that that actional must govern the dynamics of the physical COIN, and the parameters of the actional must be physical variables in the COIN, variables whose values we can modify.

3.5.3 Active Walker Models

The field of active walker models [22, 117, 118] is concerned with modeling “walkers” (be they human walkers or instead simple physical objects) crossing fields along trajectories, where those trajectories are a function of several factors, including in particular the trails already worn into the field. Often the kind of trajectories considered are those that can be cast as solutions to actional extremization problems so that the walkers can be explicitly viewed as agents optimizing a private utility.

One of the primary concerns with the field of active walker models is how the trails worn in the field change with time to reach a final equilibrium state. The problem of how to design the cement pathways in the field (and other physical features of the field) so that the final paths actually followed by the walkers will have certain desirable characteristics is then one of solving for parameters of the actional that will result in the desired worldline. This is a special instance of the inverse problem of how to design a COIN.

Using active walker models this way to design COINs, like action extremization in general, probably has limited applicability. Also, it is not clear how robust such a design approach might be, or whether it would be scalable and exempt from the need for hand-tailoring.

3.6 Other Related Subjects

This subsection presents a “catch-all” of other fields that have little in common with one another except that they bear some relation to COINs.

3.6.1 Stochastic Fields

An extremely well-researched body of work concerns the mathematical and numeric behavior of systems for which the probability distribution over possible future states conditioned on preceding states is explicitly provided. This work involves many aspects of Monte Carlo numerical algorithms [200], all of Markov Chains [95, 204, 254], and especially Markov fields, a topic that encompasses the Chapman-Kolmogorov equations [105] and its variants: Liouville’s equation, the Fokker-Plank equation, and the Detailed-balance equation in particular. Non-linear dynamics is also related to this body of work (see the synopsis of iterated function systems below and the synopsis of cellular automata above), as is Markov competitive decision processes (see the synopsis of game theory above).

Formally, one can cast the problem of designing a COIN as how to fix each of the conditional transition probability distributions of the individual elements of a stochastic field so that the aggregate behavior of the overall system is of a desired form.⁸ Unfortunately, almost all that is known in this area instead concerns the forward problem, of inferring aggregate behavior from a provided set of conditional distributions. Although such knowledge provides many “bits and pieces” of information about how to tackle the inverse problem, those pieces collectively cover only a very small subset of the entire space of tasks we might want the COIN to perform. In particular, they tell us very little about the case where the conditional distribution encapsulates RL algorithms.

⁸In contrast, in the field of Markov decision processes, discussed in [48], the full system may be a Markov field, but the system designer only sets the conditional transition probability distribution of a few of the field elements at most, to the appropriate “decision rules”. Unfortunately, it is hard to imagine how to use the results of this field to design COINs because of major scaling problems. Any decision process must accurately model likely future modifications to its own behavior — often an extremely daunting task [173]. What’s worse, if multiple such decision processes are running concurrently in the system, each such process must also model the others, potentially needing to model them in their full complexity.

3.6.2 Iterated Function Systems

The technique of iterated function systems [19] grew out of the field of nonlinear dynamics [223, 256, 262]. In such systems a function is repeatedly and recursively applied to itself. The most famous example is the logistic map, $x_{n+1} = rx_n(1 - x_n)$ for some r between 0 and 4 (so that x stays between 0 and 1). More generally the function along with its arguments can be vector-valued. In particular, we can construct such functions out of affine transformations of points in a Euclidean plane.

Iterated functions systems have been applied to image data. In this case the successive iteration of the function generically generates a fractal, one whose precise character is determined by the initial iteration-1 image. Since fractals are ubiquitous in natural images, a natural idea is to try to encode natural images as sets of iterated function systems spread across the plane, thereby potentially garnering significant image compression. The trick is to manage the inverse step of starting with the image to be compressed, and determining what iteration-1 image(s) and iterating function(s) will generate an accurate approximation of that image.

In the language of nonlinear dynamics, we have a dynamic system that consists of a set of iterating functions, together with a desired attractor (the image to be compressed). Our goal is to determine what values to set certain parameters of our dynamic system to so that the system will have that desired attractor. The potential relationship with COINs arises from this inverse nature of the problem tackled by iterated function systems. If the goal for a COIN can be cast as its relaxing to a particular attractor, and if the distributed computational elements are isomorphic to iterated functions, then the tricks used in iterated functions theory could be of use.

Although the techniques of iterated function systems might prove of use in designing COINs, they are unlikely to serve as a generally applicable approach to designing COINs. In addition, they do not involve RL algorithms, and often involve extensive hand-tuning.

3.6.3 Recurrent Neural Nets

A recurrent neural net consists of a finite set of “neurons” each of which has a real-valued state at each moment in time. Each neuron’s state is updated at each moment in time based on its current state and that of some of the other neurons in the system. The topology of such dependencies constitute the “inter-neuronal connections” of the net, and the associated parameters are often called the “weights” of the net. The dynamics can be either discrete or continuous (*i.e.*, given by difference or differential equations).

Recurrent nets have been investigated for many purposes [107, 128, 210, 293]. One of the more famous of these is associative memories. The idea is that given a pre-specified pattern for the (states of the neurons in the) net, there may exist inter-neuronal weights which result in a basin of attraction focussed on that pattern. If this is the case, then the net is equivalent to an associative memory, in that a complete pre-specified pattern across

all neurons will emerge under the net’s dynamics from any initial pattern that partially matches the full pre-specified pattern. In practice, one wishes the net to simultaneously possess many such pre-specified associative memories. There are many schemes for “training” a recurrent net to have this property, including schemes based on spin glasses [125, 126, 127] and schemes based on gradient descent [224].

As can the fields of cellular automata and iterated function systems, the field of recurrent neural nets can be viewed as concerning certain variants of COINs. Also like those other fields though, recurrent neural nets has shortcomings if one tries to view it as a general approach to a science of COINs. In particular, recurrent neural nets do not involve RL algorithms, and training them often suffers from scaling problems. More generally, in practice they can be hard to train well without hand-tailoring.

3.6.4 Network Theory

Packet routing in a data network [29, 129, 249, 270, 147, 110] presents a particularly interesting domain for the investigation of COINs. In particular, with such routing:

- (i) the problem is inherently distributed;
- (ii) for all but the most trivial networks it is impossible to employ global control ;
- (iii) the routers have only access to local information (routing tables);
- (iv) it constitutes a relatively clean and easily modified experimental testbed; and
- (v) there are potentially major bottlenecks induced by ‘greedy’ behavior on the part of the individual routers, which behavior constitutes a readily investigated instance of the Tragedy Of the Commons (TOC).

Many of the approaches to packet routing incorporate a variant on RL [41, 46, 54, 169, 175]. Q-routing is perhaps the best known such approach and is based on routers using reinforcement learning to select the best path [41]. Although generally successful, Q-routing is not a general scheme for inverting a global task. This is even true if one restricts attention to the problem of routing in data networks — there exists a global task in such problems, but that task is directly used to construct the algorithm.

A particular version of the general packet routing problem that is acquiring increased attention is the Quality of Service (QoS) problem, where different communication packets (voice, video, data) share the same bandwidth resource but have widely varying importances both to the user and (via revenue) to the bandwidth provider. Determining which packet has precedence over which other packets in such cases is not only based on priority in arrival time but more generally on the potential effects on the income of the bandwidth provider. In this context, RL algorithms have been used to determine routing policy, control call admission and maximize revenue by allocating the available bandwidth efficiently [46, 175].

Many researchers have exploited the noncooperative game theoretic understanding of the TOC in order to explain the bottleneck character of empirical data networks’ behavior and suggest potential alternatives to current routing schemes [26, 76, 153, 154,

162, 164, 206, 207, 243]. Closely related is work on various “pricing”-based resource allocation strategies in congestable data networks [171]. This work is at least partially based upon current understanding of pricing in toll lanes, and traffic flow in general (see below). All of these approaches are particularly of interest when combined with the RL-based schemes mentioned just above. Due to these factors, much of the current research on a general framework for COINs is directed toward the packet-routing domain (see next section).

3.6.5 Traffic Theory

Traffic congestion typifies the TOC public good problem: everyone wants to use the same resource, and all parties greedily trying to optimize their use of that resource not only worsens global behavior, but also worsens *their own* private utility (*e.g.*, if everyone disobeys traffic lights, everyone gets stuck in traffic jams). Indeed, in the well-known Braess’ paradox [21, 58, 59, 155], keeping everything else constant — including the number and destinations of the drivers — but opening a new traffic path can *increase* everyone’s time to get to their destination. (Viewing the overall system as an instance of the Prisoner’s dilemma, this paradox in essence arises through the creation of a novel ‘defect-defect’ option for the overall system.) Greedy behavior on the part of individuals also results in very rich global dynamic patterns, such as stop and go waves and clusters [119, 120].

Much of traffic theory employs and investigates tools that have previously been applied in statistical physics [119, 150, 151, 211, 216] (see subsection above). In particular, the spontaneous formation of traffic jams provides a rich testbed for studying the emergence of complex activity from seemingly chaotic states [119, 121]. Furthermore, the dynamics of traffic flow is particularly amenable to the application and testing of many novel numerical methods in a controlled environment [16, 30, 237]. Many experimental studies have confirmed the usefulness of applying insights gleaned from such work to real world traffic scenarios [119, 198, 197].

3.6.6 Topics from further afield

Finally, there are a number of other fields that, while either still nascent or not extremely closely related to COINs, are of interest in COIN design:

Amorphous computing: Amorphous computing grew out of the idea of replacing traditional computer design, with its requirements for high reliability of the components of the computer, with a novel approach in which widespread unreliability of those components would not interfere with the computation [1]. Some of its more speculative aspects are concerned with “how to program” a massively distributed, noisy system of components which may consist in part of biochemical and/or biomechanical components [152, 274]. Work here has tended to focus on schemes for how to robustly induce

desired geometric dynamics across the physical body of the amorphous computer — issue that are closely related to morphogenesis, and thereby lend credence to the idea that biochemical components are a promising approach. Especially in its limit of computers with very small constituent components, amorphous computing also is closely related to the fields of nanotechnology [71] and control of smart matter (see below).

Control of smart matter: As the prospect of nanotechnology-driven mechanical systems gets more concrete, the daunting problem of how to robustly control, power, and sustain protean systems made up of extremely large sets of nano-scale devices looms more important [111, 112, 122]. If this problem were to be solved one would in essence have “smart matter”. For example, one would be able to “paint” an airplane wing with such matter and have it improve drag and lift properties significantly.

Morphogenesis: How does a leopard embryo get its spots, or a zebra embryo its stripes? More generally, what are the processes underlying morphogenesis, in which a body plan develops among a growing set of initially undifferentiated cells? These questions, related to control of the dynamics of chemical reaction waves, are essentially special cases of the more general question of how ontogeny works, of how the genotype-phenotype mapping is carried out in development. The answers involve homeobox (as well as many other) genes [18, 73, 144, 91, 266]. Under the presumption that the functioning of such genes is at least in part designed to facilitate genetic changes that increase a species’ fitness, that functioning facilitates solution of the inverse problem, of finding small-scale changes (to DNA) that will result in “desired” large scale effects (to body plan) when propagated across a growing distributed system.

Self Organizing systems The concept of self-organization and self-organized criticality [15] was originally developed to help understand why many distributed physical systems are attracted to critical states that possess long-range dynamic correlations in the large-scale characteristics of the system. It provides a powerful framework for analyzing both biological and economic systems. For example, natural selection (particularly punctuated equilibrium [78, 109]) can be likened to self-organizing dynamical system, and some have argued it shares many the properties (*e.g.*, scale invariance) of such systems [63]. Similarly, one can view the economic order that results from the actions of human agents as a case of self-organization [65]. The relationship between complexity and self-organization is a particularly important one, in that it provides the potential laws that allow order to arise from chaos [145].

Small worlds (6 Degrees of Separation): In many distributed systems where each component can interact with a small number of “neighbors”, an important problem is how to propagate information across the system quickly and with minimal overhead. On the one extreme the neighborhood topology of such systems can exist on a completely regular grid-like structure. On the other, the topology can be totally random. In either case, certain nodes may be effectively ‘cut-off’ from other nodes if the information pathways between them are too long. Recent work has investigated “small worlds” networks

(sometimes called 6 degrees of separation) in which underlying grid-like topologies are “doped” with a scattering of long-range, random connections. It turns out that very little such doping is necessary to allow for the system to effectively circumvent the information propagation problem [183, 273].

Control theory: Adaptive control [6, 231], and in particular adaptive control involving locally weighted RL algorithms [8, 194], constitute a broadly applicable framework for controlling small, potentially inexactly modeled systems. Augmented by techniques in the control of chaotic systems [55, 67, 68], they constitute a very successful way of solving the “inverse problem” for such systems. Unfortunately, it is not clear how one could even attempt to scale such techniques up to the massively distributed systems of interest in COINs. The next section discusses in detail some of the underlying reasons why the purely model-based versions of these approaches are inappropriate as a framework for COINs.

4 A FRAMEWORK DESIGNED FOR COINs

Summarizing the discussion to this point, it is hard to see how any already extant scientific field can be modified to encompass systems meeting all of the requirements of COINs listed at the beginning of Section 3. This is not too surprising, since none of those fields were explicitly designed to analyze COINs. This section first motivates in general terms a framework that is explicitly designed for analyzing COINs. It then presents the formal nomenclature of that framework. This is followed by derivations of some of the central theorems of that framework.⁹ Finally, we present experiments that illustrate the power the framework provides for ensuring large world utility in a COIN.

4.1 Problems with a model-based approach

What mathematics might one employ to understand and design COINs? Perhaps the most natural approach, related to the stochastic fields work reviewed above, involves the following three steps:

1) First one constructs a detailed stochastic model of the COIN’s dynamics, a model parameterized by a vector θ . As an example, θ could fix the utility functions of the individual agents of the COIN, aspects of their RL algorithms, which agents communicate with each other and how, etc.

2) Next we solve for the function $f(\theta)$ which maps the parameters of the model to the resulting stochastic dynamics.

3) Cast our goal for the system as a whole as achieving a high expected value of some “world utility”. Then as our final step we would have to solve the inverse problem: we

⁹A much more detailed discussion, including intuitive arguments, proofs and fully formal definitions of the concepts discussed in this section, can be found in [283].

would have to search for a θ which, via f , results in a high value of $E(\text{world utility} \mid \theta)$.

Let's examine in turn some of the challenges each of these three steps entail:

I) We are primarily interested in very large, very complex systems, which are noisy, faulty, and often operate in a non-stationary environment. Moreover, our “very complex system” consists of many RL algorithms, all potentially quite complicated, all running simultaneously. Clearly coming up with a detailed model that captures the dynamics of all of this in an accurate manner will often be extraordinarily difficult. Moreover, unfortunately, given that the modeling is highly detailed, often the level of verisimilitude required of the model will be quite high. For example, unless the modeling of the faulty aspects of the system were quite accurate, the model would likely be “brittle”, and overly sensitive to which elements of the COIN were and were not operating properly at any given time.

II) Even for models much simpler than the ones called for in (I), solving explicitly for the function f can be extremely difficult. For example, much of Markov Chain theory is an attempt to broadly characterize such mappings. However as a practical matter, usually it can only produce potentially useful characterizations when the underlying models are quite inaccurate simplifications of the kinds of models produced in step (I).

III) Even if one can write down an f , solving the associated inverse problem is often impossible in practice.

IV) In addition to these difficulties, there is a more general problem with the model-based approach. We wish to perform our analysis on a “high level”. Our thesis is that due to the robust and adaptive nature of the individual agents' RL algorithms, there will be very broad, easily identifiable regions of θ space all of which result in excellent $E(\text{world utility} \mid \theta)$, and that these regions will not depend on the precise learning algorithms used to achieve the low-level tasks (cf. the list at the beginning of Section 3). To fully capitalize on this, one would want to be able to slot in and out different learning algorithms for achieving the low-level tasks without having to redo our entire analysis each time. However in general this would be possible with a model-based analysis only for very carefully designed models (if at all). The problem is that the result of step (3), the solution to the inverse problem, would have to concern aspects of the COIN that are (at least approximately) invariant with respect to the precise low-level learning algorithms used. Coming up with a model that has this property while still avoiding problems (I-III) is usually an extremely daunting challenge.

Fortunately, there is an alternative approach which avoids the difficulties of detailed modeling. Little modeling of any sort ever is used in this alternative, and what modeling does arise has little to do with dynamics. In addition, any such modeling is extremely high-level, intended to serve as a decent approximation to almost any system having “reasonable” RL algorithms, rather than as an accurate model of one particular system.

We call any framework based on this alternative a **descriptive framework**. In such a framework one identifies certain **salient characteristics** of COINs, which are

characteristics of a COIN’s entire worldline that one strongly expects to find in COINs that have large world utility. Under this expectation, one makes the assumption that if a COIN is explicitly modified to have the salient characteristics (for example in response to observations of its run-time behavior), then its world utility will benefit. So long as the salient characteristics are (relatively) easy to induce in a COIN, then this assumption provides a ready indirect way to cause that COIN to have large world utility.

An assumption of this nature is the central leverage point that a descriptive framework employs to circumvent detailed modeling. Under it, if the salient characteristics can be induced with little or no modeling (e.g., via heuristics that aren’t rigorously and formally justified), then they provide an indirect way to improve world utility without recourse to detailed modeling. In fact, since one does not use detailed modeling in a descriptive framework, it may even be that one does not have a fully rigorous mathematical proof that the central assumption holds in a particular system for one’s choice of salient characteristics. One may have to be content with reasonableness arguments not only to justify one’s scheme for inducing the salient characteristics, but for making the assumption that characteristics are correlated with large world utility in the first place.¹⁰ Of course, the trick in the descriptive framework is to choose salient characteristics that both have a beneficial relationship with world utility and that one expects to be able to induce with relatively little detailed modeling of the system’s dynamics.

4.2 Nomenclature

There exist many ways one might try to design a descriptive framework. In this subsection we present nomenclature needed for a (very) cursory overview of one of them. (See [283] for a more detailed exposition, including formal proofs.)

This overview concentrates on the four salient characteristics of intelligence, learnability, factoredness, and the wonderful life utility, all defined below. Intelligence is a quantification of how well an RL algorithm performs. We want to do whatever we can to help those algorithms achieve high values of their utility functions. Learnability is a characteristic of a utility function that one would expect to be well-correlated with how well an RL algorithm can learn to optimize it. A utility function is also factored if whenever its value increases, the overall system benefits. Finally, wonderful life utility is an example of a utility function that is both learnable and factored.

After the preliminary definitions below, this section formalizes these four salient characteristics, derives several theorems relating them, and illustrates in some computer experiments how those theorems can be used to help the system achieve high world utility.

¹⁰Despite only being implicit, such reasonableness arguments are all that underpins fields like non-Bayesian learning. See [279, 281, 280].

4.2.1 Preliminary Definitions

1) We refer to an RL algorithm by which an individual component of the COIN modifies its behavior as a **microlearning** algorithm. We refer to the initial construction of the COIN, potentially based upon salient characteristics, as the COIN **initialization**. We use the phrase **macrolearning** to refer to externally imposed run-time modifications to the COIN which are based on statistical inference concerning salient characteristics of the running COIN.

2) For convenience, we take time, t , to be discrete and confined to the integers, \mathbb{Z} . When referring to COIN initialization, we implicitly have a lower bound on t , which without loss of generality we take to be less than or equal to 0.

3) All variables that have any effect on the COIN are identified as components of Euclidean-vector-valued **states** of various discrete **nodes**. As an important example, if our COIN consists in part of a computational “agent” running a microlearning algorithm, the precise configuration of that agent at any time t , including all variables in its learning algorithm, all actions directly visible to the outside world, all internal parameters, all values observed by its probes of the surrounding environment, etc., all constitute the state vector of a node representing that agent. We define $\underline{\zeta}_{\eta,t}$ to be a vector in the Euclidean vector space $\underline{\mathbf{Z}}_{\eta,t}$, where the components of $\underline{\zeta}_{\eta,t}$ give the state of node η at time t . The i ’th component of that vector is indicated by $\underline{\zeta}_{\eta,t;i}$.

Observation 3.1: In practice, many COINs will involve variables that are most naturally viewed as discrete and symbolic. In such cases, we must exercise some care in how we choose to represent those variables as components of Euclidean vectors. There is nothing new in this; the same issue arises in modern work on applying neural nets to inherently symbolic problems. In our COIN framework, we will usually employ the same resolution of this issue employed in neural nets, namely representing the possible values of the discrete variable with a unary representation in a Euclidean space. Just as with neural nets, values of such vectors that do not lie on the vertices of the unit hypercube are not meaningful, strictly speaking. Fortunately though, just as with neural nets, there is almost always a most natural way to extend the definitions of any function of interest (like world utility) so that it is well-defined even for vectors not lying on those vertices. This allows us to meaningfully define partial derivatives of such functions with respect to the components of $\underline{\zeta}$, partial derivatives that we will evaluate at the corners of the unit hypercube.

4) For notational convenience, we define $\underline{\zeta}_t \in \underline{\mathbf{Z}}_t$ to be the vector of the states of all nodes at time t ; $\underline{\zeta}_{\eta,t} \in \underline{\mathbf{Z}}_{\eta,t}$ to be the vector of the states of all nodes other than η at time t ; and $\underline{\zeta} \equiv \underline{\zeta}_t \in \underline{\mathbf{Z}} \equiv \underline{\mathbf{Z}}'$ to be the entire vector of the states of all nodes at all times. $\underline{\mathbf{Z}}$ is infinite-dimensional in general, and usually assumed to be a Hilbert space. We will often assume that all spaces $\underline{\mathbf{Z}}_t$ over all times t are isomorphic to a space $\underline{\mathbf{Z}}^{(0)}$, i.e., $\underline{\mathbf{Z}}$ is a Cartesian product of copies of $\underline{\mathbf{Z}}^{(0)}$.

Also for notational convenience, we define gradients using ∂ -shorthand. So for exam-

ple, $\partial_{\underline{\zeta}_t} F(\underline{\zeta})$ is the vector of the partial derivative of $F(\underline{\zeta})$ with respect to the components of $\underline{\zeta}_t$. Also, we will sometimes treat the symbol “ t ” specially, as delineating a range of components of $\underline{\zeta}$. So for example an expression like “ $\underline{\zeta}_{t < t'}$ ” refers to all components $\underline{\zeta}_t$ with $t < t'$.

5) To avoid confusion with the other uses of the comma operator, we will often use $\mathbf{x} \bullet \mathbf{y}$ rather than (\mathbf{x}, \mathbf{y}) to indicate the vector formed by concatenating the two ordered sets of vector components \mathbf{x} and \mathbf{y} . For example, $\underline{\zeta}_{\eta, t < 0} \bullet \underline{\zeta}_{\eta', t > 0}$ refers to the vector formed by concatenating those components of the worldline $\underline{\zeta}$ involving node η for times less than 0 with those components involving node η' that have times greater than 0.

6) We take the universe in which our COIN operates to be completely deterministic. This is certainly the case for any COIN that operates in a digital system, even a system that emulates analog and/or stochastic processes (*e.g.*, with a pseudo-random number generator). More generally, this determinism reflects the fact that since the real world obeys (deterministic) physics, *any* real-world system, be it a COIN or something else, is, ultimately, embedded in a deterministic system.¹¹

The perspective to be kept in mind here is that of nonlinear time-series analysis. A physical time series typically reflects a few degrees of freedom that are projected out of the underlying space in which the full system is deterministically evolving, an underlying space that is actually extremely high-dimensional. This projection typically results in an illusion of stochasticity in the time series.

7) Formally, to reflect this determinism, first we bundle all variables we are not directly considering — but which nonetheless affect the dynamics of the system — as components of some catch-all **environment node**. So for example any “noise processes” and the like affecting the COIN’s dynamics are taken to be inputs from a deterministic, very high-dimensional environment that is potentially chaotic and is never directly observed [86]. Given such an environment node, we then stipulate that for all t, t' such that $t' > t$, $\underline{\zeta}_{t'}$ sets $\underline{\zeta}_t$ uniquely.

Observation 7.1: When nodes are “computational devices”, often we must be careful to specify the physical extent of those devices. Such a node may just be the associated CPU, or it may be that CPU together with the main RAM, or it may include an external storage device. Almost always, the border of the device η will end before any external system that η is “observing” begins. This means that since at time t η only knows the value of $\underline{\zeta}_{\eta, t}$, its “observational knowledge” of that external system is indirect. That knowledge reflects a coupling between $\underline{\zeta}_{\eta, t}$ and $\underline{\zeta}_{\eta, t'}$, a coupling that is induced by the dynamical evolution of the system from preceding moments up to the time t . If the dynamics does not force such a coupling, then η has no observational knowledge of the

¹¹This determinism holds even for systems with an explicitly quantum mechanical character. Quantum mechanical systems evolve according to Schrodinger’s equation, which is purely deterministic; as is now well-accepted, the “stochastic” aspect of quantum mechanics can be interpreted as an epiphenomenon of Schrodinger’s equation that arises when the Hamiltonian has an “observational” or “entangling” coupling between some of its variables [84, 297, 106], a coupling that does not obviate the underlying determinism.

outside world.

8) We express the dynamics of our system by writing $\zeta_{t' \geq t} = C(\zeta_t)$. (In this paper there will be no need to be more precise and specify the precise dependency of $C(\cdot)$ on t and/or t' .) We define $\{C\}$ to be a set of constraint equations enforcing that dynamics, and also, more generally, fixing the entire manifold C of vectors $\underline{\zeta} \in \underline{\mathbf{Z}}$ that we consider to be ‘allowed’. So C is a subset of the set of all $\underline{\zeta} \in \underline{\mathbf{Z}}$ that are consistent with the deterministic laws governing the COIN, *i.e.*, that obey $\zeta_{t' \geq t} = C(\zeta_t) \forall t, t'$. We generalize this notation in the obvious way, so that (for example) $C_{t \geq t_0}$ is the manifold consisting of all vectors $\underline{\zeta}_{t \geq t_0} \in \underline{\mathbf{Z}}_{t \geq t_0}$ that are projections of a vector in C .

Observation 8.1: Note that $C_{t \geq t_0}$ is parameterized by ζ_{t_0} , due to determinism. Note also that whereas $C(\cdot)$ is defined for any argument of the form $\underline{\zeta}_t \in \underline{\mathbf{Z}}_t$ for some t (*i.e.*, we can evolve any point forward in time), in general not all $\underline{\zeta}_t \in \underline{\mathbf{Z}}_t$ lie in C_t . In particular, there may be extra restrictions constraining the possible states of the system beyond those arising from its need to obey the relevant dynamical laws of physics. Finally, whenever trying to express a COIN in terms of the framework presented here, it is a good rule to try to write out the constraint equations explicitly to check that what one has identified as the space $\underline{\mathbf{Z}}_t$ contains all quantities needed to uniquely fix the future state of the system.

Observation 8.2: We do not want to have $\underline{\mathbf{Z}}$ be the phase space of every particle in the system. We will instead usually have $\underline{\mathbf{Z}}$ consist of variables that, although still evolving deterministically, exist at a larger scale of granularity than that of individual particles (*e.g.*, thermodynamic variables in the thermodynamic limit). However we will often be concerned with physical systems obeying entropy-driven dynamic processes that are contractive at this high level of granularity. Examples are any of the many-to-one mappings that can occur in digital computers, and, at a finer level of granularity, any of the error-correcting processes in the electronics of such a computer that allow it to operate in a digital fashion. Accordingly, although the dynamics of our system will always be deterministic, it need not be invertible.

Observation 8.3: Intuitively, in our mathematics, all behavior across time is pre-fixed. The COIN is a single fixed worldline through $\underline{\mathbf{Z}}$, with no “unfolding of the future” as the die underlying a stochastic dynamics get cast. This is consistent with the fact that we want the formalism to be purely descriptive, relating different properties of any single, fixed COIN’s history. We will often informally refer to “changing a node’s state at a particular time”, or to a microlearner’s “choosing from a set of options”, and the like. Formally, in all such phrases we are really comparing different worldlines, with the indicated modification distinguishing those worldlines.

Observation 8.4: Since the dynamics of any real-world COIN is deterministic, so is the dynamics of any component of the COIN, and in particular so is any learning algorithm running in the COIN, ultimately. However that does not mean that those deterministic components of the COIN are not allowed to be “based on”, or “motivated by” stochastic

concepts. The *motivation* behind the algorithms run by the components of the COIN does not change their underlying nature. Indeed, in our experiments below, we explicitly have the reinforcement learning algorithms that are trying to maximize private utility operate in a (pseudo-) probabilistic fashion, with pseudo-random number generators and the like.

More generally, the deterministic nature of our framework does not preclude our superimposing probabilistic elements on top of that framework, and thereby generating a stochastic extension of our framework. Exactly as in statistical physics, a stochastic nature can be superimposed on our space of deterministic worldlines, potentially by adopting a degree of belief perspective on “what probability means” [282, 139]. Indeed, the macrolearning algorithms we investigate below implicitly involve such a superimposing; they implicitly assume a probabilistic coupling between the (statistical estimate of the) correlation coefficient connecting the states of a pair of nodes and whether those nodes are in the one another’s “effect set”.

Similarly, while it does not salient characteristics that involve probability distributions, the descriptive framework does not preclude such characteristics either. As an example, the “intelligence” of an agent’s particular action, formally defined below, measures the fraction of alternative actions an agent could have taken that would have resulted in a lower utility value. To define such a fraction requires a measure across the space of such alternative actions, even if only implicitly. Accordingly, intelligence can be viewed as involving a probability distribution across the space of potential actions.

In this paper though, we concentrate on the mathematics that obtains before such probabilistic concerns are superimposed. Whereas the deterministic analysis presented here is related to game-theoretic structures like Nash equilibria, a full-blown stochastic extension would in some ways be more related to structures like correlated equilibria [9].

9) Formally, there is a lot of freedom in setting the boundary between what we call “the COIN”, whose dynamics is determined by C , and what we call “macrolearning”, which constitutes perturbations to the COIN instigated from “outside the COIN”, and which therefore is *not* reflected in C . As an example, in much of this paper, we have clearly specified microlearners which are provided fixed private utility functions that they are trying to maximize. In such cases usually we will implicitly take C to be the dynamics of the system, macrolearning and all, *for fixed private utilities* that are specified in $\underline{\zeta}$. For example, $\underline{\zeta}$ could contain, for each microlearner, the bits in an associated computer specifying the subroutine that that microlearner can call to evaluate what its private utility would be for some full worldline $\underline{\zeta}$.

Macrolearning overrides C , and in this situation it refers (for example) to any statistical inference process that modifies the private utilities at run-time to try to induce the desired salient characteristics. Concretely, this would involve modifications to the bits $\{b_i\}$ specifying each microlearner i ’s private utility, modifications that are *not* accounted for in C , and that are potentially based on variables that are not reflected in \underline{Z} . Since

C does not reflect such macrolearning, when trying to ascertain C based on empirical observation (as for example when determining how best to modify the private utilities), we have to take care to distinguish which part of the system’s observed dynamics is due to C and which part instead reflects externally imposed modifications to the private utilities.

More generally though, other boundaries between the COIN and macrolearning-based perturbations to it are possible, reflecting other definitions of \underline{Z} , and other interpretations of the elements of each $\underline{\zeta} \in \underline{Z}$. For example, say that under the perspective presented in the previous paragraph, the private utility is a function of some components s of $\underline{\zeta}$, components that do not include the $\{b_i\}$. Now modify this perspective so that in addition to the dynamics of other bits, C also encapsulates the dynamics of the bits $\{b_i\}$. Having done this, we could still view each private utility as being fixed, but rather than take the bits $\{b_i\}$ as “encoding” the subroutine that specifies the private utility of microlearner i , we would treat them as “parameters” specifying the functional dependence of the (fixed) private utility on the components of $\underline{\zeta}$. In other words, formally, they constitute an extra set of arguments to i ’s private utility, in addition to the arguments s . Alternatively, we could simply say that in this situation our private utilities are time-indexed, with i ’s private utility at time t determined by $\{b_{i,t}\}$, which in turn is determined by evolution under C . Under either interpretation of private utility, any modification under C to the bits specifying i ’s utility-evaluation subroutine constitutes dynamical laws by which the parameters of i ’s microlearner evolves in time. In this case, macrolearning would refer to some further removed process that modifies the evolution of the system in a way not encapsulated in C .

For such alternative definitions of C/\underline{Z} , we have a different boundary between the COIN and macrolearning, and we must scrutinize different aspects of the COIN’s dynamics to infer C . Whatever the boundary, the mathematics of the descriptive framework, including the mathematics concerning the salient characteristics, is restricted to a system evolving according to C , and explicitly does not account for macrolearning. This is why the strategy of trying to improve world utility by using macrolearning to try to induce salient characteristics is almost always ultimately based on an assumption rather than a proof.

10) We are provided with some Von Neumann **world utility** $G : \underline{Z} \rightarrow \mathcal{R}$ that ranks the various conceivable worldlines of the COIN. Note that since the environment node is never directly observed, we implicitly assume that the world utility is not directly (!) a function of its state. Our mathematics will not involve G alone, but rather the relationship between G and various sets of **personal utilities** $g_{\eta,t} : \underline{Z} \otimes Z \rightarrow \mathcal{R}$.

Intuitively, as discussed below, for many purposes such personal utilities are equivalent to arbitrary “virtual” versions of the private utilities mentioned above. In particular, it is only private utilities that will occur within any macrolearning computer algorithms that may be running in the COIN as manifested in C . Personal utilities are external

mathematical constructions that the COIN framework employs to analyze the behavior of the system. They can be involved in learning processes, but only as tools that are employed outside of the COIN’s evolution under C , *i.e.*, only in macrolearning. (For example, analysis of them can be used to modify the private utilities.)

Observation 10.1: These utility definitions are very broad. In particular, they do not require casting of the utilities as discounted sums. Note also that our world utility is not indexed by t . Again reflecting the descriptive, worldline character of the formalism, we simply assign a single value to an entire worldline of the system, implicitly assuming that one can always say which of two candidate worldlines are preferable. So given some “present time” t_0 , issues like which of two “potential futures” $\underline{\zeta}_{t>t_0}, \underline{\zeta}'_{t>t_0}$ is preferable are resolved by evaluating the relevant utility at two associated points $\underline{\zeta}$ and $\underline{\zeta}'$, where the $t > t_0$ components of those points are the futures indicated, and the two points share the same (usually implicit) $t \leq t_0$ “past” components.

This time-independence of G automatically avoids formal problems that can occur with general (*i.e.*, not necessarily discounted sum) time-indexed utilities, problems like having what’s optimal at one moment in time conflict with what’s optimal at other moments in time.¹² For personal utilities such formal problems are often irrelevant however. Before we begin our work, we as COIN designers must be able to rank all possible worldlines of the system at hand, to have a well-defined design task. That is why world utility cannot be time-indexed. However if a particular microlearner’s goal keeps changing in an inconsistent way, that simply means that that microlearner will grow “confused”. From our perspective as COIN designers, there is nothing *a priori* unacceptable about such confusion. It may even result in better performance of the system as a whole, in which case we would actually want to induce it. Nonetheless, for simplicity, in most of this paper we will have all $g_{\eta,t}$ be independent of t , just like world utility.

World utility is defined as that function that we are ultimately interested in optimizing. In conventional RL it is a discounted sum, with the sum starting at time t . In other words, conventional RL has a time-indexed world utility. It might seem that in this at least, conventional RL considers a case that has more generality than that of the COIN framework presented here. (It obviously has less generality in that its world utility is restricted to be a discounted sum.) In fact though, the apparent time-indexing of conventional RL is illusory, and the time-dependent discounted sum world utility of conventional RL is actually a special case of the non-time-indexed world utility of our COIN framework. To see this formally, consider any (time-independent) world utility

¹²Such conflicts can be especially troublesome when they interfere with our defining what we mean by an “optimal” set of actions by the nodes at a particular time t . The effects of the actions by the nodes, and therefore whether those actions are “optimal” or not, depends on the future actions of the nodes. However if they too are to be “optimal”, according to their world-utility, those future actions will depend on *their* futures. So we have a potentially infinite regress of differing stipulations of what “optimal” actions at time t entails.

$G(\underline{\zeta})$ that equals $\sum_{t=0}^{\infty} \gamma^t r(\underline{\zeta}_t)$ for some function $r(\cdot)$ and some positive constant γ with magnitude less than 1. Then for any $t' > 0$ and any $\underline{\zeta}'$ and $\underline{\zeta}''$ where $\underline{\zeta}'_{t < t'} = \underline{\zeta}''_{t < t'}$, $\text{sgn}[G(\underline{\zeta}') - G(\underline{\zeta}'')] = \text{sgn}[\sum_{t=0}^{\infty} \gamma^t r(\underline{\zeta}'_t) - \sum_{t=0}^{\infty} \gamma^t r(\underline{\zeta}''_t)]$. Conventional RL merely expresses this in terms of time-dependent utilities $u_{t'}(\underline{\zeta}_{t > t'}) \equiv \sum_{t=t'}^{\infty} \gamma^{t-t'} r(\underline{\zeta}_t)$ by writing $\text{sgn}[G(\underline{\zeta}') - G(\underline{\zeta}'')] = \text{sgn}[u_{t'}(\underline{\zeta}') - u_{t'}(\underline{\zeta}'')] for all t' . Since utility functions are, by definition, only unique up to the relative orderings they impose on potential values of their arguments, we see that conventional RL's use of a time-dependent discounted sum world utility $u_{t'}$ is identical to use of a particular time-independent world utility in our COIN framework.$

11) As mentioned above, there may be variables in each node's state which, under one particular interpretation, represent the “utility functions” that the associated microlearner's computer program is trying to extremize. When there are such components of $\underline{\zeta}$, we refer to the utilities they represent as **private utilities**. However even when there are private utilities, formally we allow the personal utilities to differ from them. The personal utility functions $\{g_\eta\}$ do not exist “inside the COIN”; they are not specified by components of $\underline{\zeta}$. This separating of the private utilities from the $\{g_\eta\}$ will allow us to avoid the teleological problem that one may not always be able to explicitly identify “the” private utility function reflected in $\underline{\zeta}$ such that a particular computational device can be said to be a microlearner “trying to increase the value of its private utility”. To the degree that we can couch the theorems purely in terms of personal rather than private utilities, we will have successfully adopted a purely behaviorist approach, without any need to interpret what a computational device is “trying to do”.

Despite this formal distinction though, often we will implicitly have in mind deploying the personal utilities onto the microlearners as their private utilities, in which case the terms can usually be used interchangeably. The context should make it clear when this is the case.

4.2.2 Intelligence

We will need to quantify how well the entire system performs in terms of G . To do this requires a measure of the performance of an arbitrary worldline $\underline{\zeta}$, for an arbitrary utility function, under arbitrary dynamic laws C . Formally, such a measure is a mapping from three arguments to \mathbf{R} .

Such a measure will also allow us to quantify how well each microlearner performs in purely behavioral terms, in terms of its personal utility. (In our behaviorist approach, we do not try to make specious distinctions between whether a microlearner's performance is due to its level of “innate sophistication”, or rather due to dumb luck — all that matters is the quality of its behavior as reflected in its utility value for the system's worldline.) This behaviorism in turn will allow us to avoid having private utilities explicitly arise in our theorems (although they still arise frequently in pedagogical discussion). Even when private utilities exist, there will be no formal need to explicitly identify some components

of $\underline{\zeta}$ as such utilities. Assuming a node’s microlearner is competent, the fact that it is trying to optimize some particular private utility U will be manifested in our performance measure’s having a large value at $\underline{\zeta}$ for C for that utility U .

The problem of how to formally define such a performance measure is essentially equivalent to the problem of how to quantify bounded rationality in game theory. Some of the relevant work in game theory, for example that involving ‘trembling hand equilibria’ or ‘ ϵ equilibria’ [97] is concerned with refinements or modifications of Nash equilibria (see also [161]). Rather than a behaviorist approach, such work adopts a strongly teleological perspective on rationality. In general, such work is only applicable to those situations where the rationality is bounded due to the precise causal mechanisms investigated in that work. Most of the other game-theoretic work first models (!) the microlearner, as some extremely simple computational device (*e.g.*, a deterministic finite automaton (DFA)). One then assumes that the microlearner performs perfectly for that device, so that one can measure that learner’s performance in terms of some computational capacity measure of the model (*e.g.*, for a DFA, the number of states of that DFA) [100, 202, 230].¹³ However, if taken as renditions of real-world computer-based microlearners — never mind human microlearners — the models in this approach are often extremely abstracted, with many important characteristics of the real learners absent or distorted. In addition, there is little reason to believe that any results arising from this approach would not be highly dependent on the model choice and on the associated representation of computational capacity. Yet another disadvantage is that this approach concentrates on perfect, fully rational behavior of the microlearners, within their computational restrictions.

We would prefer a less model-dependent approach, especially given our wish that the performance measure be based solely on the utility function at hand, $\underline{\zeta}$, and C . Now we don’t want our performance measure to be a “raw” utility value like $g_\eta(\underline{\zeta})$, since that is not invariant with respect to monotonic transformations of g_η . Similarly, we don’t want to penalize the microlearner for not achieving a certain utility value if that value was impossible to achieve not due to the microlearner’s shortcomings, but rather due to C and the actions of other nodes. A natural way to address these concerns is to generalize the game-theoretic concept of “best-response strategy” and consider the problem of how well η performs *given the actions of the other nodes*. Such a measure would compare the utility ultimately induced by each of the possible states of η at some particular time, which without loss of generality we can take to be 0, to that induced by the actual state $\underline{\zeta}_{\eta,0}$. In other words, we would compare the utility of the actual worldline $\underline{\zeta}$ to those of a set of alternative worldlines $\underline{\zeta}'$, where $\underline{\zeta}_{\eta,0} = \underline{\zeta}'_{\eta,0}$, and use those comparisons to quantify the quality of η ’s performance.

Now we are only concerned with comparing the effects of replacing $\underline{\zeta}$ with $\underline{\zeta}'$ on *future* contributions to the utility. But if we allow arbitrary $\underline{\zeta}'_{t<0}$, then in and of themselves

¹³Some of the more popular model-based scenarios for investigating bounded rationality, like ‘fictitious play’ (see the game theory section above), do not even stipulate one particular way to quantify that rationality.

the difference between those past components of $\underline{\zeta}'$ and those of $\underline{\zeta}$ can modify the value of the utility, regardless of the effects of any difference in the future components. Our presumption is that for many COINs of interest we can avoid this conundrum by restricting attention to those $\underline{\zeta}'$ where $\underline{\zeta}'_{t<0}$ differs from $\underline{\zeta}_{t<0}$ only in the internal parameters of η 's microlearner, differences that only at times $t \geq 0$ manifest themselves in a form the utility is concerned with. (In game-theoretic terms, such “internal parameters” encode full extensive-form strategies, and we only consider changes to the vertices at or below the $t = 0$ level in the tree of an extensive-form strategy.)

Although this solution to our conundrum is fine when we can apply it, we don't want to restrict the formalism so that it can only concern systems having computational algorithms which involve a clearly pre-specified set of extensive strategy “internal parameters” and the like. So instead, we formalize our presumption behaviorally, even for computational algorithms that do not have explicit extensive strategy internal parameters. Since changing the internal parameters doesn't affect the $t < 0$ components of $\underline{\zeta}_\eta$, *that the utility is concerned with*, and since we are only concerned with changes to $\underline{\zeta}$ that affect the utility, we simply elect to not change the $t < 0$ values of the internal parameters of $\underline{\zeta}_\eta$ at all. In other words, we leave $\underline{\zeta}_{\eta,t<0}$ unchanged. The advantage of this stipulation is that we can apply it just as easily whether η does or doesn't have any “internal parameters” in the first place.

So in quantifying the performance of η for behavior given by $\underline{\zeta}$ we compare $\underline{\zeta}$ to a set of $\underline{\zeta}'$, a set restricted to those $\underline{\zeta}'$ sharing $\underline{\zeta}$'s past: $\underline{\zeta}'_{t<0} = \underline{\zeta}_{t<0}$, $\underline{\zeta}'_{\eta,0} = \underline{\zeta}_{\eta,0}$, and $\underline{\zeta}'_{t \geq 0} \in C_{t \geq 0}$. Since $\underline{\zeta}'_{\eta,0}$ is free to vary (reflecting the possible changes in the state of η at time 0) while $\underline{\zeta}'_{t<0}$ is not, $\underline{\zeta}' \notin C$, in general. We may even wish to allow $\underline{\zeta}'_{t \geq 0} \notin C_{t \geq 0}$ in certain circumstances. (Recall that C may reflect other restrictions imposed on allowed worldlines besides adherence to the underlying dynamical laws, so simply obeying those laws does not suffice to ensure that a worldline lies on C .) In general though, our presumption is that as far as utility values are concerned, considering these dynamically impossible $\underline{\zeta}'$ is equivalent to considering a more restricted set of $\underline{\zeta}'$ with “modified internal parameters”, all of which are $\in C$.

We now present a formalization of this performance measure. Given C and a measure $d\mu(\underline{\zeta}_0)$ demarcating what points in $\mathbf{Z}_{\eta,0}$ we are interested in, we define the ($t = 0$) **intelligence** for node η of a point $\underline{\zeta}$ with respect to a utility U as follows:

$$\epsilon_{\eta,U}(\underline{\zeta}) \equiv \int d\mu(\underline{\zeta}'_0) \Theta[U(\underline{\zeta}) - U(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}'_0))] \cdot \delta(\underline{\zeta}'_{\eta,0} - \underline{\zeta}_{\eta,0}) \quad (1)$$

where $\Theta(\cdot)$ is the Heaviside theta function which equals 0 if its argument is below 0 and equals 1 otherwise, $\delta(\cdot)$ is the Dirac delta function, and we assume that $\int d\mu(\underline{\zeta}'_{\eta,0}) = 1$.

Intuitively, $\epsilon_{\eta,U}(\underline{\zeta})$ measures the fraction of alternative states of η which, if η had been in those states at time 0, would either degrade or not improve η 's performance (as measured by U). Sometimes in practice we will only want to consider changes in those components of $\underline{\zeta}_{\eta,0}$ that we consider as “free to vary”, which means in particular that

those changes are consistent with C and the state of the external world, $\underline{\zeta}_{\eta,0}$. (This consistency ensures that η 's observational information concerning the external world is correct; see Observation 7.1 above.) Such a restriction means that even though $\underline{\zeta}_0$ may not be consistent with C and $\underline{\zeta}_{t<0}$, by itself it is still consistent with C ; in quantifying the quality of a particular $\underline{\zeta}_{\eta,0}$. So we don't compare our point to other $\underline{\zeta}_0$ that are physically impossible, no matter what the past is. Any such restrictions on what changes we are considering are reflected implicitly in intelligence, in the measure $d\mu$.

As an example of intelligence, consider the situation where for each player η , the support of the measure $d\mu(\underline{\zeta}'_0)$ extends over all possible actions that η could take that affect the ultimate value of its personal utility, g_η . In this situation we recover conventional full rationality game theory involving Nash equilibria, as the analysis of scenarios in which the intelligence of each player η with respect to g_η equals 1.¹⁴ As an alternative, we could for each η restrict $d\mu(\underline{\zeta}'_0)$ to some limited "set of actions that η actively considers". This provides us with an "effective Nash equilibrium" at the point $\underline{\zeta}$ where each $\epsilon_{\eta,g_\eta}(\underline{\zeta})$ equals 1, in the sense that *as far it's concerned*, each player η has played a best possible action at such a point. As yet another alternative, we could restrict each $d\mu(\underline{\zeta}'_0)$ to some infinitesimal neighborhood about $\underline{\zeta}_0$, and thereby define a "local Nash equilibrium" by having $\epsilon_{\eta,g_\eta}(\underline{\zeta}) = 1$ for each player η .

In general, competent greedy pursuit of private utility U by the microlearner controlling node η means that the intelligence of η for personal utility U , $\epsilon_{\eta,U}(\underline{\zeta})$, is close to 1. Accordingly, we will often refer interchangeably to a capable microlearner's "pursuing private utility U ", and to its having high intelligence for personal utility U . Alternatively, if the microlearner for node η is incompetent, then it may even be that "by luck" its intelligence for some personal utility $\{g_\eta\}$ exceeds its intelligence for the different private utility that it's actually trying to maximize, U_η .

Say that we expect that a particular microlearner is "smart", in that it is more likely to have high rather than low intelligence. We can model this by saying that given a particular $\underline{\zeta}_{\eta,0}$, the conditional probability that $\underline{\zeta}_{\eta,0} = z$ is a monotonically increasing function of $\epsilon_{\eta,g_\eta}(\underline{\zeta}_{t<0} \bullet C(z \bullet \underline{\zeta}_{\eta,0}))$. Since for a given $\underline{\zeta}_{\eta,0}$ the intelligence ϵ_{η,g_η} is a monotonically increasing function of g_η , this modelling assumption means that the probability that $\underline{\zeta}_{\eta,0} = z$ is a monotonically increasing function of $g_\eta(\underline{\zeta}_{t<0} \bullet C(z \bullet \underline{\zeta}_{\eta,0}))$. An alternative weaker model is to only stipulate that the probability of having a particular pair $(\underline{\zeta}_{\eta,0}, \underline{\zeta}_{\eta,0})$ with ϵ_{η,g_η} equal to z is a monotonically increasing function of z . (This probability is an integral over a joint distribution, rather than a conditional

¹⁴As an alternative to such fully rational games, one can define a bounded rational game as one in which the intelligences equal some vector $\vec{\epsilon}$ whose components need not all equal 1. Many of the theorems of conventional game theory can be directly carried over to such bounded-rational games [284] by redefining the utility functions of the players. In other words, much of conventional full rationality game theory applies even to games with bounded rationality, under the appropriate transformation. This result has strong implications for the legitimacy of the common criticism of modern economic theory that its assumption of full rationality does not hold in the real world, implications that extend significantly beyond the Sonnenschein-Mantel-Debreu Theorem equilibrium aggregate demand theorem [177].

distribution, as in the original model.) In either case, the “better” the microlearner, the more tightly peaked the associated probability distribution over intelligence values is.

Any two utility functions that are related by a monotonically increasing transformation reflect the same preference ordering over the possible arguments of those functions. Since it is only that ordering that we are ever concerned with, we would like to remove this degeneracy by “normalizing” all utility functions. In other words, we would like to reduce any equivalence set of utility functions that are monotonic transformations of one another to a canonical member of that set. To see what this means in the COIN context, fix $\underline{\zeta}_\eta$. Viewed as a function from $\underline{\mathbf{Z}}_\eta \rightarrow \mathcal{R}$, $\epsilon_{\eta,U}(\underline{\zeta}_\eta, \cdot)$ is itself a utility function, one that is a monotonically increasing function of U . (It says how well η would have performed for all vectors $\underline{\zeta}_\eta$.) Accordingly, the integral transform taking U to $\epsilon_{\eta,U}(\underline{\zeta}_\eta, \cdot)$ is a (contractive, non-invertible) mapping from utilities to utilities. Applied to any member of a utility in U ’s equivalence set, this mapping produces the same image utility, one that is also in that equivalence set. It can be proven that any mapping from utilities to utilities that has this and certain other simple properties must be such an integral transform. In this, intelligence is the unique way of “normalizing” Von Neumann utility functions.

For those conversant with game theory, it is worth noting some of the interesting aspects that ensue from this normalizing nature of intelligences. At any point $\underline{\zeta}$ that is a Nash equilibrium in the set of personal utilities $\{g_\eta\}$, all intelligences $\epsilon_{\eta,g_\eta}(\underline{\zeta})$ must equal 1. Since that is the maximal value any intelligence can take on, a Nash equilibrium in the $\{g_\eta\}$ is a Pareto optimal point in the associated intelligences (for the simple reason that no deviation from such a $\underline{\zeta}$ can raise any of the intelligences). Conversely, if there exists at least one Nash equilibrium in the $\{g_\eta\}$, then there is not a Pareto optimal point in the $\{\epsilon_{\eta,g_\eta}(\underline{\zeta})\}$ that is not a Nash equilibrium.

Now restrict attention to systems with only a single instant of time, *i.e.*, single-stage games. Also have each of the (real-valued) components of each $\underline{\zeta}_\eta$ be a mixing component of an associated one of η ’s potential strategies for some underlying finite game. Then have $g_\eta(\underline{\zeta})$ be the associated expected payoff to η . (So the payoff to η of the underlying pure strategies is given by the values of $g_\eta(\underline{\zeta})$ when $\underline{\zeta}$ is a unit vector in the space $\underline{\mathbf{Z}}_\eta$ of η ’s possible states.) Then we know that there must exist at least one Nash equilibrium in the $\{g_\eta\}$. Accordingly, in this situation the set of Nash equilibria in the $\{g_\eta\}$ is identical to the set of points that are Pareto optimal in the associated intelligences. (See Eq. 5 in the discussion of factored systems below.)

4.2.3 Learnability

Intelligence can be a difficult quantity to work with, unfortunately. As an example, fix η , and consider any (small region centered about some) $\underline{\zeta}$ along with some utility U , where $\underline{\zeta}$ is not a local maximum of U . Then by increasing the values U takes on in that small region we will increase the intelligence $\epsilon_{\eta,U}(\underline{\zeta})$. However in doing this we will also

necessarily *decrease* the intelligence at points outside that region. So intelligence has a non-local character, a character that prevents us from directly modifying it to ensure that it is simultaneously high for any and all ζ .

A second, more general problem is that without specifying the details of a microlearner, it can be extremely difficult to predict which of two private utilities the microlearner will be better able to learn. Indeed, even *with* the details, making that prediction can be nearly impossible. So it can be extremely difficult to determine what private utility intelligence values will accrue to various choices of those private utilities. In other words, macrolearning that involves modifying the private utilities to try to directly increase intelligence with respect to those utilities can be quite difficult.

Fortunately, we can circumvent many of these difficulties by using a proxy for (private utility) intelligence. Although we expect its value usually to be correlated with that of intelligence in practice, this proxy does not share intelligence’s non-local nature. In addition, the proxy does not depend heavily on the details of the microlearning algorithms used, *i.e.*, it is fairly independent of those aspects of C . Intuitively, this proxy can be viewed as a “salient characteristic” for intelligence.

We motivate this proxy by considering having $g_\eta = G$ for all η . If we try to actually use these $\{g_\eta\}$ as the microlearners’ private utilities, particularly if the COIN is large, we will invariably encounter a very bad signal-to-noise problem. For this choice of utilities, the effects of the actions taken by node η on its utility may be “swamped” and effectively invisible, since there are so many other processes going into determining G ’s value. This makes it hard for η to discern the echo of its actions and learn how to improve its private utility. It also means that η will find it difficult to decide how best to act once learning has completed, since so much of what’s important to η is decided by processes outside of η ’s immediate purview. In such a scenario, there is nothing that η ’s microlearner can do to reliably achieve high intelligence.¹⁵

In addition to this “observation-driven” signal/noise problem, there is an “action-driven” one. For reasons discussed in Observation 7.1 above, we can define a distribution $d\mu(\underline{\zeta}'_{\eta,0})$ reflecting what η does/doesn’t know concerning the actual state of the outside world $\hat{\eta}$ at time 0. If the node η chooses its actions in a Bayes-optimal manner, then $\underline{\zeta}_{\eta,0;act} = \operatorname{argmax}_z [\int d\mu(\underline{\zeta}'_{\eta,0}) U(\underline{\zeta}_{t<0} \bullet C(z \bullet \underline{\zeta}_{\eta,0;act} \bullet \underline{\zeta}'_{\eta,0}))]$, where z runs over the allowed action components of η at time 0. Since this will differ from $\operatorname{argmax}_z U(\underline{\zeta}_{t<0} \bullet C(z \bullet \underline{\zeta}_{\eta,0;act} \bullet \underline{\zeta}_{\eta,0}))$ in general, this Bayes-optimal node’s intelligence will be less than 1 for the particular $\underline{\zeta}$ at hand, in general. Moreover, the less U ’s ultimate value (after the application of C , etc.) depends on $\underline{\zeta}_{\eta,0}$, the smaller the difference in these two argmax -based z ’s, and therefore the higher the intelligence of η , in general.¹⁶

¹⁵This “signal-to-noise” problem is actually endemic to reinforcement learning as a whole, even sometimes occurring when one has just a single reinforcement learner, and only a few random variables jointly determining the value of the rewards [287].

¹⁶In practice, due to computational limitations if nothing else, the node won’t be exactly Bayes-optimal. But incorporating such a suboptimality doesn’t affect the thrust of this argument that we want

We would like a measure of U that captures these effects, but without depending on function maximization or any other detailed aspects of how the node determines its actions. One natural way to do this is via the **(utility) learnability**: Given a measure $d\mu(\underline{\zeta}'_0)$ restricted to a manifold C , the ($t = 0$) utility learnability of a utility U for a node η at $\underline{\zeta}$ is:

$$\Lambda_{\eta,U}(\underline{\zeta}) \equiv \frac{\int d\mu(\underline{\zeta}'_0) |U(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}'_{\eta,0})) - U(\underline{\zeta})|}{\int d\mu(\underline{\zeta}'_0) |U(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}'_{\eta,0} \bullet \underline{\zeta}_{\eta,0})) - U(\underline{\zeta})|}. \quad (2)$$

Intelligence learnability is defined the same way, with $U(\cdot)$ replaced by $\epsilon_{\eta,U}(\cdot)$. Note that any affine transformation of U has no effect on either the utility learnability $\Lambda_{\eta,U}(\underline{\zeta})$ or the associated intelligence learnability, $\Lambda_{\eta,\epsilon_{\eta,U}}(\underline{\zeta})$.

The integrand in the numerator of the definition of learnability reflects how much of the change in U that results from replacing $\underline{\zeta}_0$ with $\underline{\zeta}'_0$ is due to the change in η 's $t = 0$ state (the “signal”). The denominator reflects how much of the change in U that results from replacing $\underline{\zeta}$ with $\underline{\zeta}'$ is due to the change in the $t = 0$ states of nodes other than η (the “noise”). So learnability quantifies how easy it is for the microlearner to discern the “echo” of its behavior in the utility function U . Our presumption is that the microlearning algorithm will achieve higher intelligence if provided with a more learnable private utility.

Intuitively, the (utility) **differential learnability** of U at a point $\underline{\zeta}$ is the learnability with $d\mu$ restricted to an infinitesimal ball about $\underline{\zeta}$. We formalize it as the following ratio of magnitudes of a pair of gradients, one involving η , and one involving $\hat{\eta}$:

$$\lambda_{\eta,U}(\underline{\zeta}) \equiv \frac{\|\partial_{\underline{\zeta}_{\eta,0}} U(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0}))\|}{\|\partial_{\underline{\zeta}_{\hat{\eta},0}} U(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\hat{\eta},0}))\|}. \quad (3)$$

Note that a particular value of differential utility learnability, by itself, has no significance. Simply rescaling the units of $\underline{\zeta}_{\eta,0}$ will change that value. Rather what is important is the ratio of differential learnabilities, at the same $\underline{\zeta}$, for different U 's. Such a ratio quantifies the relative preferability of those U 's.

One nice feature of differential learnability is that unlike learnability, it does not depend on choice of some measure $d\mu(\cdot)$. This independence can lead to trouble if one is not careful however, and in particular if one uses learnability for purposes other than choosing between utility functions. For example, in some situations, the COIN designer will have the option of enlarging the set of variables from the rest of the COIN that are “input” to some node η at $t = 0$ and that therefore can be used by η to decide what action to take. Intuitively, doing so will not affect the RL “signal” for η 's microlearner (the magnitude of the potential “echo” of η 's actions are not modified by changing some aspect of how it chooses among those actions). However it *will* reduce the “noise”, in that η 's microlearner now knows more about the state of the rest of the system.

U 's ultimate value to not depend on $\underline{\zeta}_{\eta,0}$.

In the full integral version of learnability, this effect can be captured by having the support of $d\mu(\cdot)$ restricted to reflect the fact that the extra inputs to η at $t = 0$ are correlated with the $t = 0$ state of the external system. In differential learnability however this is not possible, precisely because no measure $d\mu(\cdot)$ occurs in its definition. So we must capture the reduction in noise in some other fashion.¹⁷

Alternatively, if the extra variables are being input to η for all $t \geq 0$, not just at $t = 0$, and if η “pays attention” to those variables for all $t \geq 0$, then by incorporating those changes into our system C itself has changed, $\forall t \geq 0$. Hypothesize that at those t the node η is capable of modifying its actions to “compensate” for what (due to our augmentation of η ’s inputs) η now knows to be going on outside of it. Under this hypothesis, those changes in those external events will have less of an effect on the ultimate value of g_η than they would if we had not made our modification. In this situation, the noise term has been reduced, so that the differential learnability properly captures the effect of η ’s having more inputs.

Another potential danger to bear in mind concerning differential learnability is that it is usually best to consider its average over a region, in particular over points with less than maximal intelligence. It is really designed for such points; in fact, at the intelligence-maximizing $\underline{\zeta}$, $\lambda_{\eta,U}(\underline{\zeta}) = 0$.

Whether in its differential form or not, and whether referring to utilities or intelligence, learnability is not meant to capture all factors that will affect how high an intelligence value a particular microlearner will achieve. Such an all-inclusive definition is not possible, if for no other reason the fact that there are many such factors that are idiosyncratic to the particular microlearner used. Beyond this though, certain more general factors that affect most popular learning algorithms, like the curse of dimensionality, are also not (explicitly) designed into learnability. Learnability is not meant to provide a full characterization of performance — that is what intelligence is designed to do. Rather (relative) learnability is only meant to provide a *guide* for how to improve performance.

A system that has infinite (differential, intelligence) learnability for all its personal utilities is said to be “perfectly” (differential, intelligence) learnable. It is straightforward to prove that a system is perfectly learnable $\forall \underline{\zeta} \in C$ iff $\forall \eta, \forall \underline{\zeta} \in C, g_\eta(\underline{\zeta})$ can

¹⁷One way to capture this noise reduction is to replace the noise term $\partial_{\underline{\zeta}_{\eta,0}} U(\underline{\zeta}_{\eta,t < 0} \bullet C(\underline{\zeta}_0))$ occurring in the definition of differential learnability with something more nuanced. For example, one may wish to replace it with the maximum of the dot product of $\underline{u} \equiv \partial_{\underline{\zeta}} U(\underline{\zeta})$ with any \underline{Z} vector \underline{v} , subject not only to the restrictions that $\|\underline{v}\| = 1$ and $\underline{v}_{\eta,0} = \mathbf{0}$, but also subject to the restriction that \underline{v} must lie in the tangent plane of C at $\underline{\zeta}$. The first two restrictions, in concert with the extra restriction that $\underline{v}_{\eta,t < 0} = \mathbf{0}$, give the original definition of the noise term. If they are instead joined with the third, new restriction, they will enforce any applicable coupling between the state of η at time 0 and the rest of the system at time 0. Solving with Lagrange multipliers, we get $\underline{v} \propto \underline{u} - \lambda_2 \underline{\alpha} - \lambda_3 \underline{\beta}$, where $\underline{\alpha}$ is the normal to C at $\underline{\zeta}$, $\underline{\beta}_{\eta',t'} \equiv \delta_{\eta',\eta} \delta_{t',t}$, and $\lambda_2 = \frac{\underline{\alpha} \cdot \underline{u} - (\underline{\beta} \cdot \underline{u})(\underline{\alpha} \cdot \underline{\beta})}{1 - (\underline{\alpha} \cdot \underline{\beta})^2}$ while $\lambda_3 = \frac{\underline{\beta} \cdot \underline{u} - (\underline{\alpha} \cdot \underline{u})(\underline{\alpha} \cdot \underline{\beta})}{1 - (\underline{\alpha} \cdot \underline{\beta})^2}$. As a practical matter though, it is often simplest to assume that the $\underline{\zeta}_{\eta,0}$ can vary arbitrarily, independent of $\underline{\zeta}_{\eta,0}$, so that the noise term takes the form in Eq. 3.

be written as $\psi_\eta(\underline{\zeta}_{\eta,0})$ for some function $\psi_\eta(\cdot)$. (See the discussion below on the general condition for a system’s being perfectly factored.)

4.3 A descriptive framework for COINs

With these definitions in hand, we can now present (a portion of) one descriptive framework for COINs. In this subsection, after discussing salient characteristics in general, we present some theorems concerning the relationship between personal utilities and the salient characteristic we choose to concentrate on. We then discuss how to use these theorems to induce that salient characteristic in a COIN.

4.3.1 Candidate salient characteristics of a COIN

The starting point with a descriptive framework is the identification of “salient characteristics of a COIN which one strongly expects to be associated with its having large world utility”. In this chapter we will focus on salient characteristics that concern the relationship between personal and world utilities. These characteristics are formalizations of the intuition that we want COINs in which the competent greedy pursuit of their private utilities by the microlearners results in large world utility, without any bottlenecks, TOC, “frustration” (in the spin glass sense) or the like.

One natural candidate for such a characteristic, related to Pareto optimality [100, 101], is **weak triviality**. It is defined by considering any two worldlines $\underline{\zeta}$ and $\underline{\zeta}'$ both of which are consistent with the system’s dynamics (*i.e.*, both of which lie on C), where for every node η , $g_\eta(\underline{\zeta}) \geq g_\eta(\underline{\zeta}')$.¹⁸ If for any such pair of worldlines where one “Pareto dominates” the other it is necessarily true that $G(\underline{\zeta}) \geq G(\underline{\zeta}')$, we say that the system is weakly trivial. We might expect that systems that are weakly trivial for the microlearners’ private utilities are configured correctly for inducing large world utility. After all, for such systems, if the microlearners collectively change $\underline{\zeta}$ in a way that ends up helping all of them, then necessarily the world utility also rises. More formally, for a weakly trivial system, the maxima of G are Pareto-optimal points for the personal utilities (although the reverse need not be true).

As it turns out though, weakly trivial systems can readily evolve to a world utility *minimum*, one that often involves TOC. To see this, consider automobile traffic in the absence of any traffic control system. Let each node be a different driver, and say their private utilities are how quickly they each individually get to their destination. Identify world utility as the sum of private utilities. Then by simple additivity, for all $\underline{\zeta}$ and $\underline{\zeta}'$, whether they lie on C or not, if $g_\eta(\underline{\zeta}) \geq g_\eta(\underline{\zeta}') \quad \forall \eta$ it follows that $G(\underline{\zeta}) \geq G(\underline{\zeta}')$; the system is weakly trivial. However as any driver on a rush-hour freeway with no carpool lanes or metering lights can attest, every driver’s pursuing their own goal definitely does

¹⁸An obvious variant is to restrict $\underline{\zeta}'_{t < 0} = \underline{\zeta}_{t < 0}$, and require only that both of the “partial vectors” $\underline{\zeta}'_{t \geq 0}$ and $\underline{\zeta}_{t \geq 0}$ obey the relevant dynamical laws, and therefore lie in $C_{t \geq 0}$.

not result in acceptable throughput for the system as a whole; modifications to private utility functions (like fines for violating carpool lanes or metering lights) would result in far better global behavior. A system's being weakly trivial provides no assurances regarding world utility.

This does not mean weak triviality is never of use. For example, say that for a set of weakly trivial personal utilities each agent can guarantee that *regardless of what the other agents do*, its utility is above a certain level. Assume further that, being risk-averse, each agent chooses an action with such a guarantee. Say it is also true that the agents are provided with a relatively large set of candidate guaranteed values of their utilities. Under these circumstances, the system's being weakly trivial provides some assurances that world utility is not too low. Moreover, if the overhead in enforcing such a future-guaranteeing scheme is small, and having a sizable set of guaranteed candidate actions provided to each of the agents does not require an excessively centralized infrastructure, we can actually employ this kind of scheme in practice. Indeed, in the extreme case, one can imagine that every agent is guaranteed exactly what its utility would be for every one of its candidate actions. (See the discussion on General Equilibrium in the Background Section above.) In this situation, Nash equilibria and Pareto optimal points are identical, which due to weak triviality means that the point maximizing G is a Nash equilibrium. However in any less extreme situation, the system may not achieve a value of world utility that is close to optimal. This is because even for weakly trivial systems a Pareto optimal point may have poor world utility, in general.

Situations where one has guarantees of lower bounds on one's utility are not too common, but they do arise. One important example is a round of trades in a computational market (see the Background Section above). In that scenario, there is an agent-indexed set of functions $\{f_\eta(\mathbf{z} \in \underline{\mathbf{Z}}^{(0)})\}$ and the personal utility of each agent $\eta \in \{1, 2, \dots\}$ is given by $f_\eta(\underline{\zeta}_{\eta, t^*})$, where t^* is the end of the round of trades. There is also a function $F(\mathbf{z} \in \underline{\mathbf{Z}}^{(0)}) \equiv F(f_1(\mathbf{z}), f_2(\mathbf{z}), \dots)$ that is a monotonically increasing function of its arguments, and world utility G is given by $F(\underline{\zeta}_{t^*}) = F(f_1(\underline{\zeta}_{t^*}), f_2(\underline{\zeta}_{t^*}), \dots)$. So the system is weakly trivial. In turn, each $f_\eta(\mathbf{z})$ is determined solely by the "allotment of goods" possessed by η , as specified in the appropriate components of z_η . To be able to remove uncertainty about its future value of f_η in this kind of system, in determining its trading actions each agent η must employ some scheme like inter-agent contracts. This is because without such a scheme, no agent can be assured that if it agrees to a proposed trade with another agent that the full proposed transaction of that trade actually occurs. Given such a scheme, if in each trade round t each agent η myopically only considers those trades that are assured of increasing the corresponding value of f_η , then we are guaranteed that the value of the world utility is not less than the initial value $F(\underline{\zeta}_0)$.

The problem with using weak triviality as a general salient characteristic is precisely the fact that the individual microlearners *are* greedy. In a COIN, there is no system-wide incentive to replace $\underline{\zeta}$ with a different worldline that would improve everybody's private utility, as in the definition of weak triviality. Rather the incentives apply to each

microlearner individually and motivate the learners to behave in a way that may well hurt some of them. So weak triviality is, upon examination, a poor choice for the salient characteristic of a COIN.

One alternative to weak triviality follows from consideration of the stricture that we must ‘expect’ a salient characteristic to be coupled to large world utility in a running real-world COIN. What can we reasonably expect about a running real-world COIN? We cannot assume that all the private utilities will have large values — witness the traffic example. But we *can* assume that if the microlearners are well-designed, each of them will be doing close to as well it can *given the behavior of the other nodes*. In other words, within broad limits we can assume that the system is more likely to be in $\underline{\zeta}$ than $\underline{\zeta}'$ if for all η , $\epsilon_{\eta, g_\eta}(\underline{\zeta}) \geq \epsilon_{\eta, g_\eta}(\underline{\zeta}')$. We define a system to be **coordinated** iff for any such $\underline{\zeta}$ and $\underline{\zeta}'$ lying on C , $G(\underline{\zeta}) \geq G(\underline{\zeta}')$. (Again, an obvious variant is to restrict $\underline{\zeta}'_{t<0} = \underline{\zeta}_{t<0}$, and require only that both $\underline{\zeta}_{t \geq 0}$ and $\underline{\zeta}'_{t \geq 0}$ lie in $C_{t \geq 0}$.) Traffic systems are *not* coordinated, in general. This is evident from the simple fact that if all drivers acted as though there were metering lights when in fact there weren’t any, they would each be behaving with lower intelligence given the actions of the other drivers (each driver would benefit greatly by changing its behavior by no longer pretending there were metering lights, etc.). But nonetheless, world utility would be higher.

4.3.2 The Salient Characteristic of Factoredness

Like weak triviality, coordination is intimately related to the economics concept of Pareto optimality. Unfortunately, there is not room in this chapter to present the mathematics associated with coordination and its variants. We will instead discuss a third candidate salient characteristic of COINs, one which like coordination (and unlike weak triviality) we can reasonably expect to be associated with large world utility. This alternative fixes weak triviality not by replacing the personal utilities $\{g_\eta\}$ with the intelligences $\{\epsilon_{\eta, g_\eta}\}$ as coordination does, but rather by only considering worldlines whose difference at time 0 involves a single node. This results in this alternative’s being related to Nash equilibria rather than Pareto optimality.

Say that our COIN’s worldline is $\underline{\zeta}$. Let $\underline{\zeta}'$ be any other worldline where $\underline{\zeta}_{t<0} = \underline{\zeta}'_{t<0}$, and where $\underline{\zeta}'_{t \geq 0} \in C_{t \geq 0}$. Now restrict attention to those $\underline{\zeta}'$ where at $t = 0$ $\underline{\zeta}$ and $\underline{\zeta}'$ differ only for node η . If for all such $\underline{\zeta}'$

$$\text{sgn}[g_\eta(\underline{\zeta}) - g_\eta(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}'_0))] = \text{sgn}[G(\underline{\zeta}) - G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}'_0))] , \quad (4)$$

and if this is true for all nodes η , then we say that the COIN is **factored** for all those utilities $\{g_\eta\}$ (at $\underline{\zeta}$, with respect to time 0 and the utility G).

For a factored system, for any node η , *given the rest of the system*, if the node’s state at $t = 0$ changes in a way that improves that node’s utility over the rest of time, then it necessarily also improves world utility. Colloquially, for a system that is factored for a particular microlearner’s private utility, if that learner does something that improves

that personal utility, then everything else being equal, it has also done something that improves world utility. Of two potential microlearners for controlling node η (*i.e.*, two potential $\underline{\zeta}_\eta$) whose behavior until $t = 0$ is identical but which differ there, the microlearner that is smarter with respect to g will always result in a larger g , by definition of intelligence. Accordingly, for a factored system, the smarter microlearner is also the one that results in better G . So as long as we have deployed a sufficiently smart microlearner on η , we have assured a good G (given the rest of the system). Formally, this is expressed in the fact [283] that for a factored system, for all nodes η ,

$$\epsilon_{\eta, g_\eta}(\underline{\zeta}) = \epsilon_{\eta, G}(\underline{\zeta}) . \quad (5)$$

One can also prove that Nash equilibria of a factored system are local maxima of world utility. Note that in keeping with our behaviorist perspective, nothing in the definition of factored requires the existence of private utilities. Indeed, it may well be that a system having private utilities $\{U_\eta\}$ is factored, but for personal utilities $\{g_\eta\}$ that differ from the $\{U_\eta\}$.

A system's being factored does *not* mean that a change to $\underline{\zeta}_{\eta,0}$ that improves $g_\eta(\underline{\zeta})$ cannot also hurt $g_{\eta'}(\underline{\zeta})$ for some $\eta' \neq \eta$. Intuitively, for a factored system, the side effects on the rest of the system of η 's increasing its own utility do not end up decreasing world utility — but can have arbitrarily adverse effects on other private utilities. (In the language of economics, no stipulation is made that η 's “costs are endogenized.”) For factored systems, the separate microlearners successfully pursuing their separate goals do not frustrate each other *as far as world utility is concerned*.

In addition, if $g_{\eta,t'}$ is factored with respect to G , then a change to $\underline{\zeta}_{\eta,t'}$ that improves $g_{\eta,t'}(\underline{\zeta}_{t < t'}, C(\underline{\zeta}_{t'}))$ improves $G(\underline{\zeta}_{t < t'}, C(\underline{\zeta}_{t'}))$. But it may *hurt* some $g_{\eta,t'' \neq t'}(\underline{\zeta}_{t < t'}, C(\underline{\zeta}_{t'}))$ and/or $\epsilon_{(\eta,t''), g_{\eta,t''}}(\underline{\zeta}_{t < t'}, C(\underline{\zeta}_{t'}))$. (This is even true for a discounted sum of rewards personal utility, so long as $t'' > t'$.) An example of this would be an economic system cast as a single individual, η , together with an environment node, where G is a steeply discounted sum of rewards η receives over his/her lifetime, $t'' > t'$, and $\forall t, g_{\eta,t}(\underline{\zeta}) = G(\text{CL}_{t < t'}(\underline{\zeta}))$. For such a situation, it may be appropriate for η to live extravagantly at the time t' , and “pay for it” later.

As an instructive example of the ramifications of Eq. 5, say node η is a conventional computer. We want $\epsilon_{\eta, G}(\underline{\zeta})$ to be as high as possible, *i.e.*, given the state of the rest of the system at time 0, we want computer η 's state then to be the best possible, as far as the resultant value of G is concerned. Now a computer's “state” consists of the values of all its bits, including its code segment, *i.e.*, including the program it is running. So for a factored personal utility g_η , if the program running on the computer is better than most others as far as g_η is concerned, then it is also better than most other programs as far as G is concerned.

Our task as COIN designers engaged in COIN initialization or macrolearning is to find such a program and such an associated g_η . One way to approach this task is to restrict

attention to programs that consist of RL algorithms with private utility specified in the bits $\{b_i\}$ of η . This reduces the task to one of finding a private utility $\{b_i\}$ (and thereby fully specifying $\zeta_{\eta,0}$) such that our RL algorithm working with that private utility has high ϵ_{η,g_η} , i.e., such that that algorithm outperforms most other programs as far as the personal utility g_η is concerned.

Perhaps the simplest way to address this reduced task is to exploit the fact that for a good enough RL algorithm $\epsilon_{\eta,\{b_i\}}$ will be large, and therefore adopt such an RL algorithm and fix the private utility to equal g_η . In this way we further reduce the original task, which was to search over all personal utilities g_η and all programs R to find a pair such that both g_η is factored with respect to G and there are relatively few programs that outperform R , as far as g_η . The task is now instead to search over all private utilities $\{b_i\}$ such that both $\{b_i\}$ is factored with respect to G and such that there are few programs (*of any sort*, RL-based or not) that outperform our RL algorithm working on $\{b_i\}$, as far as that self-same private utility is concerned. The crucial assumption being leveraged in this approach is that our RL algorithm is “good enough”, and the reason we want learnable $\{b_i\}$ is to help effect this assumption.

In general though, we can’t have both perfect learnability and perfect factoredness. As an example, say that $\forall t, \underline{Z}_{\eta,t} = \underline{Z}_{\eta,t} = \mathcal{R}$, and that the dynamics is the identity operator: $\forall t, C(\underline{\zeta}_0)_t = \underline{\zeta}_0$. Then if $G(\underline{\zeta}_0) = \underline{\zeta}_{\eta,0} \cdot \underline{\zeta}_{\eta,0}$ and the system is perfectly learnable, it is not perfectly factored. This is because perfect learnability requires that $\forall \underline{\zeta} \in C, g_\eta(\underline{\zeta}) = \psi_\eta(\underline{\zeta}_{\eta,0})$ for some function $\psi_\eta(\cdot)$. However any change to $\underline{\zeta}_{\eta,0}$ that improves such a g_η will either help or *hurt* $G(\underline{\zeta})$, depending on the sign of $\underline{\zeta}_{\eta,0}$. For the “wrong” sign of $\underline{\zeta}_{\eta,0}$, this means the system is actually “anti-factored”. Due to such incompatibility between perfect factoredness and perfect learnability, we must usually be content with having high degree of factoredness and high learnability. In such situations, the emphasis of the macrolearning process should be more and more on having high degree of factoredness as we get closer and closer to a Nash equilibrium. This way the system won’t relax to an incorrect local maximum.

In practice of course, a COIN will often not be perfectly factored. Nor in practice are we always interested only in whether the system is factored at one particular point (rather than across a region say). These issues are discussed in [283], where in particular a formal definition of the **degree of factoredness** of a system is presented.

If a system is factored for utilities $\{g_\eta\}$, then it is also factored for any utilities $\{g'_\eta\}$ where for each η g'_η is a monotonically increasing function of g_η . More generally, the following result characterizes the set of all factored personal utilities:

Theorem 1: A system is factored at all $\underline{\zeta} \in C$ iff for all those $\underline{\zeta}$, $\forall \eta$, we can write

$$g_\eta(\underline{\zeta}) = \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta})) \quad (6)$$

for some function $\Phi_\eta(\cdot, \cdot, \cdot)$ such that $\partial_G \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G) > 0$ for all $\underline{\zeta} \in C$ and associated G values. (The form of the $\{g_\eta\}$ off of C is arbitrary.)

Proof: For fixed $\underline{\zeta}_{\eta,0}$ and $\underline{\zeta}_{t<0}$, any change to $\underline{\zeta}_{\eta,0}$ which keeps $\underline{\zeta}_{t\geq 0}$ on C and which at the same time increases $G(\underline{\zeta}) = G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}))$ must increase $\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta}))$, due to the restriction on $\partial_G \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G)$. This establishes the backwards direction of the proof.

For the forward direction, write $g_\eta(\underline{\zeta}) = g_\eta(\underline{\zeta}, G(\underline{\zeta})) = g_\eta(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}), G(\underline{\zeta})) \quad \forall \underline{\zeta} \in C$. Define this formulation of g_η as $\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta}))$, which we can re-express as $\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}, G(\underline{\zeta}))$. Now since the system is factored, $\forall \underline{\zeta} \in C, \forall \underline{\zeta}'_{t\geq 0} \in C_{t\geq 0}$,

$$\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}, G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}))) = \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0} \bullet \underline{\zeta}'_{\eta,0}, G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}'_{\eta,0})))$$

$$\Longleftrightarrow$$

$$G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0})) = G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}'_{\eta,0})) .$$

So consider any situation where the system is factored, and the values of G , $\underline{\zeta}_{t<0}$, and $\underline{\zeta}_{\eta,0}$ are specified. Then we can find *any* $\underline{\zeta}_{\eta,0}$ consistent with those values (*i.e.*, such that our provided value of G equals $G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}))$), evaluate the resulting value of $\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0} \bullet \underline{\zeta}_{\eta,0}, G)$, and know that we would have gotten the same value if we had found a different consistent $\underline{\zeta}_{\eta,0}$. This is true for all $\underline{\zeta} \in C$. Therefore the mapping $(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G) \rightarrow \Phi_\eta$ is single-valued, and we can write $\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta}))$. **QED.**

By Thm. 1, we can ensure that the system is factored without any concern for C , by having each $g_\eta(\underline{\zeta}) = \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta})) \quad \forall \underline{\zeta} \in \underline{\mathbf{Z}}$. Alternatively, by only requiring that $\forall \underline{\zeta} \in C$ does $g_\eta(\underline{\zeta}) = \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta}))$ (*i.e.*, does $g_\eta(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0})) = \Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_{\eta,0})))$), we can access a broader class of factored utilities, a class that *does* depend on C . Loosely speaking, for those utilities, we only need the projection of $\partial_{\underline{\zeta}_{t\geq 0}} G(\underline{\zeta})$ onto $C_{\eta,0}$ to be parallel to the projection of $\partial_{\underline{\zeta}_{t\geq 0}} g_\eta(\underline{\zeta})$ onto $C_{\eta,0}$. Given G and C , there are infinitely many $\partial_{\underline{\zeta}_{t\geq 0}} g_\eta(\underline{\zeta})$ having this projection (the set of such $\partial_{\underline{\zeta}_{t\geq 0}} g_\eta(\underline{\zeta})$ form a linear subspace of $\underline{\mathbf{Z}}$). The partial differential equations expressing the precise relationship are discussed in [283].

As an example of the foregoing, consider a ‘team game’ (also known as an ‘exact potential game’ [83, 193]) in which $g_\eta = G$ for all η . Such COINs are factored, trivially, regardless of C ; if g_η rises, then G must as well, by definition. (Alternatively, to confirm that team games are factored just take $\Phi_\eta(\underline{\zeta}_{t<0}, \underline{\zeta}_{\eta,0}, G) = G \quad \forall \eta$ in Thm. 1.) On the other hand, as discussed below, COINs with ‘wonderful life’ personal utilities are also factored, but the definition of such utilities depends on C .

4.3.3 Wonderful life utility

Due to their often having poor learnability and requiring centralized communication (among other infelicities), in practice team game utilities often are poor choices for

personal utilities. Accordingly, it is often preferable to use some other set of factored utilities. To present an important example, first define the $(t = 0)$ **effect set** of node η at $\underline{\zeta}$, $C_\eta^{eff}(\underline{\zeta})$, as the set of all components $\underline{\zeta}_{\eta',t}$ for which $\partial_{\underline{\zeta}_{\eta,0}}(C(\underline{\zeta}_0))_{\eta',t} \neq \vec{0}$. Define the effect set C_η^{eff} with no specification of $\underline{\zeta}$ as $\cup_{\underline{\zeta} \in C} C_\eta^{eff}(\underline{\zeta})$. (We take this latter definition to be the default meaning of “effect set”.) We will also find it useful to define C_η^{eff} as the set of components of the space \underline{Z} that are not in C_η^{eff} .

Intuitively, η ’s effect set is the set of all components $\underline{\zeta}_{\eta',t}$ which would be affected by a change in the state of node η at time 0. (They may or may not be affected by changes in the $t = 0$ states of the other nodes.) Note that the effect sets of different nodes may overlap. The extension of the definition of effect sets for times other than 0 is immediate. So is the modification to have effect sets only consist of those components $\underline{\zeta}_{\eta,t;i}$ that vary with the state of node η at time 0, rather than consist of the full vectors $\underline{\zeta}_{\eta,t}$ possessing such a component. These modifications will be skipped here, to minimize the number of variables we must keep track of.

Next for any set σ of components (η', t) , define $CL_\sigma(\underline{\zeta})$ as the “virtual” vector formed by clamping the σ -components of $\underline{\zeta}$ to an arbitrary fixed value. (In this paper, we take that fixed value to be $\vec{0}$ for all components listed in σ .) Consider in particular a **Wonderful Life** set σ . The value of the **wonderful life utility** (WLU for short) for σ at $\underline{\zeta}$ is defined as:

$$WLU_\sigma(\underline{\zeta}) \equiv G(\underline{\zeta}) - G(CL_\sigma(\underline{\zeta})) . \quad (7)$$

In particular, the WLU for the effect set of node η is $G(\underline{\zeta}) - G(CL_{C_\eta^{eff}}(\underline{\zeta}))$, which for $\underline{\zeta} \in C$ can be written as $G(\underline{\zeta}_{t < 0} \bullet C(\underline{\zeta}_0)) - G(CL_{C_\eta^{eff}}(\underline{\zeta}_{t < 0} \bullet C(\underline{\zeta}_0)))$.

We can view η ’s effect set WLU as analogous to the change in world utility that would have arisen if node η “had never existed”. (Hence the name of this utility - cf. the Frank Capra movie.) Note however, that CL is a purely “fictional”, counter-factual operation, in the sense that it produces a new $\underline{\zeta}$ without taking into account the system’s dynamics. Indeed, no assumption is even being made that $CL_\sigma(\underline{\zeta})$ is consistent with the dynamics of the system. The sequence of states the node η is clamped to in the definition of the WLU need not be consistent with the dynamical laws embodied in C .

This dynamics-independence is a crucial strength of the WLU. It means that to evaluate the WLU we do *not* try to infer how the system would have evolved if node η ’s state were set to 0 at time 0 and the system evolved from there. So long as we know $\underline{\zeta}$ extending over all time, and so long as we know G , we know the value of WLU. This is true even if we know nothing of the dynamics of the system.

An important example is effect set wonderful life utilities when the set of all nodes is partitioned into ‘subworlds’ in such a way that all nodes in the same subworld ω share substantially the same effect set. In such a situation, all nodes in the same subworld ω will have essentially the same personal utilities, exactly as they would if they used team game utilities with a “world” given by ω . When all such nodes have large intelligence values, this sharing of the personal utility will mean that all nodes in the same subworld

are acting in a coordinated fashion, loosely speaking.

The importance of the WLU arises from the following results:

Theorem 2: i) A system is factored at all $\underline{\zeta} \in C$ iff for all those $\underline{\zeta}$, $\forall \eta$, we can write

$$g_\eta(\underline{\zeta}) = \hat{\Phi}_\eta(\underline{\zeta}_{C_\eta^{eff}}, G(\underline{\zeta})) \quad (8)$$

for some function $\hat{\Phi}_\eta(\cdot, \cdot)$ such that $\partial_G \hat{\Phi}_\eta(\underline{\zeta}_{C_\eta^{eff}}, G) > 0$ for all $\underline{\zeta} \in C$ and associated G values. (The form of the $\{g_\eta\}$ off of C is arbitrary.)

ii) In particular, a COIN is factored for personal utilities set equal to the associated effect set wonderful life utilities.

Proof: To prove (i), first write $\underline{\zeta}_{C_\eta^{eff}} = \underline{\zeta}_{t < 0} \bullet \underline{\zeta}_{\eta, 0} \bullet \underline{\zeta}_{(C_\eta^{eff})_{t > 0}}$. For all $\underline{\zeta} \in C$, $\underline{\zeta}_{(C_\eta^{eff})_{t > 0}}$ is independent of $\underline{\zeta}_{\eta, 0}$, and so by definition of $C(\cdot)$ it is a single-valued function of $\underline{\zeta}_{\eta, 0}$ for such $\underline{\zeta}$. Therefore $\underline{\zeta}_{C_\eta^{eff}} = \underline{\zeta}_{t < 0} \bullet \underline{\zeta}_{\eta, 0} \bullet f(\underline{\zeta}_{\eta, 0})$ for some function $f(\cdot)$. Accordingly, by Thm. 1, for $\{g_\eta\}$ of the form stipulated in (i), the system is factored. Going the other way, if the system is factored, then by Thm. 1 it can be written as $\Phi_\eta(\underline{\zeta}_{t < 0}, \underline{\zeta}_{\eta, 0}, G(\underline{\zeta}))$. Since both $\underline{\zeta}_{t < 0}$ and $\underline{\zeta}_{\eta, 0} \notin C_\eta^{eff}$, we can rewrite this as $\Phi_\eta([\underline{\zeta}_{C_\eta^{eff}}]_{t < 0}, [\underline{\zeta}_{C_\eta^{eff}}]_{\eta, 0}, G(\underline{\zeta}))$. **QED.**

Part (ii) of the theorem follows immediately from part (i). For pedagogical value though, here we instead derive it directly. First, since $\text{CL}_{C_\eta^{eff}}(\underline{\zeta})$ is independent of $\underline{\zeta}_{\eta', t}$ for all $(\eta', t) \in C_\eta^{eff}$, so is the $\underline{\mathbf{Z}}$ vector $\text{CL}_{C_\eta^{eff}}(\underline{\zeta}_{t < 0} \bullet C(\underline{\zeta}_0))$, i.e., $\partial_{\underline{\zeta}_{\eta, 0}}[\text{CL}_{C_\eta^{eff}}(\underline{\zeta}_{t < 0} \bullet C(\underline{\zeta}_0))]_{\eta', t} = \vec{0} \quad \forall (\eta', t) \in C_\eta^{eff}$. This means that viewed as a $\underline{\zeta}_{t < 0}$ -parameterized function from C_0 to $\underline{\mathbf{Z}}$, $\text{CL}_{C_\eta^{eff}}(\underline{\zeta}_{t < 0} \bullet C(\cdot))$ is a single-valued function of the $\underline{\zeta}_{\eta, 0}$ components. Therefore $G(\text{CL}_{C_\eta^{eff}}(\underline{\zeta}_{t < 0} \bullet C(\underline{\zeta}_0)))$ can only depend on $\underline{\zeta}_{t < 0}$ and the non- η components of $\underline{\zeta}_0$. Accordingly, the WLU for C_η^{eff} is just G minus a term that is a function of $\underline{\zeta}_{t < 0}$ and $\underline{\zeta}_{\eta, 0}$. By choosing $\Phi_\eta(\cdot, \cdot, \cdot)$ in Thm. 1 to be that difference, we see that η 's effect set WLU is of the form necessary for the system to be factored. **QED.**

As a generalization of (ii), the system is factored if each node η 's personal utility is (a monotonically increasing function of) the WLU for a set σ_η that contains C_η^{eff} .

For conciseness, except where explicitly needed, for the remainder of this subsection we will suppress the argument " $\underline{\zeta}_{t < 0}$ ", taking it to be implicit. The next result concerning the practical importance of effect set WLU is the following:

Theorem 3: Let σ be a set containing C_η^{eff} . Then

$$\frac{\lambda_{\eta, WLU_\sigma}(\underline{\zeta})}{\lambda_{\eta, G}(\underline{\zeta})} = \frac{\|\partial_{\underline{\zeta}_{\eta, 0}} G(C(\underline{\zeta}_0))\|}{\|\partial_{\underline{\zeta}_{\eta, 0}} G(C(\underline{\zeta}_0)) - \partial_{\underline{\zeta}_{\eta, 0}} G(\text{CL}_\sigma(C(\underline{\zeta}_0)))\|}.$$

Proof: Writing it out,

$$\lambda_{\eta, WLU_\sigma}(\underline{\zeta}) = \frac{\|\partial_{\underline{\zeta}_{\eta,0}} G(C(\underline{\zeta}_0)) - \partial_{\underline{\zeta}_{\eta,0}} G(\text{CL}_\sigma(C(\underline{\zeta}_0)))\|}{\|\partial_{\underline{\zeta}_{\eta,0}} G(C(\underline{\zeta}_0)) - \partial_{\underline{\zeta}_{\eta,0}} G(\text{CL}_\sigma(C(\underline{\zeta}_0)))\|}.$$

The second term in the numerator equals 0, by definition of effect set. Dividing by the similar expression for $\lambda_{\eta, G}(\underline{\zeta})$ then gives the result claimed. **QED.**

So if we expect that ratio of magnitudes of gradients to be large, effect set WLU has much higher learnability than team game utility — while still being factored, like team game utility. As an example, consider the case where the COIN is a very large system, with η being only a relatively minor part of the system (*e.g.*, a large human economy with η being a “typical John Doe living in Peoria Illinois”). Often in such a system, for the vast majority of nodes $\eta' \neq \eta$, how G varies with $\underline{\zeta}_{\eta'}$, will be essentially independent of the value $\underline{\zeta}_{\eta,0}$. (For example, how GDP of the US economy varies with the actions of our John Doe from Peoria, Illinois will be independent of the state of some Jane Smith living in Los Angeles, California.) In such circumstances, Thm. 3 tells us that the effect set wonderful life utility for η will have a far larger learnability than does the world utility.

For any fixed σ , if we change the clamping operation (*i.e.*, change the choice of the “arbitrary fixed value” we clamp each component to), then we change the mapping $\underline{\zeta}_0 \rightarrow \text{CL}_\sigma(C(\underline{\zeta}_0))$, and therefore change the mapping $(\underline{\zeta}_{\eta,0}, \underline{\zeta}_{\eta,0}) \rightarrow G(\text{CL}_\sigma(C(\underline{\zeta}_0)))$. Accordingly, changing the clamping operation can affect the value of $\partial_{\underline{\zeta}_{\eta,0}} G(\text{CL}_\sigma(C(\underline{\zeta}_0)))$ evaluated at some point $\underline{\zeta}_0$. Therefore, by Thm. 3, changing the clamping operation can affect $\lambda_{\eta, WLU_\sigma}(\underline{\zeta})$. So properly speaking, for any choice of σ , if we are going to use WLU_σ , we should set the clamping operation so as to maximize learnability. For simplicity though, in this paper we will ignore this phenomenon, and simply set the clamping operation to the more or less “natural” choice of $\mathbf{0}$, as mentioned above.

Next consider the case where, for some node η , we can write $G(\underline{\zeta})$ as $G_1(\underline{\zeta}_{C_\eta^{eff}}) + G_2(\underline{\zeta}_{t < 0} \bullet \underline{\zeta}_{C_\eta^{eff}})$. Say it is also true that η ’s effect set is a small fraction of the set of all components. In this case it is often true that the values of $G(\cdot)$ are much larger than those of $G_1(\cdot)$, which means that partial derivatives of $G(\cdot)$ are much larger than those of $G_1(\cdot)$. In such situations the effect set WLU is far more learnable than the world utility, due to the following results:

Theorem 4: If for some node η there is a set σ containing C_η^{eff} , a function $G_1(\underline{\zeta}_\sigma \in \mathbf{Z}_\sigma)$, and a function $G_2(\underline{\zeta}_\sigma \in \mathbf{Z}_\sigma)$, such that $G(\underline{\zeta}) = G_1(\underline{\zeta}_\sigma) + G_2(\underline{\zeta}_\sigma)$, then

$$\frac{\lambda_{\eta, WLU_\sigma}(\underline{\zeta})}{\lambda_{\eta, G}(\underline{\zeta})} = \frac{\|\partial_{\underline{\zeta}_{\eta,0}} G(C(\underline{\zeta}_0))\|}{\|\partial_{\underline{\zeta}_{\eta,0}} G(\text{CL}_\sigma(C(\underline{\zeta}_0)))\|}.$$

Proof: For brevity, write G_1 and G_2 both as functions of full $\underline{\zeta} \in \mathbf{Z}$, just such func-

tions that are only allowed to depend on the components of $\underline{\zeta}$ that lie in σ and those components that do not lie in σ , respectively. Then the σ WLU for node η is just $g_\eta(\underline{\zeta}) = G_1(\underline{\zeta}) - G_1(\text{CL}_\sigma(\underline{\zeta}))$. Since in that second term we are clamping all the components of $\underline{\zeta}$ that $G_1(\cdot)$ cares about, for this personal utility $\partial_{\underline{\zeta}_0} g_\eta(C(\underline{\zeta}_0)) = \partial_{\underline{\zeta}_0} G_1(C(\underline{\zeta}_0))$. So in particular $\partial_{\underline{\zeta}_{\eta,0}} g_\eta(C(\underline{\zeta}_0)) = \partial_{\underline{\zeta}_{\eta,0}} G_1(C(\underline{\zeta}_0)) = \partial_{\underline{\zeta}_{\eta,0}} G(\text{CL}_\sigma(C(\underline{\zeta}_0)))$. Now by definition of effect set, $\partial_{\underline{\zeta}_{\eta,0}} G_2(\underline{\zeta}_{t<0} \bullet C(\underline{\zeta}_0)) = \vec{0}$, since $\hat{\sigma}$ does not contain C_η^{eff} . So $\partial_{\underline{\zeta}_{\eta,0}} G(C(\underline{\zeta}_0)) = \partial_{\underline{\zeta}_{\eta,0}} G_1(C(\underline{\zeta}_0)) = \partial_{\underline{\zeta}_{\eta,0}} g_\eta(C(\underline{\zeta}_0))$. **QED.**

The obvious extensions of Thm.'s 3 and 4 to effect sets with respect to times other than 0 can also be proven [283].

An important special case of Thm. 4 is the following:

Corollary 1: If for some node η we can write

$$\text{i) } G(\underline{\zeta}) = G_1(\underline{\zeta}_\sigma) + G_2([\underline{\zeta}_\sigma]_{t \geq 0}) + G_3(\underline{\zeta}_{t < 0})$$

for some set σ containing C_η^{eff} , and if

$$\text{ii) } \|\partial_{\underline{\zeta}_{\eta,0}} G(C(\underline{\zeta}_0))\| \gg \|\partial_{\underline{\zeta}_{\eta,0}} G_1([C(\underline{\zeta}_0)]_\sigma)\|,$$

then

$$\lambda_{\eta, WLU_\sigma}(\underline{\zeta}) \gg \lambda_{\eta, G}(\underline{\zeta}).$$

In practice, to assure that condition (i) of this corollary is met might require that σ be a proper superset of C_η^{eff} . Countervailingly, to assure that condition (ii) is met will usually force us to keep σ as small as possible.

One can often remove elements from an effect set and still have the results of this section hold. Most obviously, if $(\eta', t) \in C_\eta^{eff}$ but $\partial_{\underline{\zeta}_{\eta',t}} G(\underline{\zeta}) = \mathbf{0}$, we can remove (η', t) from C_η^{eff} without invalidating our results. More generally, if there is a set $\sigma' \in C_\eta^{eff}$ such that for each component $(\eta, 0; i)$ the chain rule term $\sum_{(\eta', t) \in \sigma'} [\partial_{\underline{\zeta}_{\eta',t}} G(\underline{\zeta})] \cdot [\partial_{\underline{\zeta}_{\eta,0;i}} [C(\underline{\zeta}_0)]_{\eta',t}] = 0$, then the effects on G of changes to $\underline{\zeta}_{\eta,0}$ that are “mediated” by the members of σ' cancel each other out. In this case we can usually remove the elements of σ' from C_η^{eff} with no ill effects.

4.3.4 Inducing our salient characteristic

Usually the mathematics of a descriptive framework — a formal investigation of the salient characteristics — will not provide theorems of the sort, “If you modify the COIN the following way at time t , the value of the world utility will increase.” Rather it provides theorems that relate a COIN’s salient characteristics with the general properties of the COIN’s entire history, and in particular with those properties embodied in C . In particular, the salient characteristic that we are concerned with in this chapter is that the system be highly intelligent for personal utilities for which it is factored, and

our mathematics concerns the relationship between factoredness, intelligence, personal utilities, effect sets, and the like.

More formally, the desideratum associated with our salient characteristic is that we want the COIN to be at a ζ for which there is some set of $\{g_\eta\}$ (not necessarily consisting of private utilities) such that (a) ζ is factored for the $\{g_\eta\}$, and (b) $\epsilon_{\eta,g_\eta}(\zeta)$ is large for all η . Now there are several ways one might try to induce the COIN to be at such a point. One approach is to have each algorithm controlling η explicitly try to “steer” the worldline towards such a point. In this approach η needn’t even have a private utility in the usual sense. (The overt “goal” of the algorithm controlling η involves finding a ζ with a good associated extremum over the class of all possible g_η , independent of any private utilities.) Now initialization of the COIN, *i.e.*, fixing of ζ_0 , involves setting the algorithm controlling η , in this case to the steering algorithm. Accordingly, in this approach to initialization, we fix ζ_0 to a point for which there is some special g_η such that both $C(\zeta_0)$ is factored for g_η , and $\epsilon_{\eta,g_\eta}(C(\zeta_0))$ is large. There is nothing peculiar about this. What is odd though is that in this approach we do not know what that “special” g_η is when we do that initialization; it’s to be determined, by the unfolding of the system.

In this chapter we concentrate on a different approach, which can involve either initialization or macrolearning. In this alternative we deploy the $\{g_\eta\}$ as the microlearners’ private utilities at some $t < 0$, in a process not captured in C , so as to induce a factored COIN that is as intelligent as possible. (It is with that “deploying of the $\{g_\eta\}$ ” that we are trying to induce our salient characteristic in the COIN.) Since in this approach we are using private utilities, we can replace intelligence with its surrogate, learnability. So our task is to choose $\{g_\eta\}$ which are as learnable as possible while still being factored.

Solving for such utilities can be expressed as solving a set of coupled partial differential equations. Those equations involve the tangent plane to the manifold C , a functional trading off (the differential versions of) degree of factoredness and learnability, and any communication constraints on the nodes we must respect. While there is not space in the current chapter to present those equations, we can note that they are highly dependent on the correlations among the components of $\zeta_{\eta,t}$. So in this approach, in COIN initialization we use some preliminary guesses as to those correlations to set the initial $\{g_\eta\}$. For example, the effect set of a node constitutes all components $\zeta_{\eta',t>0}$ that have non-zero correlation with $\zeta_{\eta,0}$. Furthermore, by Thm. 2 the system is factored for effect set WLU personal utilities. And by Coroll. 1, for small effect sets, the effect set WLU has much greater differential utility learnability than does G . Extending the reasoning behind this result to all ζ (or at least all likely ζ), we see that for this scenario, the descriptive framework advises us to use Wonderful Life private utilities based on (guesses for) the associated effect sets rather than the team game private utilities, $g_\eta = G \ \forall \eta$.

In macrolearning we must instead run-time estimate an approximate solution to our

partial differential equations, based on statistical inference.¹⁹ As an example, we might start with an initial guess as to η 's effect set, and set its private utility to the associated WLU. But then as we watch the system run and observe the correlations among the components of $\underline{\zeta}$, we might modify which components we think comprise η 's effect set, and modify η 's personal utility accordingly.

4.4 Illustrative Simulations of our Descriptive Framework

As implied above, often one can perform reasonable COIN initialization and/or macrolearning without writing down the partial differential equations governing our salient characteristic explicitly. Simply “hacking” one’s way to the goal of maximizing both degree of factoredness and intelligibility, for example by estimating effect sets, often results in dramatic improvement in performance. This is illustrated in the experiments recounted in the next two subsections.

4.4.1 COIN Initialization

Even if we don’t exactly know the effect set of each node η , often we will be able to make a reasonable guess about which components of $\underline{\zeta}$ comprise the “preponderance” of η 's effect set. We call such a set a **guessed effect set**. As an example, often the primary effects of changes to η 's state will be on the future state of η , with only relatively minor effects on the future states of other nodes. In such situations, we would expect to still get good results if we approximated the effect set WLU of each node η with a WLU based on the guessed effect set $\underline{\zeta}_{\eta, t \geq 0}$. In other words, we would expect to be able to replace $\text{WLU}_{C_{\eta}^{eff}}$ with $\text{WLU}_{\underline{\zeta}_{\eta, t \geq 0}}$ and still get good performance.

This phenomenon was borne out in the experiments recounted in [288] that used COIN initialization for distributed control of network packet routing. In a conventional approach to packet routing, each router runs what it believes (based on the information available to it) to be a shortest path algorithm (SPA), *i.e.*, each router sends its packets in the way that it surmises will get those packets to their destinations most quickly. Unlike with an approach based on our COIN framework, with SPA-based routing the routers have no concern for the possible deleterious side-effects of their routing decisions on the global performance (*e.g.*, they have no concern for whether they induce bottlenecks). We performed simulations in which we compared such a COIN-based routing system to an SPA-based system. For the COIN-based system G was global throughput and no macrolearning was used. The COIN initialization was to have each router’s private

¹⁹Recall that in the physical world, it is often useful to employ devices using algorithms that are based on probabilistic concepts, even though the underlying system is ultimately deterministic. (Indeed, theological Bayesians invoke a “degree of belief” interpretation of probability to *demand* such an approach — see [282] for a discussion of the legitimacy of this viewpoint.) Similarly, although we take the underlying system in a COIN to be deterministic, it is often useful to use microlearners or — as here — macrolearners that are based on probabilistic concepts.

utility be a WLU based on an associated guessed effect set generated *a priori*. In addition, the COIN-based system was realistic in that each router’s reinforcement algorithm had imperfect knowledge of the state of the system. On the other hand, the SPA was an idealized “best-possible” system, in which each router knew exactly what the shortest paths were at any given time. Despite the handicap that this disparity imposed on the COIN-based system, it achieved significantly better global throughput in our experiments than did the perfect-knowledge SPA-based system, and in particular, avoided the Braess’ Paradox that was built-in to some of those systems [265].

The experiments in [288] were primarily concerned with the application of packet-routing. To concentrate more precisely on the issue of COIN initialization, we ran subsequent experiments on variants of Arthur’s famous “El Farol bar problem” (see Section 3). To facilitate the analysis we modified Arthur’s original problem to be more general, and since we were not interested in directly comparing our results to those in the literature, we used a more conventional (and arguably “dumber”) machine learning algorithm than the ones investigated in [4, 47, 50, 64].

In this formulation of the bar problem [289], there are N agents, each of whom picks one of seven nights to attend a bar the following week, a process that is then repeated. In each week, each agent’s pick is determined by its predictions of the associated rewards it would receive. These predictions in turn are based solely upon the rewards received by the agent in preceding weeks. An agent’s “pick” at week t (*i.e.*, its node’s state at that week) is represented as a unary seven-dimensional vector. (See the discussion in the definitions subsection of our representing discrete variables as Euclidean variables.) So η ’s zeroing its state in some week, as in the $\text{CL}_{\underline{\zeta}_{\eta,t}}$ operation, essentially means it elects not to attend any night that week.

The world utility is

$$G(\underline{\zeta}) = \sum_t R(\underline{\zeta}_t),$$

where: $R(\underline{\zeta}_t) = \sum_{k=1}^7 \gamma_k(x_k(\underline{\zeta}_t))$; $x_k(\underline{\zeta}_t)$ is the total attendance on night k at week t ; $\gamma_k(y) \equiv \alpha_k y \exp(-y/c)$; and c and each of the $\{\alpha_k\}$ are real-valued parameters. Intuitively, the “world reward” R is the sum of the global “rewards” for each night in each week. It reflects the effects in the bar as the attendance profile of agents changes. When there are too few agents attending some night, the bar suffers from lack of activity and therefore the global reward for that night is low. Conversely, when there are too many agents the bar is overcrowded and the reward for that night is again low. Note that $\gamma_k(\cdot)$ reaches its maximum when its argument equals c .

In these experiments we investigate two different $\vec{\alpha}$ ’s. One treats all nights equally; $\vec{\alpha} = [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]$. The other is only concerned with one night; $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$. In our experiments, $c = 6$ and N is chosen to be 4 times larger than the number of agents necessary to have c agents attend the bar on each of the seven nights, *i.e.*, there are $4 \cdot 6 \cdot 7 = 168$ agents (this ensures that there are no trivial solutions and that for the world utility to be maximized, the agents have to “cooperate”).

As explained below, our microlearning algorithms worked by providing a real-valued “reward” signal to each agent at each week t . Each agent’s reward function is a surrogate for an associated utility function for that agent. The difference between the two functions is that the reward function only reflects the state of the system at one moment in time (and therefore is potentially observable), whereas the utility function reflects the agent’s ultimate goal, and therefore can depend on the full history of that agent across time.

We investigated three agent reward functions. One was based on effect set WLU. The other two were “natural” rewards included for comparison purposes. With d_η the night selected by η , the three rewards are:

$$\begin{aligned} \text{Uniform Division (UD): } r_\eta(\underline{\zeta}_t) &\equiv \gamma_{d_\eta}(x_{d_\eta}(\underline{\zeta}_t))/x_{d_\eta}(\underline{\zeta}_t) \\ \text{Global (G): } r_\eta(\underline{\zeta}_t) &\equiv R(\underline{\zeta}_t) = \sum_{k=1} \gamma_k(x_k(\underline{\zeta}_t)) \\ \text{Wonderful Life (WL): } r_\eta(\underline{\zeta}_t) &\equiv R(\underline{\zeta}_t) - R(\text{CL}_{\underline{\zeta}_{\eta,t}}(\underline{\zeta}_t)) \\ &= \gamma_{d_\eta}(x_{d_\eta}(\underline{\zeta}_t)) - \gamma_{d_\eta}(x_{d_\eta}(\text{CL}_{\underline{\zeta}_{\eta,t}}(\underline{\zeta}_t))) \end{aligned}$$

The conventional UD reward is a natural “naive” choice for the agents’ reward; the total reward on each night gets uniformly divided among the agents attending that night. If we take $g_\eta(\underline{\zeta}) = \sum_t r_\eta(\underline{\zeta}_t)$ (*i.e.*, η ’s utility is an undiscounted sum of its rewards), then for the UD reward $G(\underline{\zeta}) = \sum_\eta g_\eta(\underline{\zeta})$, so that the system is weakly trivial. The original version of the bar problem in the physics literature [50] is the special case where UD reward is used but there are only two “nights” in the week (one of which corresponds to “staying at home”); $\vec{\alpha}$ is uniform; and $\gamma_k(x_k) = x_k \Theta(c_k N - x_k)$ for some vector \vec{c} , taken to equal (.6, .4) in the very original papers. So the reward to agent η is 1 if it attends the bar and attendance is below capacity, or if it stays at home and the bar is over capacity. Reward is 0 otherwise. (In addition, unlike in our COIN-based systems, in the original work on the bar problem the microlearners work by explicitly predicting the bar attendance, rather than by directly modifying behavior to try to increase a reward signal.)

In contrast to the UD reward, providing the G reward at time t to each agent results in all agents receiving the same reward. This is the team game reward function, investigated for example in [60]. For this reward function, the system is automatically factored if we define $g_\eta(\underline{\zeta}) \equiv \sum_t r_\eta(\underline{\zeta}_t)$. However, evaluation of this reward function requires centralized communication concerning all seven nights. Furthermore, given that there are 168 agents, G is likely to have poor learnability as a reward for any individual agent.

This latter problem is obviated by using the WL reward, where the subtraction of the clamped term removes some of the “noise” of the activity of all other agents, leaving only the underlying “signal” of how the agent in question affects the utility. So one would expect that with the WL reward the agents can readily discern the effects of their actions on their rewards. Even though the conditions in Coroll. 1 don’t hold²⁰, this

²⁰The $t = 0$ elements of C_η^{eff} are just $\underline{\zeta}_{\eta,t=0}$, but the contributions of $\underline{\zeta}_{\eta,t=0}$ to G cannot be written as

reasoning accords with the implicit advice of Coroll. 1 under the approximation of the $t = 0$ effect set as $C_\eta^{eff} \approx \underline{\zeta}_{\eta, t \geq 0}$. In other words, it agrees with that corollary's implicit advice under the identification of $\underline{\zeta}_{\eta, t \geq 0}$ as η 's $t = 0$ guessed effect set.

In fact, in this very simple system, we can explicitly calculate the ratio of the WL reward's learnability to that of the G reward, by recasting the system as existing for only a single instant so that $C_\eta^{eff} = \underline{\zeta}_{\eta, 0}$ exactly and then applying Thm. 3. So for example, say that all $\alpha_k = 1$, and that the number of nodes N is evenly divided among the seven nights. The numerator term in Thm. 3 is a vector whose components are some of the partials of G evaluated when $x_k(\underline{\zeta}_0) = N/7$. This vector is $7(N - 1)$ dimensional, one dimension for each of the 7 components of (the unary vector comprising) each node in $\hat{\eta}$. For any particular $\eta' \neq \eta$ and night i , the associated partial derivative is $\sum_k [e^{-x_k(\underline{\zeta}_0)/c} (1 - x_k(\underline{\zeta}_0)/c) \cdot \partial_{\underline{\zeta}_{\eta', 0; i}}(x_k(\underline{\zeta}_0))]$, where as usual " $\underline{\zeta}_{\eta', 0; i}$ " indicates the i 'th component of the unary vector $\underline{\zeta}_{\eta', 0}$. Since $\partial_{\underline{\zeta}_{\eta', 0; i}}(x_k(\underline{\zeta}_0)) = \delta_{i, k}$, for any fixed i and η' , this sum just equals $e^{(-N/7c)} (1 - N/7c)$. Since there are $7(N - 1)$ such terms, after taking the norm we obtain $|e^{(-N/7c)} [1 - N/7c] \sqrt{7(N - 1)}|$.

The denominator term in Thm. 3 is the difference between the gradients of the global reward and the clamped reward. These differ on only $N - 1$ terms, one term for that component of each node $\eta' \neq \eta$ corresponding to the night η attends. (The other $6N - 6$ terms are identical in the two partials and therefore cancel.) This yields $|e^{(-N/7c)} [1 - N/7c] [1 - e^{1/c}(1 - \frac{7}{N-7c})] \sqrt{N-1}|$. Combining with the result of the previous paragraph, our ratio is $|\sqrt{7} \frac{N-7c}{(N-7c)(1-e^{1/c})+7e^{1/c}}| \simeq 11$.

In addition to this learnability advantage of the WL reward, to evaluate its WL reward each agent only needs to know the total attendance on the night it attended, so no centralized communication is required. Finally, although the system won't be perfectly factored for this reward (since in fact the effect set of η 's action at t would be expected to extend a bit beyond $\underline{\zeta}_{\eta, t}$), one might expect that it is close enough to being factored to result in large world utility.

Each agent keeps a seven dimensional Euclidean vector representing its estimate of the reward for attending each night of the week. At the end of each week, the component of this vector corresponding to the night just attended is proportionally adjusted towards the actual reward just received. At the beginning of the succeeding week, the agent picks the night to attend using a Boltzmann distribution with energies given by the components of the vector of estimated rewards, where the temperature in the Boltzmann distribution decays in time. (This learning algorithm is equivalent to Claus and Boutilier's [57] independent learner algorithm for multi-agent reinforcement learning.) We used the same parameters (learning rate, Boltzmann temperature, decay rates, etc.) for all three reward functions. (This is an *extremely* primitive RL algorithm which we only chose for its pedagogical value; more sophisticated RL algorithms are crucial for eliciting high intelligence levels when one is confronted with more complicated learning problems.)

a sum of a $\underline{\zeta}_{\eta, t=0}$ contribution and a $\underline{\zeta}_{\eta, t=0}$ contribution.




Figure 1: Average world reward when $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$ (left) and when $\vec{\alpha} = [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]$ (right). In both plots the top curve is WL, middle is G, and bottom is UD.

Figure 1 presents world reward values as a function of time, averaged over 50 separate runs, for all three reward functions, for both $\vec{\alpha} = [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]$ and $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$. The behavior with the G reward eventually converges to the global optimum. This is in agreement with the results obtained by Crites [60] for the bank of elevators control problem. Systems using the WL reward also converged to optimal performance. This indicates that for the bar problem our approximations of effects sets are sufficiently accurate, *i.e.*, that ignoring the effects one agent’s actions will have on future actions of other agents does not significantly diminish performance. This reflects the fact that the only interactions between agents occurs indirectly, via their affecting each others’ reward values.

However since the WL reward is more learnable than the G reward, convergence with the WL reward should be far quicker than with the G reward. Indeed, when $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$, systems using the G reward converge in 1250 weeks, which is 5 times worse than the systems using WL reward. When $\vec{\alpha} = [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]$ systems take 6500 weeks to converge with the G reward, which is more than *30 times* worse than the time with the WL reward.

In contrast to the behavior for reward functions based on our COIN framework, use of the conventional UD reward results in very poor world reward values, values that deteriorated as the learning progressed. This is an instance of the TOC. For example, for the case where $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$, it is in every agent’s interest to attend the same night — but their doing so shrinks the world reward “pie” that must be divided among all agents. A similar TOC occurs when $\vec{\alpha}$ is uniform. This is illustrated in fig. 2 which shows a typical example of daily attendance figures ($\{x_k(\zeta_t)\}$) for each of the three reward functions for $t = 2000$. In this example optimal performance (achieved with the WL reward) has 6 agents each on 6 separate nights, (thus maximizing the reward on 6 nights), and the remaining 132 agents on one night.

Figure 3 shows how $t = 2000$ performance scales with N for each of the reward signals




Figure 2: Typical daily attendance when $\vec{\alpha} = [1 \ 1 \ 1 \ 1 \ 1 \ 1 \ 1]$ for WL (left), G (center), and UD (right).




Figure 3: Behavior of each reward function with respect to the number of agents for $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$.

for $\vec{\alpha} = [0 \ 0 \ 0 \ 7 \ 0 \ 0 \ 0]$. Systems using the UD reward perform poorly regardless of N . Systems using the G reward perform well when N is low. As N increases however, it becomes increasingly difficult for the agents to extract the information they need from the G reward. (This problem is significantly worse for uniform $\vec{\alpha}$.) Because of their superior learnability, systems using the WL reward overcome this signal-to-noise problem (*i.e.*, because the WL reward is based on the *difference* between the actual state and the state where one agent is clamped, it is much less affected by the total number of agents).

4.5 Macrolearning

In the experiments recounted above, the agents were sufficiently independent that assuming they did not affect each other's actions (when forming guesses for effect sets) allowed the resultant WL reward signals to result in optimal performance. In this section we investigate the contrasting situation where we have initial guesses of effect sets that are quite poor and that therefore result in bad global performance when used with WL rewards. In particular, we investigate the use of macrolearning to correct those guessed effect sets at run-time, so that with the corrected guessed effect sets WL rewards will instead give optimal performance. This models real-world scenarios where the system designer's initial guessed effect sets are poor approximations of the actual associated effect sets and need to be corrected adaptively.

In these experiments the bar problem is significantly modified to incorporate constraints designed to result in poor G when the WL reward is used with certain initial guessed effect sets. To do this we forced the nights actually attended by some of the agents (followers) to agree with those attended by other agents (leaders), regardless of what night those followers “picked” via their microlearning algorithms. (For leaders, picked and actually attended nights were always the same.) We then had the world utility be the sum, over all leaders, of the values of a triply-indexed reward matrix whose indices are the nights that each leader-follower set attends: $G(\underline{\zeta}) = \sum_t \sum_i R_{l_i(t), f1_i(t), f2_i(t)}$ where $l_i(t)$ is the night the i^{th} leader attends in week t , and $f1_i(t)$ and $f2_i(t)$ are the nights attended by the followers of leader i , in week t (in this study, each leader has two followers). We also had the states of each node be one of the integers $\{0, 1, \dots, 6\}$ rather than (as in the bar problem) a unary seven-dimensional vector. This was a bit of a contrivance, since constructions like $\partial_{\underline{\zeta}_{\eta,0}}$ aren’t meaningful for such essentially symbolic interpretations of the possible states $\underline{\zeta}_{\eta,0}$. As elaborated below, though, it was helpful for constructing a scenario in which guessed effect set WLU results in poor performance, *i.e.*, a scenario in which we can explore the application of macrolearning.

To see how this setup can result in poor world utility, first note that the system’s dynamics is what restricts all the members of each triple $(l_i(t), f1_i(t), f2_i(t))$ to equal the night picked by leader i for week t . So $f1_i(t)$ and $f2_i(t)$ are both in leader i ’s actual effect set at week t — whereas the initial guess for i ’s effect set may or may not contain nodes other than $l_i(t)$. (For example, in the bar problem experiments, the guessed effect set does not contain any nodes beyond $l_i(t)$.) On the other hand, G and R are defined for all possible triples $(l_i(t), f1_i(t), f2_i(t))$. So in particular, R is defined for the dynamically unrealizable triples that can arise in the clamping operation. This fact, combined with the leader-follower dynamics, means that for certain R ’s there exist guessed effect sets such that the dynamics assures poor world utility when the associated WL rewards are used. This is precisely the type of problem that macrolearning is designed to correct.

As an example, say each week only contains two nights, 0 and 1. Set $R_{111} = 1$ and $R_{000} = 0$. So the contribution to G when a leader picks night 1 is 1, and when that leader picks night 0 it is 0, independent of the picks of that leader’s followers (since the actual nights they attend are determined by their leader’s picks). Accordingly, we want to have a private utility for each leader that will induce that leader to pick night 1. Now if a leader’s guessed effect set includes both of its followers (in addition to the leader itself), then clamping all elements in its effect set to 0 results in an R value of $R_{000} = 0$. Therefore the associated guessed effect set WLU will reward the leader for choosing night 1, which is what we want. (For this case WL reward equals $R_{111} - R_{000} = 1$ if the leader picks night 1, compared to reward $R_{000} - R_{000} = 0$ for picking night 0.)

However consider having two leaders, i_1 and i_2 , where i_1 ’s guessed effect set consists of i_1 itself together with the two followers of i_2 (rather than together with the two followers of i_1 itself). So neither of leader i_1 ’s followers are in its guessed effect set, while i_1 itself is. Accordingly, the three indices to i_1 ’s R need not have the same value. Similarly, clamping

the nodes in its guessed effect set won't affect the values of the second and third indices to i_1 's R , since the values of those indices are set by i_1 's followers. So for example, if i_2 and its two followers go to night 0 in week 0, and i_1 and its two followers go to night 1 in that week, then the associated guessed effect set wonderful life reward for i_1 for week 0 is $G(\zeta_{t=0}) - G(\text{CL}_{l_{i_1}(0), f_{1i_2}(0), f_{2i_2}(0)}(\zeta_{t=0})) = R_{l_{i_1}(0), f_{1i_1}(0), f_{2i_1}(0)} + R_{l_{i_2}(0), f_{1i_2}(0), f_{2i_2}(0)} - [R_{0, f_{1i_1}(0), f_{2i_1}(0)} + R_{l_{i_2}(0), 0, 0}]$. This equals $R_{111} + R_{000} - R_{011} - R_{000} = 1 - R_{011}$. Simply by setting $R_{011} > 1$ we can ensure that this is negative. Conversely, if leader i_1 had gone to night 0, its guessed effect WLU would have been 0. So in this situation leader i_1 will get a greater reward for going to night 0 than for going to night 1. In this situation, leader i_1 's using its guessed effect set WLU will lead it to make the wrong pick.

To investigate the efficacy of the macrolearning, two sets of separate experiments were conducted. In the first one the reward matrix R was chosen so that if each leader is maximizing its WL reward, but for guessed effect sets that contain none of its followers, then the system evolves to *minimal* world reward. So if a leader incorrectly guesses that some σ is its effect set even though σ doesn't contain both of that leader's followers, and if this is true for all leaders, then we are assured of worst possible performance. In the second set of experiments, we investigated the efficacy of macrolearning for a broader spectrum of reward matrices by generating those matrices randomly. We call these two kinds of reward matrices *worst-case* and *random* reward matrices, respectively. In both cases, if it can modify the initial guessed effect sets of the leaders to include their followers, then macrolearning will induce the system to be factored.

The microlearning in these experiments was the same as in the bar problem. All experiments used the WL personal reward with some (initially random) guessed effect set. When macrolearning was used, it was implemented starting after the microlearning had run for a specified number of weeks. The macrolearner worked by estimating the correlations between the agents' selections of which nights to attend. It did this by examining the attendances of the agents over the preceding weeks. Given those estimates, for each agent η the two agents whose attendances were estimated to be the most correlated with those of agent η were put into agent η 's guessed effect set. Of course, none of this macrolearning had any effect on global performance when applied to follower agents, but the macrolearning algorithm cannot know that ahead of time; it applied this procedure to each and every agent in the system.

Figure 4 presents averages over 50 runs of world reward as a function of weeks using the worst-case reward matrix. For comparison purposes, in both plots the top curve represents the case where the followers are in their leader's guessed effect sets. The bottom curve in both plots represents the other extreme where no leader's guessed effect set contains either of its followers. In both plots, the middle curve is performance when the leaders' guessed effect sets are initially random, both with (right) and without (left) macrolearning turned on at week 500.

The performance for random guessed effect sets differs only slightly from that of hav-




Figure 4: Leader-follower problem with worst case reward matrix. In both plots, every follower is in its leader's guessed effect set in the top curve, no follower is in its leader's guessed effect set in the bottom curve, and followers are randomly assigned to guessed effect sets of the leaders in the middle curve. The two plots are without (left) and with (right) macrolearning at 500 weeks.

ing leaders' guessed effect sets contain none of their followers; both start with poor values of world reward that deteriorates with time. However, when macrolearning is performed on systems with initially random guessed effect sets, the system quickly rectifies itself and converges to optimal performance. This is reflected by the sudden vertical jump through the middle of the right plot at 500 weeks, the point at which macrolearning changed the guessed effect sets. By changing those guessed effect sets macrolearning results in a system that is factored for the associated WL reward function, so that those reward functions quickly induced the maximal possible world reward.




Figure 5: Leader-follower problem for random reward matrices. The ordering of the plots is exactly as in Figure 4. Macrolearning is applied at 2000 weeks, in the right plot.

Figure 5 presents performance averaged over 50 runs for world reward as a function of

weeks using a spectrum of reward matrices selected at random. The ordering of the plots is exactly as in Figure 4. Macrolearning is applied at 2000 weeks, in the right plot. The simulations in Figure 5 were lengthened from those in Figure 4 because the convergence time of the full spectrum of reward matrices case was longer.

In figure 5 the macrolearning resulted in a transient degradation in performance at 2000 weeks followed by convergence to the optimal. Without macrolearning the system’s performance no longer varied after 2000 weeks. Combined with the results presented in Figure 4, these experiments demonstrate that macrolearning induces optimal performance by aligning the agents’ guessed effect sets with those agents that they actually do influence the most.

5 CONCLUSION

Many distributed computational tasks cannot be addressed by direct modeling of the underlying dynamics, or are at best poorly addressed that way due to robustness and scalability concerns. Such tasks should instead be addressed by model-independent machine learning techniques. In particular, Reinforcement Learning (RL) techniques are often a natural choice for how to address such tasks. When — as is often the case — we cannot rely on centralized control and communication, such RL algorithms have to be deployed locally, throughout the system.

This raises the important and profound question of how to configure those algorithms, and especially their associated utility functions, so as to achieve the (global) computational task. In particular we must ensure that the RL algorithms do not “work at cross-purposes” as far as the global task is concerned, lest phenomena like tragedy of the commons occur. How to initialize a system to do this is a novel kind of inverse problem, and how to adapt a system at run-time to better achieve such a global task is a novel kind of learning problem. We call any distributed computational system analyzed from the perspective of such an inverse problem a Collective INtelligence (COIN).

As discussed in the literature review section of this chapter, there are many approaches/fields that address aspects of COINs. These range from multi-agent systems through conventional economics and on to computational economics. (Human economies are a canonical model of a functional COIN.) They range onward to game theory, various aspects of distributed biological systems, and on through physics, active walker models, and recurrent neural nets. Unfortunately, none of these fields seems appropriate as a general approach to understanding COINs.

After this literature review we present a mathematical theory for COINs. We then present experiments on two test problems that validate the predictions of that theory for how best to design a COIN to achieve a global computational task. The first set of experiments involves a variant of Arthur’s famous El Farol Bar problem. The second set instead considers a leader-follower problem that is hand-designed to cause maximal

difficulty for the advice of our theory on how to initialize a COIN. This second set of experiments is therefore a test of the on-line learning aspect of our approach to COINs. In both experiments the procedures derived from our theory, procedures using only local information, vastly outperformed natural alternative approaches, even such approaches that exploited global information. Indeed, in both problems, following the theory summarized in this chapter provides good solutions even when the exact conditions required by the associated theorems hold only approximately.

There are many directions in which future work on COINs will proceed; it is a vast and rich area of research. We are already successfully applying our current understanding of COINs, tentative as it is, to internet packet routing problems. We are also investigating COINs in a more general optimization context where economics-inspired market mechanisms are used to guide some of the interactions among the agents of the distributed system. The goal in this second body of work is to parallelize and solve numerical optimization problems where the concept of an “agent” may not be in the natural definition of the problem. We also intend to try to apply our current COIN framework to the problem of designing high-occupancy toll lanes in vehicular traffic, and to help understand the “design space” necessary for distributed biochemical entities like pre-genomic cells.

Acknowledgements: The authors would like to thank Ann Bell, Michael New, Peter Norvig and Joe Sill for their comments.

References

- [1] H. Abelson and N. Forbes. Morphous-computing techniques may lead to intelligent materials. *Computers in Physics*, 12(6):520–522, 1998.
- [2] M. R. Anderson and T. W. Sandholm. Leveled commitment contracts with myopic and strategic agents. In *Proceedings of the Fifteenth National Conference on Artificial Intelligence*, pages 36–45, 1998.
- [3] K. Arrow and G. Debreu. The existence of an equilibrium for a competitive equilibrium. *Econometrica*, 22:265–290, 1954.
- [4] W. B. Arthur. Complexity in economic theory: Inductive reasoning and bounded rationality. *The American Economic Review*, 84(2):406–411, May 1994.
- [5] W. Ashcroft and N. D. Mermin. *Solid State Physics*. W. B. Saunders, Philadelphia, 1976.
- [6] J.J. Astrom and B. Wittenmark. *Adaptive Control*. Addison–Wesley, 1994.
- [7] C. G. Atkeson. Nonparametric model-based reinforcement learning. In *Advances in Neural Information Processing Systems - 10*, pages 1008–1014. MIT Press, 1998.
- [8] C. G. Atkeson, S. A. Schaal, and A. W. Moore. Locally weighted learning. *Artificial Intelligence Review*, 11:11–73, 1997.

- [9] R. J. Aumann. Correlated equilibrium as an expression of Bayesian rationality. *Econometrica*, 55(1):1–18, 1987.
- [10] R.J. Aumann and S. Hart. *Handbook of Game Theory with Economic Applications, Volumes I and II*. North-Holland Press, 1992.
- [11] R. Axelrod. *The Evolution of Cooperation*. Basic Books, NY, 1984.
- [12] R. Axelrod. *The Complexity of Cooperation: Agent-Based Models of Competition and Collaboration*. Princeton University Press, NJ, 1997.
- [13] P. Bak, S. F. Norrelykke, and M. Shubik. The dynamics of money. (pre-print), 1998.
- [14] P. Bak and K. Sneppen. Punctuated equilibrium and criticality in a simple model of evolution. *Physical Review Letters*, 71(24):4083–4086, 1993.
- [15] P. Bak, C. Tang, and K. Wiesenfeld. Self-organized criticality. *Physical Review A*, 38:364, 1988.
- [16] M. Bando, K. Hasebe, A. Nakayama, A. Shibata, and Y. Sugiyama. Dynamical model of traffic congestion and numerical simulation. *Physical Review E*, 51(2):1035–1042, 1995.
- [17] S. Banks. Exploring the foundations of artificial societies: Experiments in evolving solutions to the iterated N-player prisoner’s dilemma. In R. Brooks and P. Maes, editors, *Artificial Life IV*, pages 337–342. MIT Press, 1994.
- [18] J. B. L. Bard. *Morphogenesis : The Cellular and Molecular Processes of Developmental Anatomy*. Cambridge University Press, 1995.
- [19] M. F. Barnsely, editor. *Chaotic Dynamics and Fraactals*. Academic Press, 1986.
- [20] T. Basar and G.J. Olsder. *Dynamic Noncooperative Game Theory*. Academic Press, 1982.
- [21] T. Bass. Road to ruin. *Discover*, pages 56–61, May 1992.
- [22] M. Batty. Predicting where we walk. *Nature*, 388:19–20, July 1997.
- [23] E. Baum. Toward a model of mind as a laissez-faire economy of idiots. In L. Saitta, editor, *Proceedings of the 13th International Conference on Machine Learning*, pages 28–36. Morgan Kaufman, 1996.
- [24] E. Baum. Toward a model of mind as an economy of agents. *Machine Learning*, 1999. (in Press).
- [25] M. Begon, D.J. Thompshon, and M. Mortimer, editors. *Population Ecology : A Unified Study of Animals and Plants*. Blackwell Science Inc., 1996.

- [26] A. M. Bell and W. A. Sethares. The El Farol problem and the internet: Congestion and coordination failure. In *Fifth International Conference of the Society for Computational Economics*, Boston, MA, 1999.
- [27] J. Bendor and P. Swistak. The evolutionary advantage of conditional cooperation. *Complexity*, 4(2):15–18, 1996.
- [28] J. Berg and A. Engel. Matrix games, mixed strategies, and statistical mechanics. *Physics Review Letters*, 81:4999–5002, 1998. preprint cond-mat/9809265.
- [29] D. Bertsekas and R. Gallager. *Data Networks*. Prentice Hall, Englewood Cliffs, NJ, 1992.
- [30] O. Biham and A. A. Middleton. Self-organization and a dynamical transition in traffic-flow models. *Physical Review A*, 46(10):R6124–6127, 1992.
- [31] K. Binmore. *Fun and Games: A Text on Game Theory*. D. C. Heath and Company, Lexington, MA, 1992.
- [32] L. E. Blume and D. Easley. Optimality and natural selection in markets. pre-print: econ-wpa 9712003.pdf, 1997.
- [33] E. Bonabeau, F. Henaux, S. Guerin, D. Snyders, P. Kuntz, and G. Theraulaz. Routing in telecommunications networks with “smart” and-like agents. (pre-print), 1999.
- [34] E. Bonabeau, A. Sobkowski, G. Theraulaz, and J.-L. Deneubourg. Adaptive task allocation inspired by a model of division of labor of social insects. (pre-print), 1999.
- [35] J. P. Boon. Lattice gas automata: Theory, simulation, implementation. *Journal of Statistical Physics*, 68(3/4), 1991.
- [36] V. S. Borkar, S. Jain, and G. Rangarajan. Collective behaviour and diversity in economic communities: Some insights from an evolutionary game. In *Proceedings of the Workshop on Econophysics*, Budapest, Hungary, 1997.
- [37] V. S. Borkar, S. Jain, and G. Rangarajan. Dynamics of individual specialization and global diversification in communities. *Complexity*, 3(3):50–56, 1998.
- [38] C. Boutilier. Learning conventions in multiagent stochastic domains using likelihood estimates. (pre-print), 1999.
- [39] C. Boutilier. Planning, learning and coordination in multiagent decision processes. (pre-print), 1999.
- [40] C. Boutilier, Y. Shoham, and M. P. Wellman. Editorial: Economic principles of multi-agent systems. *Artificial Intelligence Journal*, 94:1–6, 1997.

- [41] J. Boyan and M. Littman. Packet routing in dynamically changing networks: A reinforcement learning approach. In *Advances in Neural Information Processing Systems - 6*, pages 671–678. Morgan Kaufmann, 1994.
- [42] J. M. Bradshaw, editor. *Software Agents*. MIT Press, 1997.
- [43] J. Breyer, J. Ackermann, and J. McCaskill. Evolving reaction–diffusion ecosystems with self-assembling structures in thin films. *Artificial Life*, 4:25–40, 1998.
- [44] R. A. Brooks. Intelligence without reason. In *Proceedings of the Twelveth International Joint Conference on Artificial Intelligence*, pages 569–595, 1991.
- [45] R. A. Brooks. Intelligence without representation. *Artificial Intelligence*, 47:139–159, 1991.
- [46] T. X. Brown, H. Tong, and S. Singh. Optimizing admission control while ensuring quality of service in multimedia networks via reinforcement learning. In *Advances in Neural Information Processing Systems - 11*. MIT Press, 1999.
- [47] G. Caldarelli, M. Marsili, and Y. C. Zhang. A prototype model of stock exchange. *Europhys. Letters*, 40:479–484, 1997.
- [48] A. R. Cassandra, L. P. Kaelbling, and M. L. Littman. Acting optimally in partially observable stochastic domains. In *Proceedings of the 12th National Conference on Artificial Intelligence*, 1994.
- [49] A. Cavagna. Irrelevance of memory in the minority game. preprint cond-mat/9812215, December 1998.
- [50] D. Challet and Y. C. Zhang. Emergence of cooperation and organization in an evolutionary game. *Physica A*, 246(3-4):407, 1997.
- [51] D. Challet and Y. C. Zhang. On the minority game: Analytical and numerical studies. *Physica A*, 256:514, 1998.
- [52] J. Cheng. The mixed strategy equilibria and adaptive dynamics in the bar problem. Technical report, Santa Fe Institute Computational Economics Workshop, 1997.
- [53] D. R. Cheriton and K. Harty. A market approach to operating system memory allocation. In S.E. Clearwater, editor, *Market-Based Control: A Paradigm for Distributed Resource Allocation*. World Scientific, 1995.
- [54] S. P. M. Choi and D. Y. Yeung. Predictive Q-routing: A memory based reinforcement learning approach to adaptive traffic control. In D. S. Touretzky, M. C. Mozer, and M. E. Hasselmo, editors, *Advances in Neural Information Processing Systems - 8*, pages 945–951. MIT Press, 1996.

- [55] D. J. Christini and J. J. Collins. Using noise and chaos control to control nonchaotic systems. *Physical Review E*, 52(6):5806–5809, 1995.
- [56] A. Church. The calculi of lambda-conversion. In *Annals of Mathematics Studies*, no. 6. Princeton University Press, 1941.
- [57] C. Claus and C. Boutilier. The dynamics of reinforcement learning cooperative multiagent systems. In *Proceedings of the Fifteenth National Conference on Artificial Intelligence*, pages 746–752, June 1998.
- [58] J. E. Cohen and C. Jeffries. Congestion resulting from increased capacity in single-server queueing networks. *IEEE/ACM Transactions on Networking*, 5(2):305–310, 1997.
- [59] J. E. Cohen and F. P. Kelly. A paradox of congestion in a queueing network. *Journal of Applied Probability*, 27:730–734, 1990.
- [60] R. H. Crites and A. G. Barto. Improving elevator performance using reinforcement learning. In D. S. Touretzky, M. C. Mozer, and M. E. Hasselmo, editors, *Advances in Neural Information Processing Systems - 8*, pages 1017–1023. MIT Press, 1996.
- [61] J. P. Crutchfield and J. E. Hanson. Turbulent pattern bases for cellular automata. *Physica D*, 69:279–301, 1993.
- [62] R. Das, M. Mitchell, and J. P. Crutchfield. A genetic algorithm discovers particle-based computation in cellular automata. In Y. Davidor, H.-P. Schwefel, and R. Manner, editors, *Parallel Problem Solving from Nature III*, pages 344–353. Springer-Verlag, 1998.
- [63] J. de Boer, B. Derrida, H. Flyvberg, A. D. Jackson, and T. Wettig. Simple model of self-organized biological evolution. *Physical Review Letters*, 73(6):906–909, 1994.
- [64] M. A. R. de Cara, O. Pla, and F. Guinea. Competition, efficiency and collective behavior in the “El Farol” bar model. preprint cond-mat/9811162 (to appear in European Physics Journal B), November 1998.
- [65] A. de Vany. The emergence and evolution of self-organized coalitions. In M. Gilli, editor, *Computational Methods in Economics*. Kluwer Scientific Publishers, 1999. (to appear).
- [66] D. W. Deamer and J. Oro. The role of lipids in prebiotic structures. *Biosystems*, 12:167–175, 1980.
- [67] W. L. Ditto, S. N. Rauseo, and M. L. Spano. Experimental control of chaos. *Phys. Rev. Let.*, 65:3211, 1990.
- [68] W. L. Ditto and K. Showalter. Introduction: Control and synchronization of chaos. *Chaos*, 7(4):509–511, 1997.

- [69] M. Dorigo and L. M. Gambardella. Ant colonies for the travelling salesman problem. *Biosystems*, 39, 1997.
- [70] M. Dorigo and L. M. Gambardella. Ant colony systems: A cooperative learning approach to the travelling salesman problem. *IEEE Transactions on Evolutionary Computation*, 1(1):53–66, 1997.
- [71] K. E. Drexler. *Nanosystems: Molecular Machinery, Manufacturing, and Computation*. John Wiley and Sons, 1992.
- [72] B. Drossel. A simple model for the formation of a complex organism. preprint adap-org/9811002, November 1998.
- [73] D. Duboule, editor. *Guidebook to the Homeobox Genes*. Oxford University Press, 1995.
- [74] J. Duffy and J. Ochs. Emergence of money as a medium of exchange: an experimental study. Technical report, University of Pittsburgh, 1998.
- [75] J. Eatwell, Milgate M., and Newman P. *The New Palgrave Game Theory*. Macmillan Press, 1989.
- [76] A. A. Economides and J. A. Silvester. Multi-objective routing in integrated services networks: A game theory approach. In *IEEE Infocom '91: Proceedings of the Conference on Computer Communication*, volume 3, 1991.
- [77] L. Edwards, Y. Peng, and J. A. Reggia. Computational models for the formation of protocell structures. *Artificial Life*, 4:61–77, 1998.
- [78] N. Eldredge and S. J. Gould. Punctuated equilibria: An alternative to phyletic gradualism. In J. M. Schopf, editor, *Models in Paleobiology*, pages 82–115. Greeman, Cooper, 1972.
- [79] C. M. Ellison. The Utah TENEX scheduler. *Proceedings of the IEEE*, 63:940–945, 1975.
- [80] J. M. Epstein. Zones of cooperation in demographic prisoner’s dilemma. *Complexity*, 4(2):36–48, 1996.
- [81] J. M. Epstein. *Nonlinear Dynamics, Mathematical Biology, and Social Science*. Addison Wesley, Reading, MA, 1997.
- [82] J. M. Epstein and R. Axtell. *Growing Artificial Societies: Social Sciences from the Bottom Up*. MIT Press, Cambridge, MA, 1996.
- [83] Y. M. Ermoliev and S. D. Flam. Learning in potential games. Technical Report IR-97-022, International Institute for Applied Systems Analysis, June 1997.

- [84] H. Everett III. “relative state” formulation of quantum mechanics. *Reviews of Modern Physics*, 29:424–462, 1957.
- [85] Codd E. F. *Cellular Automata*. Academic Press, New York, 1968.
- [86] J. D. Farmer and J. J. Sidorowich. Exploiting chaos to predict the future and reduce noise. In *Evolution, Learning, and Cognition*, page 277. World Scientific Press, 1988.
- [87] N. Feltovich. Equilibrium and reinforcement learning with private information: An experimental study. pre-print, Dept. of Economics, U. of Houston, July 1997.
- [88] J. Ferber. Reactive distributed artificial intelligence: Principles and applications. In G. O-Hare and N. Jennings, editors, *Foundations of Distributed Artificial Intelligence*, pages 287–314. Wiley, 1996.
- [89] D. F. Ferguson, C. Nikolaou, and Y. Yemini. An economy for flow control in computer networks. In *IEEE Infocom ’89*, pages 110–118, 1989.
- [90] J. Filar and K. Vrieze. *Competitive Markov Decision Processes*. Springer-Verlag, 1997.
- [91] R. Fleischmajer. *Morphogenesis : Cellular Interactions (Annals of the New York Academy of Sciences, V. 857)*. New York Academy of Sciences, 1995.
- [92] W. Fontana. Algorithmic chemistry. In C. G. Langton, C. Taylor, J. D. Farmer, and S. Rasmussen, editors, *Artificial Life II*, pages 159–209. Addison Wesley, 1992.
- [93] C. L. Forgy. RETE: A fast algorithm for the many pattern/many object pattern match problem. *Artificial Intelligence*, 19(1):17–37, 1982.
- [94] S. Forrest, A. S. Perelson, L. Allen, and R. Cherukuri. Self-nonsel self discrimination in a computer. In *Proceedings of the IEEE Symposium on Research in Security and Privacy*, pages 202–212. IEEE Computer Society Press, 1994.
- [95] D. Freedman. *Markov Chains*. Springer-Verlag, 1983.
- [96] B. M. Friedman and F. H. Hahn, editors. *The Handbook of Monetary Economics Vol. I*. North-Holland, Amsterdam, 1990.
- [97] J. W. Friedman. *Game Theory with Applications to Economics*. Oxford University Press, NY, NY, 1986.
- [98] N. Friedman, D. Koller, and A. Pfeffer. Structured representation of complex stochastic systems. In *Proceedings of the 15th National Conference on Artificial Intelligence (AAAI), Madison, Wisconsin*. AAAI Press, 1998.
- [99] D. Fudenberg and D. K. Levine. Steady state learning and Nash equilibrium. *Econometrica*, 61(3):547–573, 1993.

- [100] D. Fudenberg and D. K. Levine. *The Theory of Learning in Games*. MIT Press, Cambridge, MA, 1998.
- [101] D. Fudenberg and J. Tirole. *Game Theory*. MIT Press, Cambridge, MA, 1991.
- [102] Gabora. Autocatalytic closure in a cognitive system: A tentative scenario for the origin of culture. *Psychology*, 9(67), December 1998.
- [103] V. V Gafychuk. Distributed self-regulation induced by negative feedbacks in ecological and economic systems. pre-print, adap-org/98110011, November 1998.
- [104] S. Galam. Spontaneous coalition forming: A model from spin glass. pre-print cond-mat/9901022, January 1999.
- [105] C.W. Gardiner. *Handbook of Stochastic Methods*. Springer-Verlag, New York, NY, 1985.
- [106] M. Gell-Mann and J. B. Hartle. *Physics review D*, 47:3345, 1993.
- [107] C. L. Giles, G. M. Kuhn, and R. J. Williams. Dynamic recurrent networks: theory and applications. *IEEE Transactions on Neural Networks*, 5:153–156, March 1994.
- [108] C. V. Goldman and J. S. Rosenschein. Emergent coordination through the use of cooperative state-changing rules. (pre-print), 1999.
- [109] S. J. Gould and N. Eldredge. Punctuated equilibria: the tempo and mode of evolution reconsidered. *Paleobiology*, 3:115–151, 1977.
- [110] W. Grover. Self organizing broad band transport networks. *Proceedings of the IEEE*, 85(10):1582–1611, 1997.
- [111] O. Guenther, T. Hogg, and B. A. Huberman. Learning in multiagent control of smart matter. In *AAAI-97 Workshop on Multiagent Learning*, 1997.
- [112] O. Guenther, T. Hogg, and B. A. Huberman. Market organizations for controlling smart matter. In *Proceedings of the International Conference on Computer Simulation and Social Sciences*, 1997.
- [113] E. A. Hansen, A. G. Barto, and S. Zilberstein. Reinforcement learning for mixed open-loop and closed loop control. In *Advances in Neural Information Processing Systems - 9*, pages 1026–1032. MIT Press, 1998.
- [114] I Hanski. Be diverse, be predictable. *Nature*, 390:440–441, 1997.
- [115] G. Hardin. The tragedy of the commons. *Science*, 162:1243–1248, 1968.
- [116] A. Hastings. *Population Biology : Concepts and Models*. Springer-Verlag, 1997.
- [117] D. Helbing, J. Keltsch, and P. Molnar. Modelling the evolution of the human trail systems. *Nature*, 388:47–49, July 1997.

- [118] D. Helbing, F. Schweitzer, J. Keltsch, and P. Molnar. Active walker model for the formation of human and animal trail systems. *Physical Review E*, 56(3):2527–2539, 1997.
- [119] D. Helbing and M. Treiber. Jams, waves, and clusters. *Science*, 282:200–201, December 1998.
- [120] D. Helbing and M. Treiber. Phase diagram of traffic states in the presence of inhomogeneities. *Physics Review Letters*, 81:3042, 1998.
- [121] M. Herrmann and B. S. kerner. Local cluster effect in different traffic flow models. *Physica A*, 225:163–168, 1998.
- [122] T. Hogg and B. A. Huberman. Controlling smart matter. *Smart Materials and Structures*, 7:R1–R14, 1998.
- [123] J. Holland and J. H. Miller. Artificial adaptive agents in economic theory. *American Economic Review*, 81:365–370, May 1991.
- [124] J. H. Holland, editor. *Adaptation in Natural and Artificial Systems*. MIT Press, 1993.
- [125] J. J. Hopfield. Neural networks and physical systems with emergent collective computational abilities. *Proceedings National Academy of Science*, 79:2554–2558, April 1982.
- [126] J. J. Hopfield and David W. Tank. Neural computation of decisions in optimization problems. *Biological Cybernetics*, 52:141–152, 1985.
- [127] J. J. Hopfield and David W. Tank. Collective computation with continuous variables. *Disordered Systems and Biological Organization*, 1986.
- [128] B. G. Horne and C. Lee Giles. An experimental comparison of recurrent neural networks. In G. Tesauro, D. S. Touretzky, and T. K. Leen, editors, *Advances in Neural Information Processing Systems - 7*, pages 697–704. MIT Press, 1995.
- [129] M.-T. T. Hsiao and A. A. Lazar. Optimal flow control of multi-class queueing networks with decentralized information. In *IEEE Infocom '89*, pages 652–661, 1987.
- [130] J. Hu and M. P. Wellman. Self-fulfilling bias in multiagent learning. In *Proceedings of the Second International Conference on Multiagent Systems*, pages 118–125, 1996.
- [131] J. Hu and M. P. Wellman. Multiagent reinforcement learning: Theoretical framework and an algorithm. In *Proceedings of the Fifteenth International Conference on Machine Learning*, pages 242–250, June 1998.

- [132] J. Hu and M. P. Wellman. Online learning about other agents in a dynamic multi-agent system. In *Proceedings of the Second International Conference on Autonomous Agents*, pages 239–246, May 1998.
- [133] M. Huber and R. A. Grupen. Learning to coordinate controllers – reinforcement learning on a control basis. In *Proceedings of the 15th International Conference of Artificial Intelligence*, volume 2, pages 1366–1371, 1997.
- [134] B. A. Hubermann, editor. *The Ecology of Computation*. North-Holland, Amsterdam, 1988.
- [135] B. A. Hubermann and S. H. Clearwater. A multi-agent system for controlling building environments. In *Proceedings of the International Conference on Multiagent Systems*, pages 171–176, 1995.
- [136] B. A. Hubermann and T. Hogg. The behavior of computational ecologies. In *The Ecology of Computation*, pages 77–115. North-Holland, 1988.
- [137] M. E. Huhns, editor. *Distributed Artificial Intelligence*. Pittman, London, 1987.
- [138] R. V. Iyer and S. Ghosh. DARYN, a distributed decision-making algorithm for railway networks: Modeling and simulation. *IEEE transaction of Vehicular Technology*, 44(1):180–191, 1995.
- [139] E. T. Jaynes. Information theory and statistical mechanics. *Physical Review*, 106:620, 1957.
- [140] N. R. Jennings, K. Sycara, and M. Wooldridge. A roadmap of agent research and development. *Autonomous Agents and Multi-Agent Systems*, 1:7–38, 1998.
- [141] N. F. Johnson, S. Jarvis, R. Jonson, P. Cheung, Y. R. Kwong, and P. M. Hui. Volatility and agent adaptability in a self-organizing market. preprint cond-mat/9802177, February 1998.
- [142] L. P. Kaelbling, M. L. Littman, and A. W. Moore. Reinforcement learning: A survey. *Journal of Artificial Intelligence Research*, 4:237–285, 1996.
- [143] E. Kalai and E. Lehrer. Rational learning leads to Nash equilibrium. *Econometrica*, 61(5):1019–1045, 1993.
- [144] C. Kappen, K. Schughart, and F.H. Ruddle. Early evolutionary origin of major homeodomain sequence classes. *Genomics*, 18:54–70, 1993.
- [145] S. A. Kauffman. *At Home in the Universe: The Search for the Laws of Self-Organization and Complexity*. Oxford University Press, 1995.
- [146] L. Keller and H. K. Reeve. Familiarity breeds cooperation. *Nature*, 394:121–122, 1998.

- [147] F. P. Kelly. Modeling communication networks, present and future. *Philosophical Trends Royal Society of London A*, 354:437–463, 1996.
- [148] J. O. Kephart. A biologically inspired immune system for computers. In R. Brooks and P. Maes, editors, *Artificial Life IV*, pages 130–139. MIT Press, 1994.
- [149] J. O. Kephart, J. E. Hanson, and J. Sairamesh. Price and niche wars in a free-market economy of software agents. *Artificial Life*, 4:1–13, 1998.
- [150] B. S. Kerner, P. Konhauser, and M. Schilke. Deterministic spontaneous appearance of traffic jams in slightly inhomogeneous traffic flow. *Physical Review E*, 51(6):6243–6246, 1995.
- [151] B. S. Kerner and H. Rehborn. Experimental properties of complexity in traffic flow. *Physical Review E*, 53(5):R4275–4278, 1996.
- [152] T. F. Knight and G. J. Sussman. Cellular gate technology. In *Proceedings of the First International Conference on Unconventional Models of Computation*, Auckland, NZ, January 1998.
- [153] Y. A. Korilis, A. A. Lazar, and A. Orda. Achieving network optima using Stackelberg routing strategies. *IEEE/ACM Transactions on Networking*, 5(1):161–173, 1997.
- [154] Y. A. Korilis, A. A. Lazar, and A. Orda. Capacity allocation under noncooperative routing. *IEEE Transactions on Automatic Control*, 1998.
- [155] Y. A. Korilis, A. A. Lazar, and A. Orda. Avoiding the Braess paradox in noncooperative networks. *Journal of Applied Probability*, 1999. (to appear).
- [156] S. Kraus. Negotiation and cooperation in multi-agent environments. *Artificial Intelligence*, pages 79–97, 1997.
- [157] V. Krishna and P. Motty. Efficient mechanism design. (pre-print), 1997.
- [158] P. Krugman. *Peddling Prosperity: Economic Sense and Nonsense in an Age of Diminished Expectations*. Norton, 1994.
- [159] J. F. Kurose and R. Simha. A microeconomic approach to optimal resource allocation in distributed computer systems. *IEEE Transactions on Computers*, 35(5):705–717, 1989.
- [160] M. Kurz. On the structure and diversity of rational beliefs. *Economic Theory*, 4:877–900, 1994.
- [161] M. Kurz. Rational beliefs and endogenous uncertainty. *Economic Theory*, 8:383–397, 1996.

- [162] R. J. La and V. Anantharam. Optimal routing control: Game theoretic approach. (submitted to IEEE transactions on Automatic Control), 1999.
- [163] C. G. Langton. Introduction. In C. G. Langton, C. Taylor, J. D. Farmer, and S. Rasmussen, editors, *Artificial Life II*, pages 3–25. Addison Wesley, 1992.
- [164] A. A. Lazar, A. Orda, and D. E. Pendarakis. Capacity allocation under noncooperative routing. *IEEE Transactions on Networking*, 5(6):861–871, 1997.
- [165] A. A. Lazar and N. Semret. Design, analysis and simulation of the progressive second price auction for network bandwidth sharing. Technical Report 487-98-21 (Rev 2.10), Columbia University, April 1998.
- [166] T. S. Lee, S. Ghosh, J. Liu, X. Ge, and A. Nerode. A mathematical framework for asynchronous, distributed, decision-making systems with semi-autonomous entities: Algorithm synthesis, simulation, and evaluation. In *Fourth International Symposium on Autonomous Decentralized Systems*, Tokyo, Japan, 1999.
- [167] T. M. Lenton. Gaia and natural selection. *Nature*, 394:439–447, 1998.
- [168] M. L. Littman. Markov games as a framework for multi-agent reinforcement learning. In *Proceedings of the 11th International Conference on Machine Learning*, pages 157–163, 1994.
- [169] M. L. Littman and J. Boyan. A distributed reinforcement learning scheme for network routing. In *Proceedings of the 1993 International Workshop on Applications of Neural Networks to Telecommunications*, pages 45–51, 1993.
- [170] R. D. Luce and H. Raiffa. *Games and Decisions*. Dover Press, 1985.
- [171] J. K. MacKie-Mason and R. V. Hal. Pricing congestible network resources. *IEEE Journal on Selected Areas of Communications*, 13(7):1141–1149, 1995.
- [172] W. G. Macready and D. H. Wolpert. What makes an optimization problem hard? *Complexity*, 5:40–46, 1996.
- [173] W. G. Macready and D. H. Wolpert. Bandit problems and the exploration/exploitation tradeoff. *IEEE Transactions on Evolutionary Computation*, 2:2–22, 1998.
- [174] P. Maes. *Designing Autonomous Agents*. MIT Press, Cambridge, MA, 1990.
- [175] P. Marbach, O. Mihatsch, M. Schulte, and J. Tsisiklis. Reinforcement learning for call admission control and routing in integrated service networks. In *Advances in Neural Information Processing Systems - 10*, pages 922–928. MIT Press, 1998.
- [176] M. Marsili and Y.-C. Zhang. Stochastic dynamics in game theory. preprint cond-mat/9801309, January 1998.

- [177] A. Mas-Colell, M. D. Whinston, and J. R. Green. *Microeconomic Theory*. Oxford University Press, New York, 1995.
- [178] J. Maynard Smith. *Evolution and the Theory of Games*. Cambridge University Press, 1982.
- [179] D. McFarland. Toward robot cooperation. In *From Animals to Animats 3: Proceedings of the Third International Conference on Simulation of Adaptive Behavior*, pages 440–443. MIT Press, 1994.
- [180] B. McMullin and F. J. Varela. Rediscovering computational autopoiesis. pre-print SFI Working paper 07-02-012, 1999.
- [181] M. Mesterton-Gibbons, J.H. Marden, and L.A. Dugatkin. On wars of attrition without assessment. *Journal of Theoretical Biology*, 181:65–83, 1992.
- [182] D. A. Meyer and T. A. Brown. Statistical mechanics of voting. *Physics Review Letters*, 81(8):1718–1721, 1998.
- [183] S. Milgram. The small world problem. *Psychology Today*, 2:60–67, 1967.
- [184] J. H. Miller. The coevolution of automata in the repeated prisoner’s dilemma. *Journal of Economic Behavior and Organization*, 29(1):87–112, 1996.
- [185] J. H. Miller. Evolving information processing organizations. (pre-print), 1996.
- [186] J. H. Miller and J. Andreoni. Auctions with adaptive artificial agents. *Journal of Games and Economic Behavior*, 10:39–64, 1995.
- [187] J. H. Miller, C. Butts, and D. Rode. Communication and cooperation. (pre-print), 1998.
- [188] M. Minsky. *The Society of Mind*. Simon and Schuster, 1988.
- [189] J. Mirrlees. An exploration in the theory of optimal income taxation. *Review of Economic Studies*, 38:175–208, 1974.
- [190] M. Mitchell. *An Introduction to Genetic Algorithms*. MIT Press, Cambridge, MA, 1996.
- [191] M. Mitchell. Computation in cellular automata: A selected review. In T. Gramss, editor, *Nonstandard Computation*. Weinheim, 1998.
- [192] M. Mitchell, J. P. Crutchfield, and P. T. Hraber. Evolving cellular automata to perform computations: Mechanisms and impediments. *Physica D*, 75:361–391, 1994.
- [193] D. Monderer and L. S. Sharpley. Potential games. *Games and Economic Behavior*, 14:124–143, 1996.

- [194] A. W. Moore, C. G. Atkeson, and S. Schaal. Locally weighted learning for control. *Artificial Intelligence Review*, 11:75–113, 1997.
- [195] D. E. Moriarty and P. Langley. Learning cooperative lane selection strategies for highways. In *Proceedings of the Fifteenth National Conference on Artificial Intelligence*, Madison, WI, 1998.
- [196] R. Munos and P. Bourgin. Reinforcement learning for continuous stochastic control problems. In *Advances in Neural Information Processing Systems - 10*, pages 1029–1035. MIT Press, 1998.
- [197] K. Nagel. Experiences with iterated traffic microsimulations in dallas. pre-print adap-org/9712001, December 1997.
- [198] K. Nagel, P. Stretz, M. Pieck, S. Leckey, R. Donnelly, and C. Barrett. TRANSIMS traffic flow characteristics. pre-print adap-org/9710003, October 1997.
- [199] J. F. Nash. Equilibrium points in N -person games. *Proceedings of the National Academy of Sciences of the United States of America*, 36(48-49), 1950.
- [200] R. M. Neal. *Bayesian Learning for Neural Networks, Lecture Notes in Statistics, No. 118*. Springer-Verlag, New York, 1996.
- [201] M. New and A. Pohorille. An inherited efficiencies model for non-genomic evolution. In *Proceedings of the 1st Conference on Modelling and Simulation In Biology, Medicine and Biomedical Engineering*, 1999.
- [202] A. Neyman. Bounded complexity justifies cooperation in the finitely repeated prisoner’s dilemma. *Economics Letters*, 19:227–230, 1985.
- [203] S. I. Nishimura and T. Ikegami. Emergence of collective strategies in a prey-predator game model. *Artificial Life*, 3:243–360, 1997.
- [204] J. Norris. *Markov Chains*. Cambridge University Press, 1998.
- [205] M. A. Nowak and K. Sigmund. Evolution of indirect reciprocity by image scoring. *Nature*, 393:573–577, 1998.
- [206] S. Olafsson. Games on networks. *Proceedings of the IEEE*, 85(10):1556–1562, 1997.
- [207] A. Orda, R. Rom, and M. Sidi. Minimum delay routing in stochastic networks. *IEEE/ACM Transactions on Networking*, 1(2):187–198, 1993.
- [208] J. Oro, E. Sherwood, J. Eichberg, and D. Epps. Formation of phospholipids under primitive earth conditions and the role of membranes in prebiological evolution. In *Light Transducing Membranes, Structure, Function and Evolution*, pages 1–21. Academic Press, New York, 1978.

- [209] J. M. Ostroy and R. M. Starr. The transaction role of money. In B. M. Friedman and F. H. Hahn, editors, *The Handbook of Monetary Economics Vol. I*. North-Holland, Amsterdam, 1990.
- [210] B. A. Pearlmutter. Learning state space trajectories in recurrent neural networks. *Neural Computation*, 1(2):263–269, 1989.
- [211] L. A. Pipes. An operational analysis of traffic dynamics. *Journal of Applied Physics*, 24(3):274–281, 1953.
- [212] A. Pohorille, C. Chipot, M. New, and M.A. Wilson. Molecular modeling of protocellular functions. In L. Hunter and T.E. Klein, editors, *Pacific Symposium on Biocomputing '96*, pages 550–569. World Scientific, 1996.
- [213] G. A. Polls. Stability is woven by complex webs. *Nature*, 395:744–745, 1998.
- [214] M. Potters, R. Cont, and J.-P. Bouchaud. Financial markets as adaptive ecosystems. preprint cond-mat/9609172 v2, June 1997.
- [215] D. Prokhorov and D. Wunsch. Adaptive critic design. *IEEE Transactions on Neural Networks*, 8(5):997–1007, 1997.
- [216] Z. Qu, F. Xie, and G. Hu. Spatiotemporal on-off intermittency by random driving. *Physical Review E*, 53(2):R1301–1304, 1996.
- [217] T. S. Ray. An approach to the synthesis of life. In C. G. Langton, C. Taylor, J. D. Farmer, and S. Rasmussen, editors, *Artificial Life II*, pages 371–408. Addison Wesley, 1992.
- [218] T. S. Ray. Evolving multi-cellular artificial life. In R. Brooks and P. Maes, editors, *Artificial Life IV*, pages 283–288. MIT Press, 1994.
- [219] T. S. Ray. Evolution, complexity, entropy and artificial life. *Physica D*, 1995.
- [220] F. Reif. *Fundamentals of Statistical and Thermal Physics*. McGraw-Hill, 1965.
- [221] E. Rich and K. Knight. *Artificial Intelligence*. McGraw-Hill, Inc., 2 edition, 1991.
- [222] A.E. Roth and I. Erev. Learning in extensive-form games: experimental data and simple dynamic models in the intermediate term. *Games and economic behavior*, 8:164–212, 1995.
- [223] G. Rozenberg and A. Salomaa. *Lindenmeyer Systems: Impacts on Theoretical Computer Science, Computer Graphics and Developmental Biology*. Springer-Verlag, 1992.
- [224] D. E. Rumelhart and J. L. McClelland, editors. *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*. Bradford Books/MIT Press, Cambridge, MA, 1986.

- [225] S. Russell. Learning agents for uncertain environments. In *Proceedings of Computational Learning Theory '98*. ACM Press, 1998.
- [226] A. Samuel. Some studies in machine learning using the game of checkers. *IBM Journal of Research and Development*, 3:210–229, 1959.
- [227] T. Sandholm and R. Crites. Multiagent reinforcement learning in the iterated prisoner’s dilemma. *Biosystems*, 37:147–166, 1995.
- [228] T. Sandholm, K. Larson, M. Anderson, O. Shehory, and F. Tohme. Anytime coalition structure generation with worst case guarantees. In *Proceedings of the Fifteenth National Conference on Artificial Intelligence*, pages 46–53, 1998.
- [229] T. Sandholm and V. R. Lesser. Issues in automated negotiations and electronic commerce: extending the contract net protocol. In *Proceedings of the Second International Conference on Multi-Agent Systems*, pages 328–335. AAAI Press, 1995.
- [230] T. Sandholm and V. R. Lesser. Coalitions among computationally bounded agents. *Artificial Intelligence*, 94:99–137, 1997.
- [231] S. Sastry and M. Bodson. *Adaptive Control, Stability, Convergence, and Robustness*. Prentice Hall, 1989.
- [232] R. Savit, R. Manuca, and R. Riolo. Adaptive competition, market efficiency, phase transitions and spin-glasses. preprint cond-mat/9712006, December 1997.
- [233] A. Schaerf, Y. Shoham, and M. Tennenholtz. Adaptive load balancing: A study in multi-agent learning. *Journal of Artificial Intelligence Research*, 162:475–500, 1995.
- [234] J. Schmidhuber, J. Zhao, and N. N. Schraudolph. Reinforcement learning with self-modifying policies. In S. Thrun and L. Pratt, editors, *Learning to Learn*, pages 293–309. Kluwer, 1997.
- [235] J. Schmidhuber, J. Zhao, and M. Wiering. Shifting inductive bias with success-story algorithm, adaptive Levin search, and incremental self-improvement. *Machine Learning*, 28:105–130, 1997.
- [236] R. Schoonderwoerd, O. Holland, and J. Bruten. Ant-like agents for load balancing in telecommunication networks. In *Autonomous Agents 97*, pages 209–216. MIT Press, 1997.
- [237] M. Schreckenberg, A. Schadschneider, K. Nagel, and N. Ito. Discrete stochastic models for traffic flow. *Physical Review E*, 51(4):2939–2949, 1995.
- [238] J. Schull. Are species intelligent? *Behavioral and Brain Sciences*, 13:63–108, 1990.

- [239] S. Sen. *Multi-Agent Learning: Papers from the 1997 AAAI Workshop (Technical Report WS-97-03)*. AAAI Press, Menlo Park, CA, 1997.
- [240] S. Sen, M. Sekaran, and J. Hale. Learning to coordinate without sharing information. (pre-print), 1999.
- [241] W. A. Sethares and A. M. Bell. An adaptive solution to the El Farol problem. In *Proceedings of the Thirty-Sixth Annual Allerton Conference on Communication, Control, and Computing*, Allerton, IL, 1998. (Invited).
- [242] R. Sethi. Stability of equilibria in games with procedural rational players. pre-print, Dept of Economics, Columbia University, November 1998.
- [243] S. J. Shenker. Making greed work in networks: A game-theoretic analysis of switch service disciplines. *IEEE Transactions on Networking*, 3(6):819–831, 1995.
- [244] Y. Shoham and K. Tanaka. A dynamic theory of incentives in multi-agent systems. In *Proceedings of the International Joint Conference on Artificial Intelligence*, 1997.
- [245] J. Sidel, P. M. Aoki, S. Barr, A. Sah, C. Staelin, M. Stonebreaker, and Yu A. Data replication in mariposa. In *Proceedings of the 12th International Conference on Data Engineering*, 1996.
- [246] S. Sinha and N. Gupte. Adaptive control of spatially extended systems: targeting spatiotemporal patterns and chaos. *Physical Review E*, 58(5):R5221–5224, 1998.
- [247] A. R. Smith. Simple nontrivial self-reproducing machines. In C. G. Langton, C. Taylor, J. D. Farmer, and S. Rasmussen, editors, *Artificial Life II*, pages 709–725. Addison Wesley, 1992.
- [248] A. Somayagi, S. Hofmeyr, and S. Forrest. Principles of computer immune system. *New Security Paradigms Workshop*, 1999. in press.
- [249] W. Stallings. *Data and Computer Communications*. MacMillan Publishing Co., New York, 1994.
- [250] R. M. Starr. *General Equilibrium Theory*. Cambridge University Press, Cambridge, UK, 1997.
- [251] R. M. Starr and M. B. Stinchcombe. Exchange in a network of trading posts. In K. J. Arrow and G. Chichilnisky, editors, *Markets, Information and Uncertainty: Essays in Economic Theory in Honor of Kenneth Arrow*. Cambridge University Press, 1998.
- [252] J. Stein. Critical exponents of the $u(n)$ vector spin glasses. *Europhysics Letters*, 34(9):717–721, 1996.

- [253] J. Stein. Critical properties of a spin glass with anisotropic dzyaloshinskii-moriya interactio. *J. Phys. A*, 29:963–971, 1996.
- [254] W. J. Stewart. *Introduction to the Numerical Solution of Markov Chains*. Princeton University Press, 1995.
- [255] M. Stonebreaker, P. M. Aoki, R. Devine, W. Litwin, and M. Olson. Mariposa: A new architecture for distributed data. In *Proceedings of the 10th International Conference on Data Engineering*, 1994.
- [256] S. H. Strogatz. *Nonlinear Dynamics and Chaos : With Applications in Physics, Biology, Chemistry, and Engineering*. Perseus Press, 1994.
- [257] D. Subramanian, P. Druschel, and J. Chen. Ants and reinforcement learning: A case study in routing in dynamic networks. In *Proceedings of the Fifteenth International Conference on Artificial Intelligence*, pages 832–838, 1997.
- [258] R. S. Sutton. Learning to predict by the methods of temporal differences. *Machine Learning*, 3:9–44, 1988.
- [259] R. S. Sutton and A. G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, Cambridge, MA, 1998.
- [260] K. Sycara. Multiagent systems. *AI Magazine*, 19(2):79–92, 1998.
- [261] G. Szabo and C. Toke. Evolutionary prisoner’s dilemma game on a square lattice. *Physical Review E*, 58(1):69–73, 1998.
- [262] M. Tabor. *Chaos and Integrability in Nonlinear Dynamics : An Introduction*. John Wiley and Sons, 1989.
- [263] G. Tesauro. Practical issues in temporal difference learning. *Machine Learning*, 8:33–53, 1992.
- [264] P. Tucker and F. Berman. On market mechanisms as a software techniques. Technical Report CS96–513, University of California, San Diego, December 1996.
- [265] K. Tumer and D. H. Wolpert. Avoiding Braess’ paradox through collective intelligence. (in preparation), 1999.
- [266] A. M. Turing. The chemical basis of morphogenesis. *Phil. Trans. Royal Society*, B237:37–72, 1952.
- [267] W. Vickrey. Counterspeculation, auctions and competitive sealed tenders. *Journal of Finance*, 16:8–37, 1961.
- [268] J. von Neuman. *The Theory of Self-Replicating Automata (Ed. Burks, A. W.)*. University of Illinois Press, Urbana, IL, 1966. (Work performed by J. von Neuman in 1952-53).

- [269] C. A. Waldspurger, T. Hogg, B. A. Huberman, J. O. Kephart, and W. S. Stornetta. Spawn: A distributed computational economy. *IEEE transactions of Software engineering*, 18(2):103–117, 1992.
- [270] J. Walrand and P. Varaiya. *High-Performance Communication Networks*. Morgan Kaufmann, San Fransisco, CA, 1996.
- [271] L. Walras. *Elements d’Economie Politique Pure*. L. Corbaz, Lausanne, 1874. (translated by W. Jaffe as *Elements of Pure Economics*, Homewood, IL”, 1954).
- [272] C. Watkins and P. Dayan. Q-learning. *Machine Learning*, 8(3/4):279–292, 1992.
- [273] D. J. Watts and S. H. Strogatz. Collective dynamics of ‘small world’ networks. *Nature*, 393:440442, 1998.
- [274] R. Weiss, G. Homsy, and R. Nagpal. Programming biological cells. In *Proceedings of the 8th International Conference on Architectural Support for Programming Languages and Operating Systems*, San Jose, NZ, 1998.
- [275] M. P. Wellman. A market-oriented programming environment and its application to distributed multicommodity flow problems. In *Journal of Artificial Intelligence Research*, 1993.
- [276] M. P. Wellman. A computational market model for distributed configuration design. In *Proceedings of the 12th National Conference on Artificial Intelligence*, 1994.
- [277] S. Wolfram. Statistical mechanics of cellular automata. *Reviews of Modern Physics*, 55:601–644, 1983.
- [278] S. Wolfram. *Cellular Automata and Complexity – Collected Papers*. Addison Wesley, 1994.
- [279] D. H. Wolpert. The relationship between PAC, the statistical physics framework, the Bayesian framework, and the VC framework. In *The Mathematics of Generalization*, pages 117–215. Addison–Wesley, 1995.
- [280] D. H. Wolpert. The existence of a priori distinctions between learning algorithms. *Neural Computation*, 8:1391–1420, 1996.
- [281] D. H. Wolpert. The lack of a priori distinctions between learning algorithms. *Neural Computation*, 8:1341–1390, 1996.
- [282] D. H. Wolpert. Reconciling Bayesian and non-Bayesian analysis. In *Maximum Entropy and Bayesian Methods*, pages 79–86. Kluwer Acedemic Publishers, 1996.
- [283] D. H. Wolpert. Collective intelligence. (pre-print), 1999.

- [284] D. H. Wolpert. A mathematics of bounded rationality. (in preparation), 1999.
- [285] D. H. Wolpert and W. G. Macready. No free lunch theorems for optimization. *IEEE Transactions on Evolutionary Computation*, 1(1):67–82, 1997. **Best Paper Award.**
- [286] D. H. Wolpert and W. G. Macready. An empirically observable measure of complexity. In *New England Complex Systems Institute Inaugural Conference Proceeding*, 1999. in press.
- [287] D. H. Wolpert, M. H. New, and A. M. Bell. Distorting reward functions to improve reinforcement learning. (pre-print), 1999.
- [288] D. H. Wolpert, K. Tumer, and J. Frank. Using collective intelligence to route internet traffic. In *Advances in Neural Information Processing Systems - 11*. MIT Press, 1999.
- [289] D. H. Wolpert, K. Wheeler, and K. Tumer. General principles of learning-based multi-agent systems. In *Proceedings of the Third International Conference of Autonomous Agents*, pages 77–83, 1999.
- [290] S. Wright. The roles of mutation, inbreeding, crossbreeding and selection in evolution. *Proceedings of the XI International Congress of Genetics*, 8:209–222, 1932.
- [291] H. P. Young. The evolution of conventions. *Econometrica*, 61(1):57–84, 1993.
- [292] E. Zambrano. Rationalizable bounded rational behavior. (pre-print), 1999.
- [293] D. H. Zanette and A. S. Mikhailov. Coherence and clustering in ensembles of neural networks. (pre-print), 1999.
- [294] W. Zhang and T. G. Dietterich. Solving combinatorial optimization tasks by reinforcement learning: A general methodology applied to resource-constrained scheduling. *Journal of Artificial Intelligence Research*, 1991.
- [295] Y. C. Zhang. Modeling market mechanism with evolutionary games. preprint cond-mat/9803308, 1998.
- [296] G. Zlotkin and J. S. Rosenschein. Coalition, cryptography, and stability: Mechanisms for coalition formation in task oriented domains. (pre-print), 1999.
- [297] W. H. Zurek and J. P. Paz. Decoherence, chaos, and the second law. *Physics Review Letters*, 72:2508, 1994.